



## Lecture 5



Topic1: System Call Execution

Topic 2: Threads

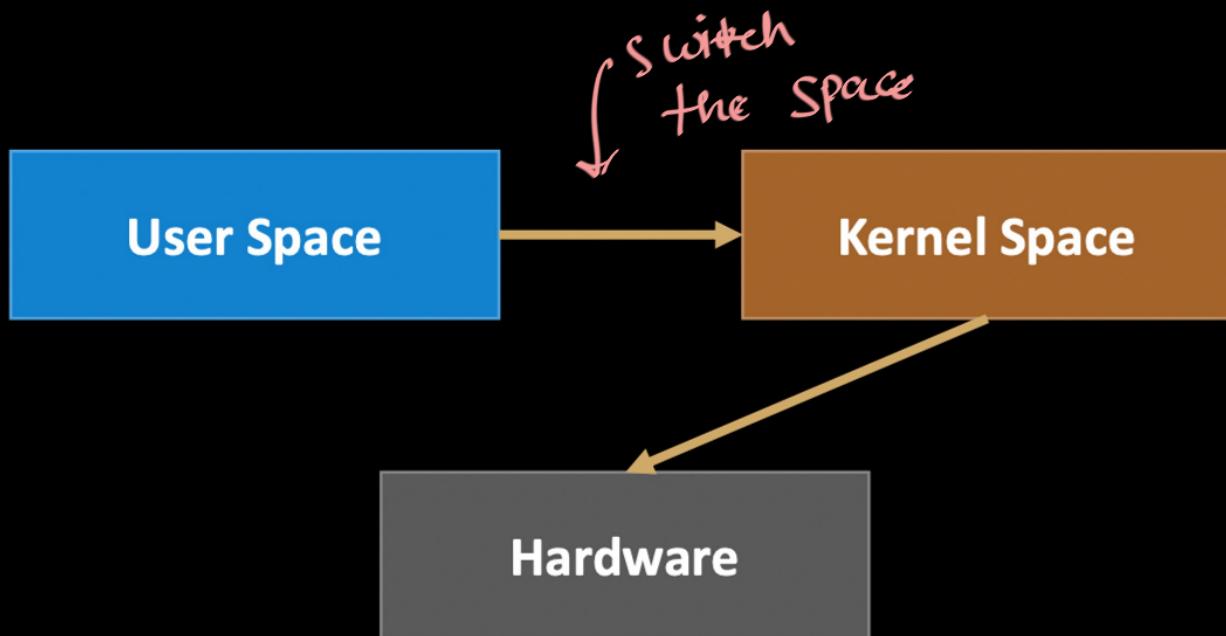


## Kernel Space vs User Space

- User space
  - Less privileged memory space where user processes execute
- Kernel space
  - Privileged memory space where the OS main process resides
  - No User application should be able to change



# Operating Systems





## What is a system call?





## What is a system call?

- Various possible answers, from different perspectives





# Operating Systems

- Various possible answers, from different perspectives
- Answer 1: **request to kernel to perform a service**
  - Open a file
  - Execute a new program
  - Create a new process
  - Send a message to another process



- Various possible answers, from different perspectives
- Answer 1: **request to kernel to perform a service**
  - Open a file
  - Execute a new program
  - Create a new process
  - Send a message to another process
- Answer 2 (programmer's perspective): “call a function”
  - `fd = open("myfile.txt");`
  - System call **looks like** any other function call



- Answer 3: entry point providing **controlled mechanism to execute kernel code**
- User-space programs **can't** call functions inside kernel
- Syscall = one of few mechanisms by which program can ask to execute kernel code
- Set of system calls is:
  - Operating-system specific
    - Can't run Linux binaries on another OS, and vice versa
  - Limited/strictly defined by OS
    - Linux kernel provides 400+ syscalls

Mov Load

hard



## System Call Execution



# System Call Execution

- Each system call has a unique numeric identifier
- OS has a system call table that maps numbers to functionality requested

fork()

57

57	address
----	---------



# Operating Systems

; System call numbers are listed in "unistd.h"

Here we're calling the "fork();", which creates a new process.

```
mov rax, 57; the system call number of "fork".  
syscall; Issue the system call
```

Assembly

fork();

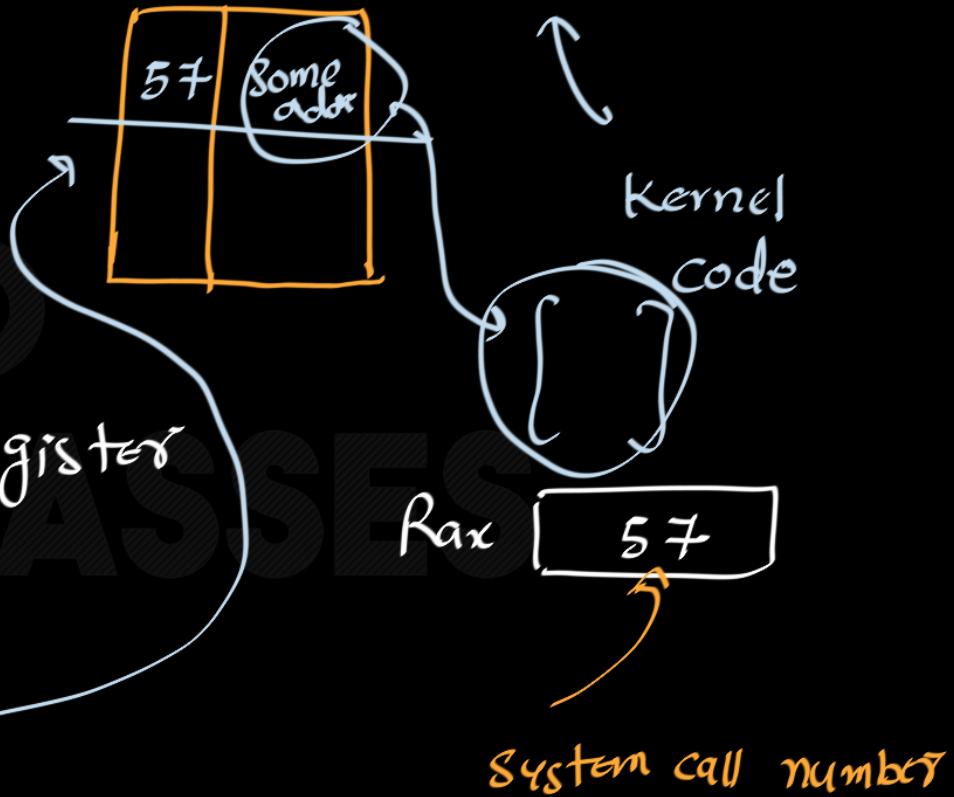
Compiler

Some CPU register

mov rax, 57;  
syscall;

we have system call here.

System call table



```
mov rax, 62;  
Mov R1 "hello"  
syscall;
```



printf("hello");

we will put in the stack

ARGUMENTS

sys\_write( )



{  
}  
3  
}

put the return value in  
specific location

Sat�am Naik to Everyone 8:12 PM

SN

store hello in some other register  
before doing syscall



## System Call Execution (Contd..)

- Each system call has a unique numeric identifier
- OS has a system call table that maps numbers to functionality requested
- When invoking a system call, user places system call number and associated parameters in an specific register, then executes the syscall instruction



## Steps in system call execution





# Operating Systems

① Program calls wrapper function in C library      *printf, fork*

② Wrapper function

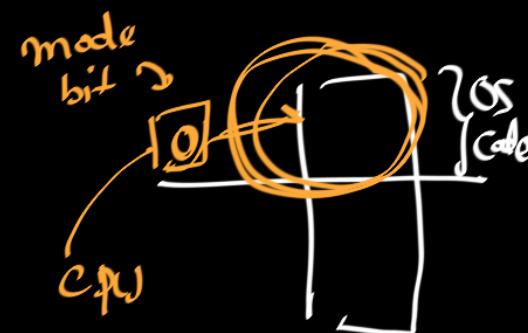
- Packages syscall arguments into registers *or stack*
- Puts (unique) syscall number into a register

③ Wrapper flips CPU to kernel mode (*user-mode*  $\Rightarrow$  *kernel-mode*)

- Execute special machine instruction (e.g., syscall )
- Main effect: CPU can now touch memory marked as accessible only in kernel mode

④ Kernel executes syscall handler:

- Invokes **service routine** corresponding to syscall number
  - **Do the real work**, generate result status
- Places return value from service routine in a register
- Switches back to user mode, passing control back to wrapper
  - (*kernel-mode*  $\Rightarrow$  *user-mode*)

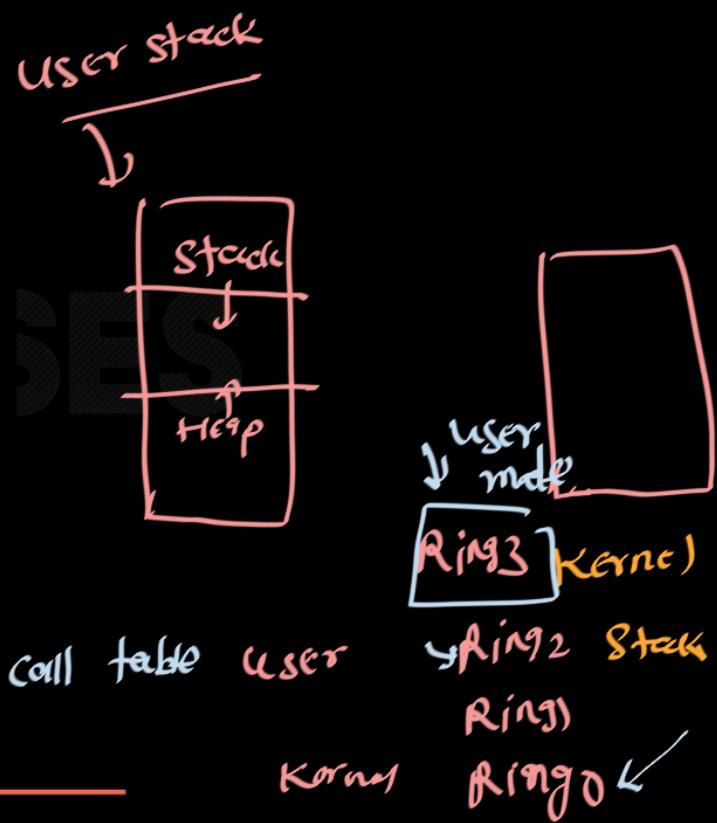


# Inside the SYSCALL instruction

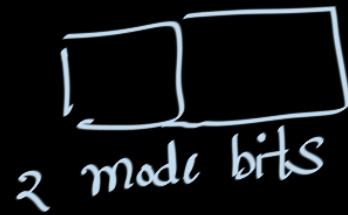
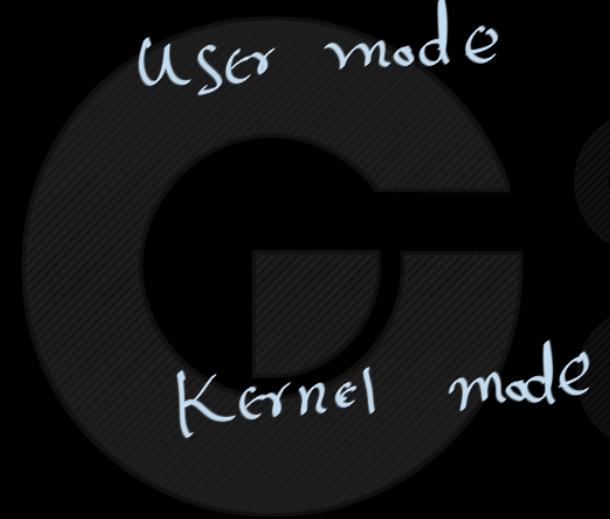
- Saves the old (user) SP value
- Switches the SP to the kernel stack
- Saves the old (user) PC value (= return addr)
- Saves the old privilege mode
- Sets the new privilege mode to 0
- Sets the new PC to the kernel syscall handler

↳ using the system call table user

<http://www.cs.cornell.edu/courses/cs3410/2018fa/schedule/slides/14-ecf.pdf>



optional



$$a = b + c$$

Ring3

Ring2 Open file

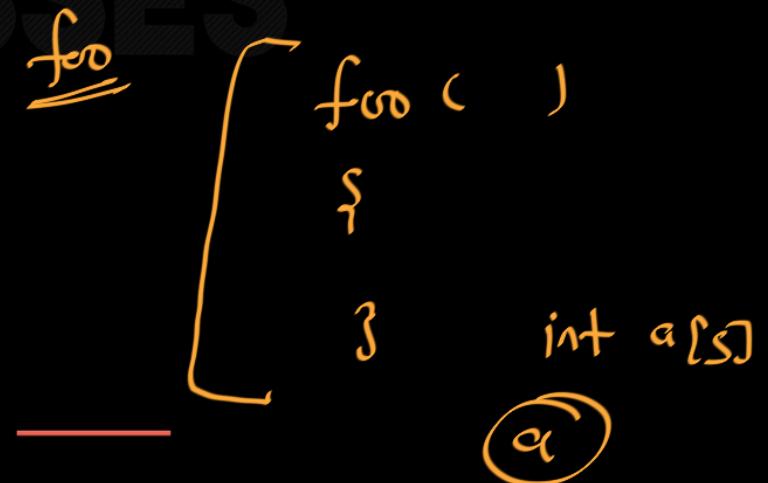
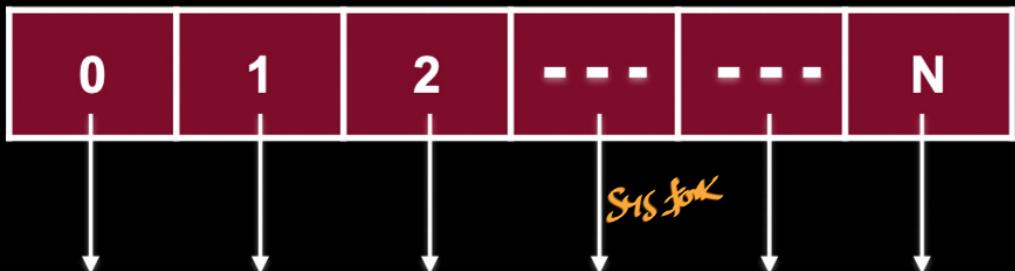
Ring1

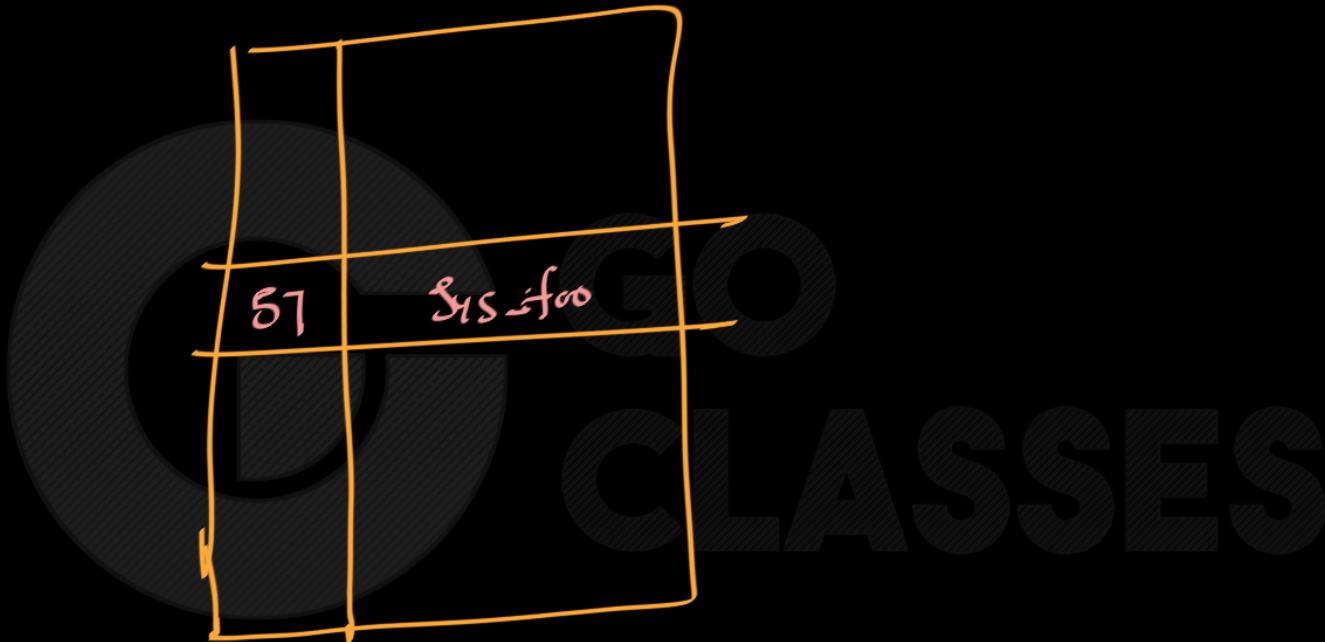
Ring0 point on monitor



# System Call table

- Protected entry points into the kernel for each system call  
We don't want application to randomly jump into any part of the OS code.
- Syscall table is usually implemented as an array of function pointers, where each function implements one system call
- Syscall table is indexed via system call number





## Summary

when ever we see "Syscall" instruction then

what are the steps?

printf( ) ← wrapped function → essentially calling  
syscall / trap / int

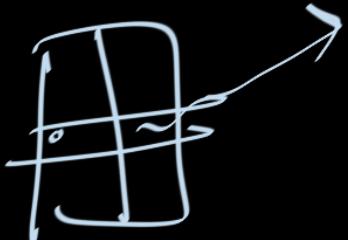
CPU is in infinite cycle of fetch-decode-execute

ISA ↗

Cpu will pause the current execution (after

decoding "Syscall")

Interrupt service routine



Page fault

exception  
/ O }

} Run time  
error  
(segmentation  
fault)

Cpu will pause the current execution (after

decoding "Syscall")

in user code

IVT

(interrupt vector  
table)

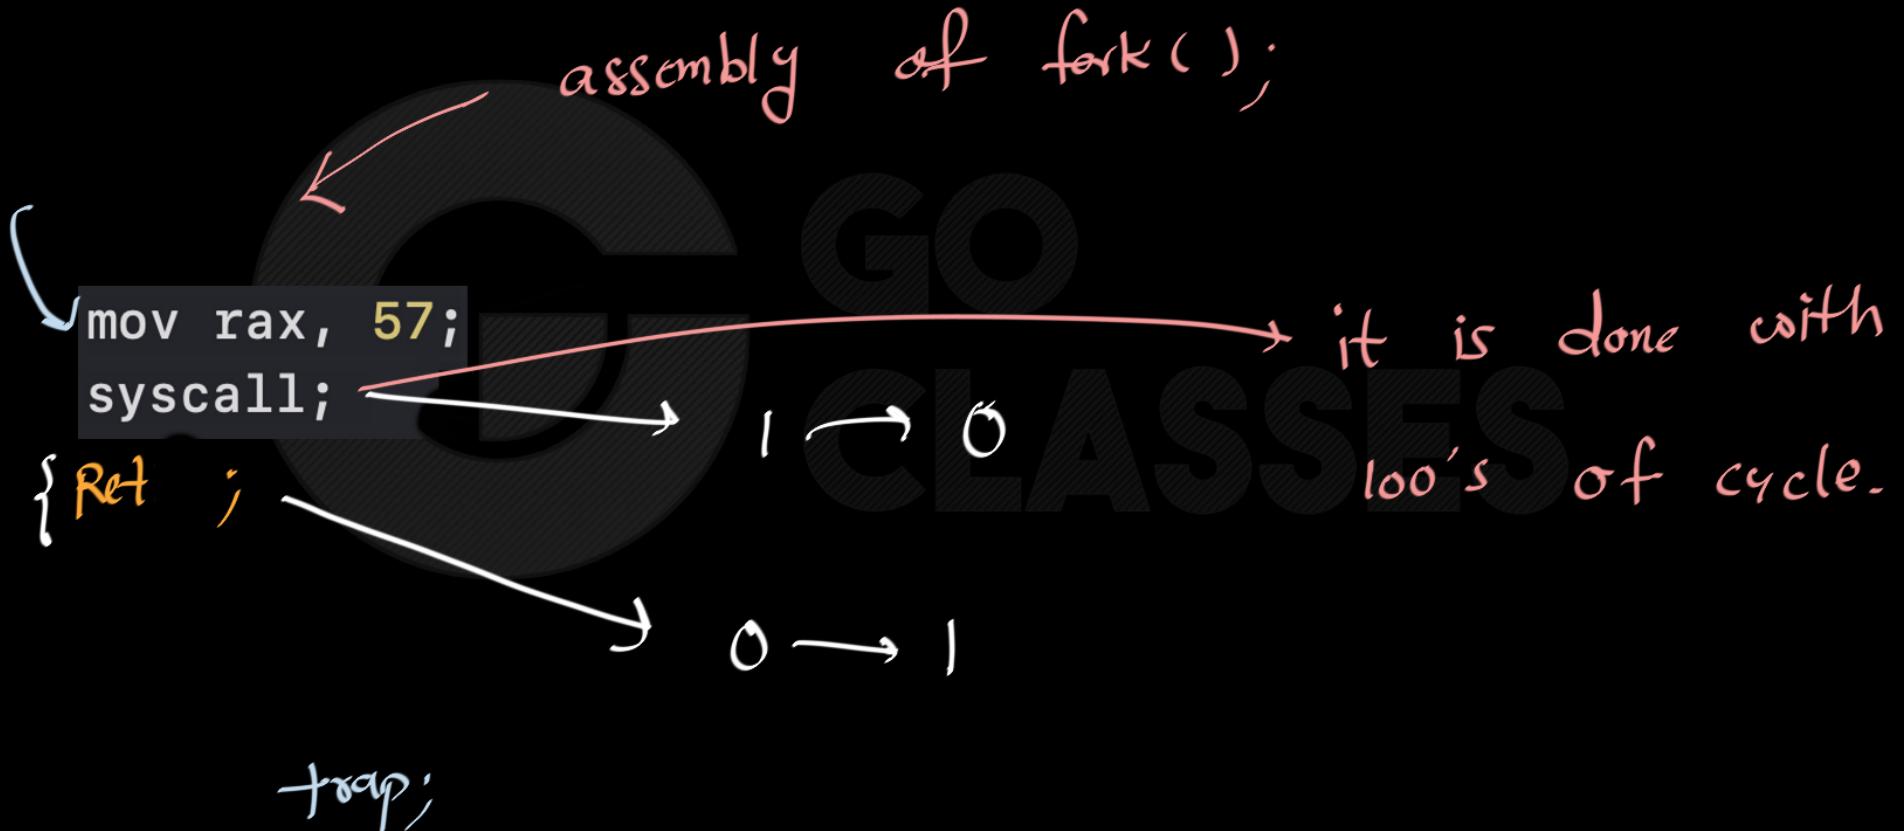
- ⇒ Save the registers value
- ⇒ Switch the mode
- ⇒ load Rax value and go to system call table
- ⇒ find the address of kernel code ← Interrupt service routine (ISR)

- if Kernel function requires the arguments then we check specific location for the arguments
- execute the Kernel Code (interrupt service routine code)
- Placing the return values in some specific location
- ⇒ Reti ← it will switch mode bit from 0 to 1  
in user code

array name of system call table is

Sys\_table

$pc = Sys_table [ \text{Rax} ]$





## Overhead

System calls **are very slow**.

They can take tens to **hundreds of thousands** of clock cycles.

This is due to:

- Changing protection domains    
- Validating arguments
- Adjusting memory mappings
- Cache effects
- ...

ES

Programs should **make fewer system calls when practical**.



<https://cse.buffalo.edu/~eblanton/course/cse220-2022-2f/materials/37-kernel.pdf>



# Operating Systems

**QUESTION 13D.** True or false? System calls are just as fast as function calls. Explain briefly.



<https://cs61.seas.harvard.edu/site/2018/Midterm/>





# Operating Systems

**QUESTION 13D.** True or false? System calls are just as fast as function calls. Explain briefly.

3pts False. System calls are slower because they transfer to kernel. *and many other overheads*

[Hide solution](#)



<https://cs61.seas.harvard.edu/site/2018/Midterm/>



## Types of System Calls

- Process control
  - create process, terminate process
  - end, abort
  - load, execute
  - get process attributes, set process attributes
- File management
  - create file, delete file
  - open, close file
  - read, write

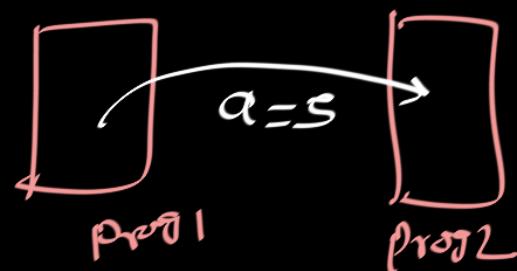


## Types of System Calls (Contd..)

- Device management
  - request device, release device
- Information maintenance
  - get time or date, set time or date

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Process id of current executing process





we can check coA interpret questions  
and try to solve all questions related  
to interrupt execution.

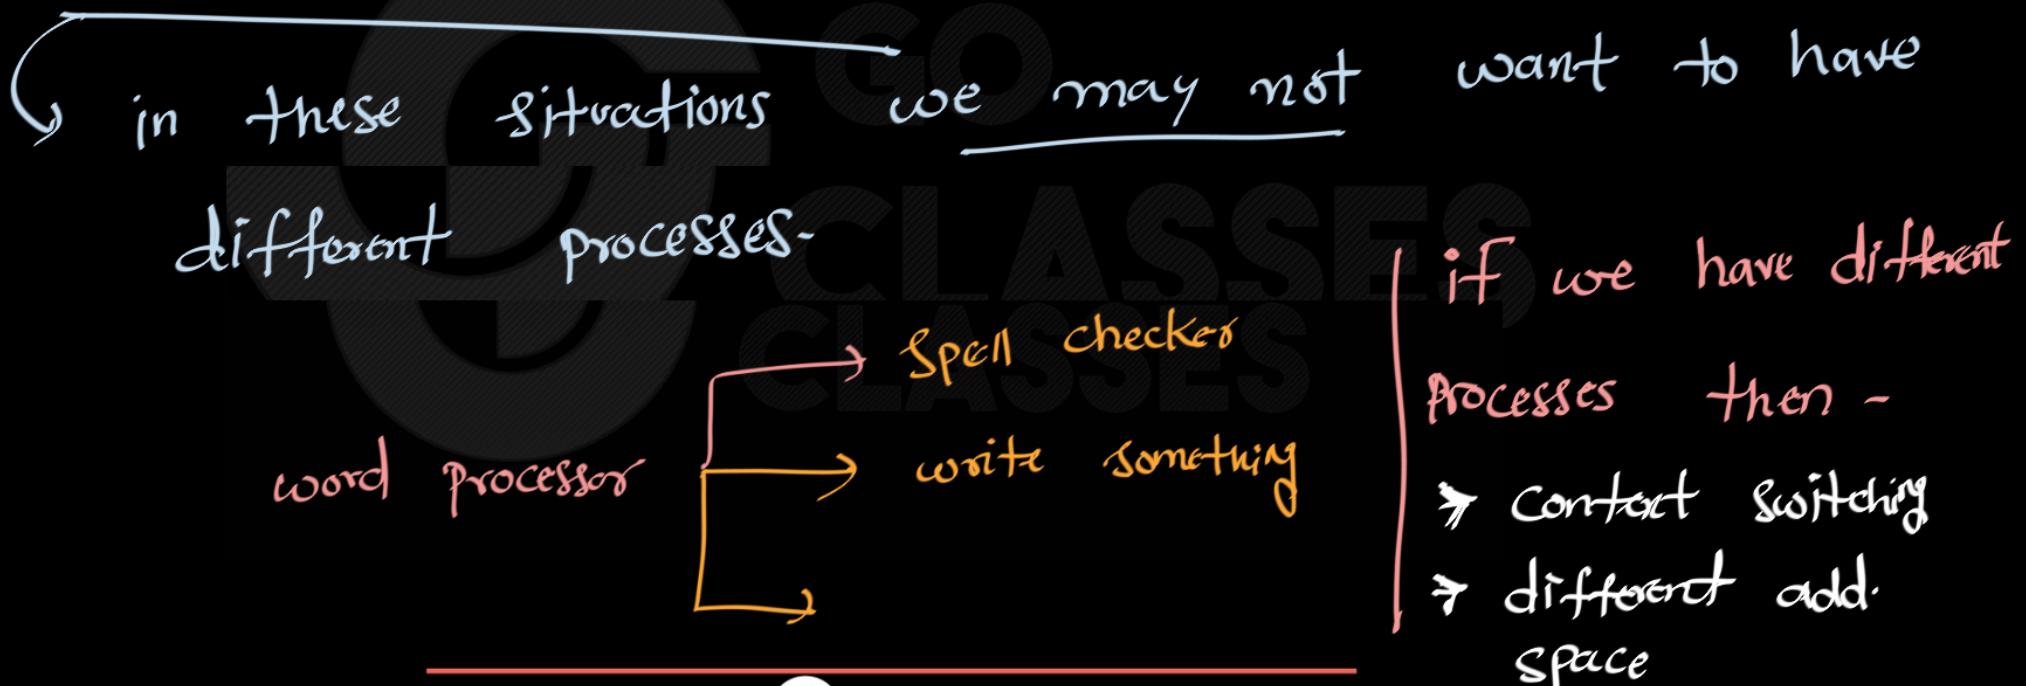


New Topic:  
Threads



# Threads

In certain situations, a single application may be required to perform several similar tasks.



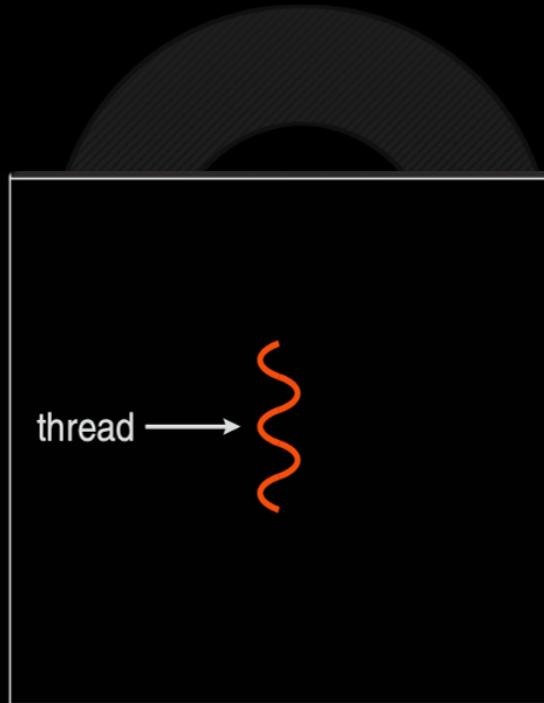
for the similar tasks, if we have different processes (instead of threads) then

Ashutosh Jadhav to Everyone 9:04 PM

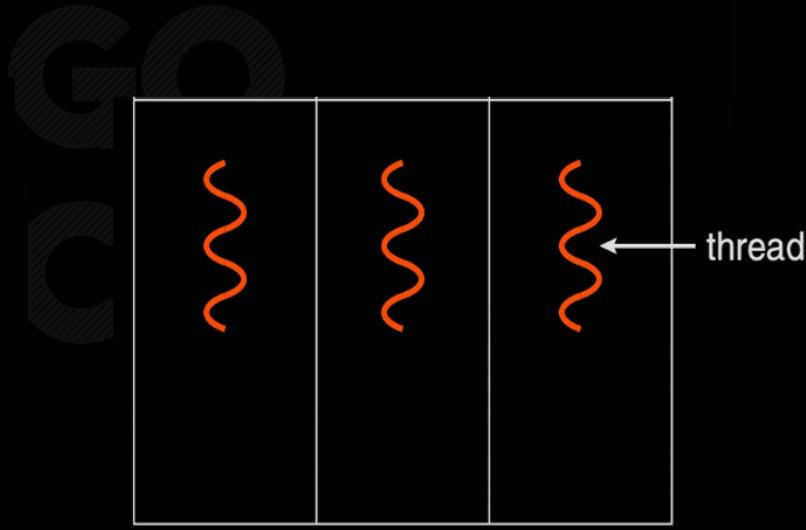
AJ inter process communication required?

IPC (inter process comm)  
is required

process context switching  
is bit costlier than  
thread context switching



single-threaded process



multithreaded process

CPU



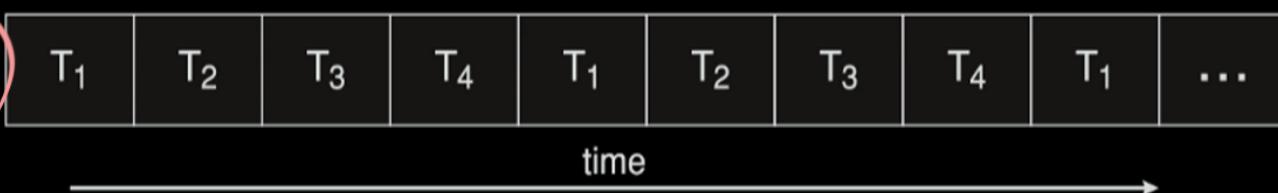
## Thread: Concurrency vs. Parallelism

- Threads are about concurrency and parallelism



# Thread: Concurrency vs. Parallelism

- Threads are about **concurrency** and **parallelism**

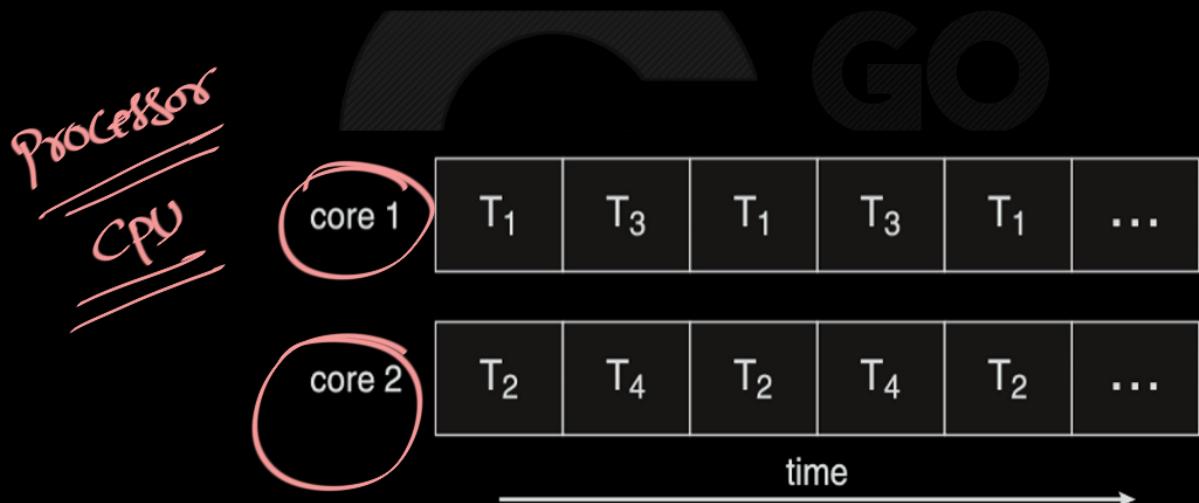


$T_1, T_2, T_3, T_4$

- Concurrent execution on **single-core system**

# Thread: Concurrency vs. Parallelism

- Threads are about **concurrency** and **parallelism**



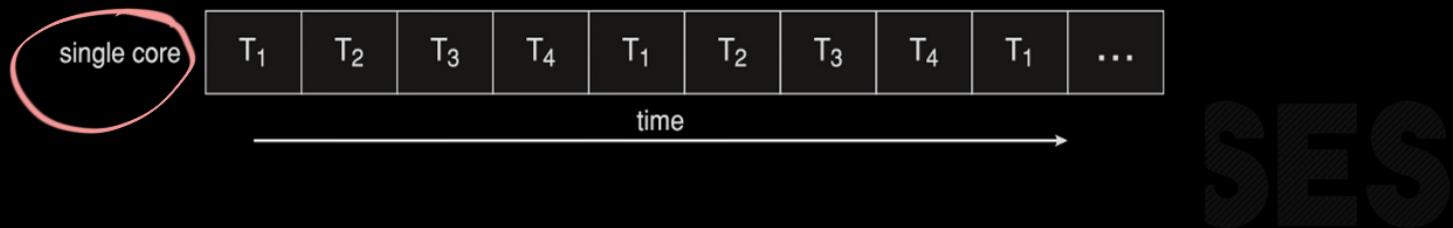
- Parallelism on a multi-core system

if we have two  
processors then  
we can run  
at most 2 threads  
at a time

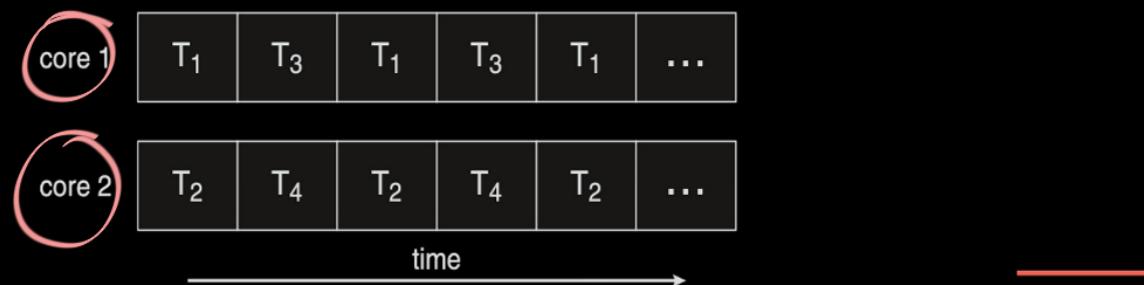
# Thread: Concurrency vs. Parallelism

- Threads are about **concurrency** and **parallelism**

- Concurrent execution on single-core system:**



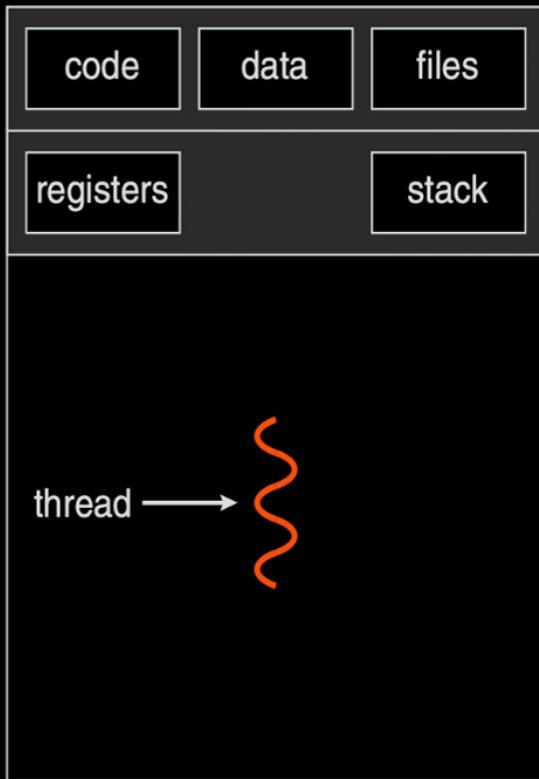
- Parallelism on a multi-core system:**



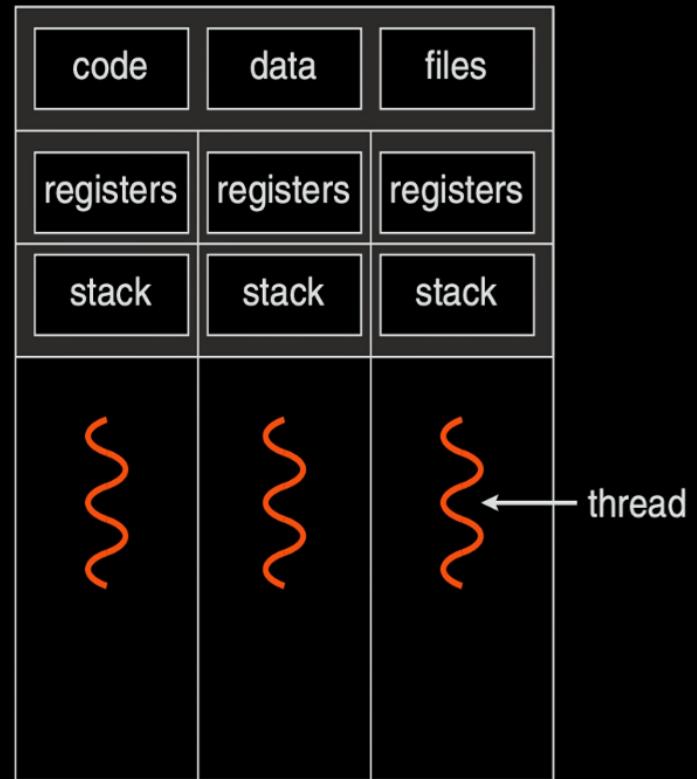


# Operating Systems

## Single and Multi threaded process

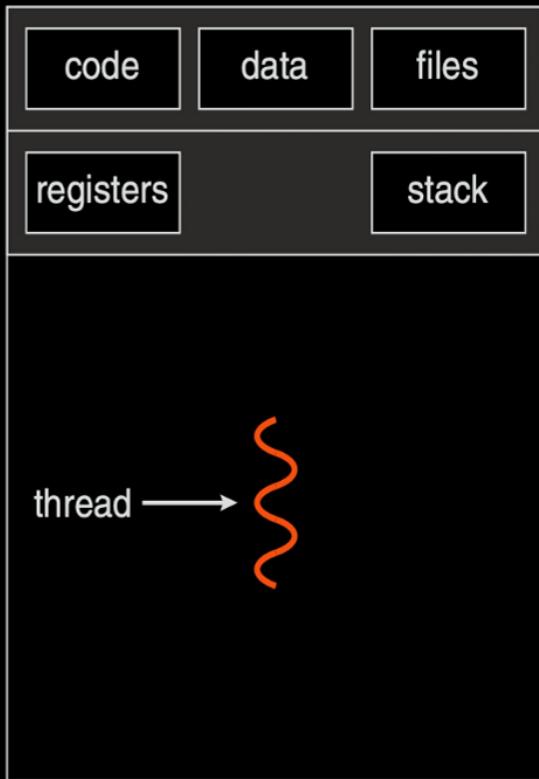


single-threaded process

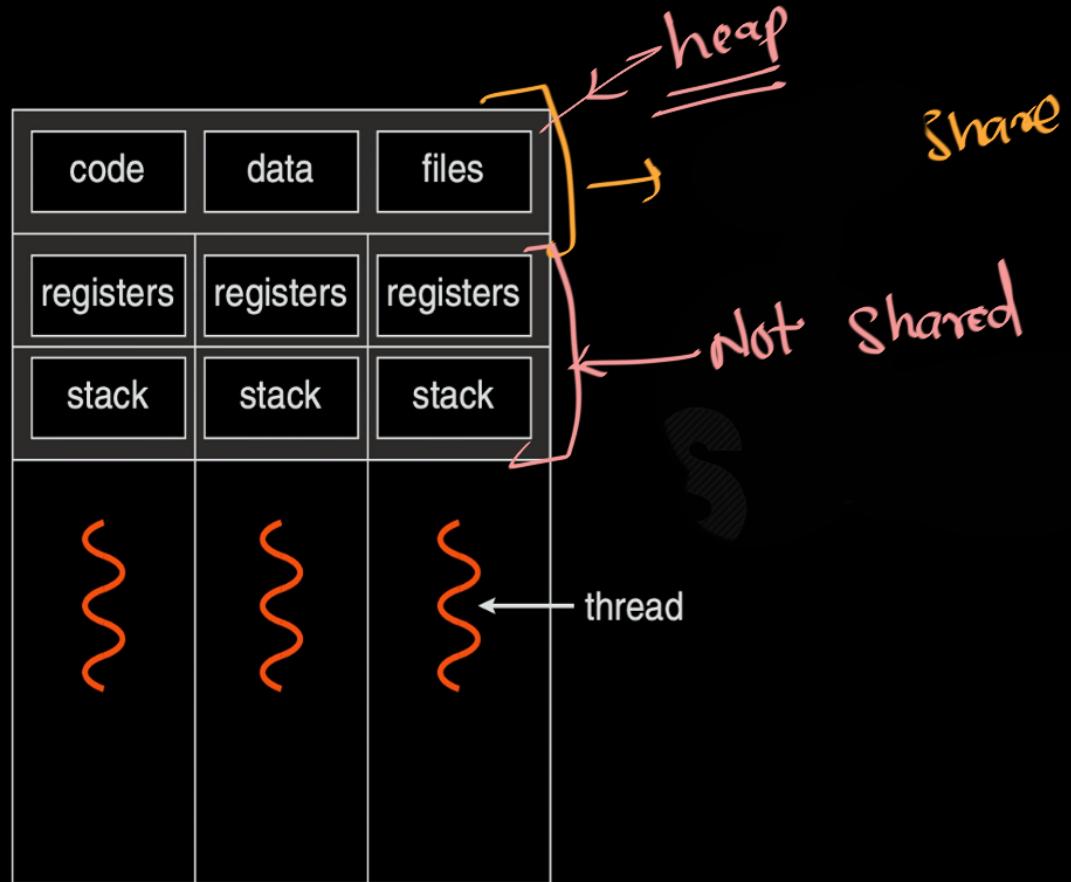


multithreaded process

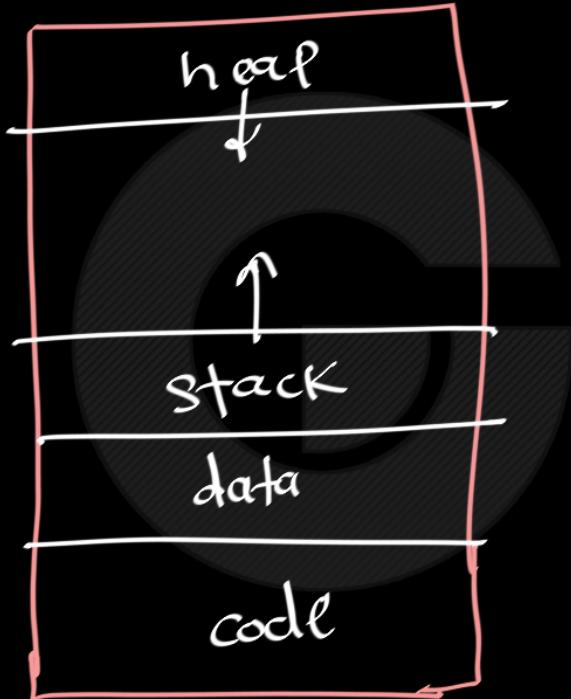
## Single and Multi threaded process



single-threaded process



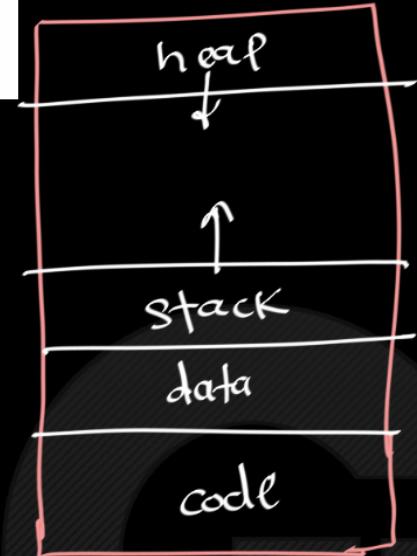
multithreaded process



Process

Content

→ might have opened  
some file  
→ some register values  
PC , R<sub>1</sub> , R<sub>2</sub>



Content

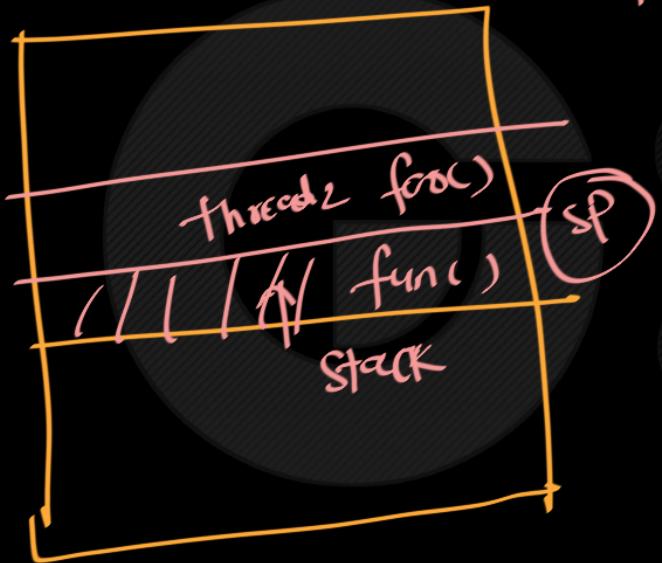
→ might have opened  
some file

→ some register values  
PC, R<sub>1</sub>, R<sub>2</sub>

Should Share

Code  
data

- ⇒ what are the things that threads  
should NOT share.  
→ Stack,
- Register values  
(PC, SP, R<sub>1</sub>PR)



thread1 int a = 5;

for ( )  
{  
    a = a + 1  
}

thread2 → {  
    func ( )  
    \$      a = 2 \* 9  
    \$



Distinguishing between these kinds of memory matters because each thread will have its own stack.

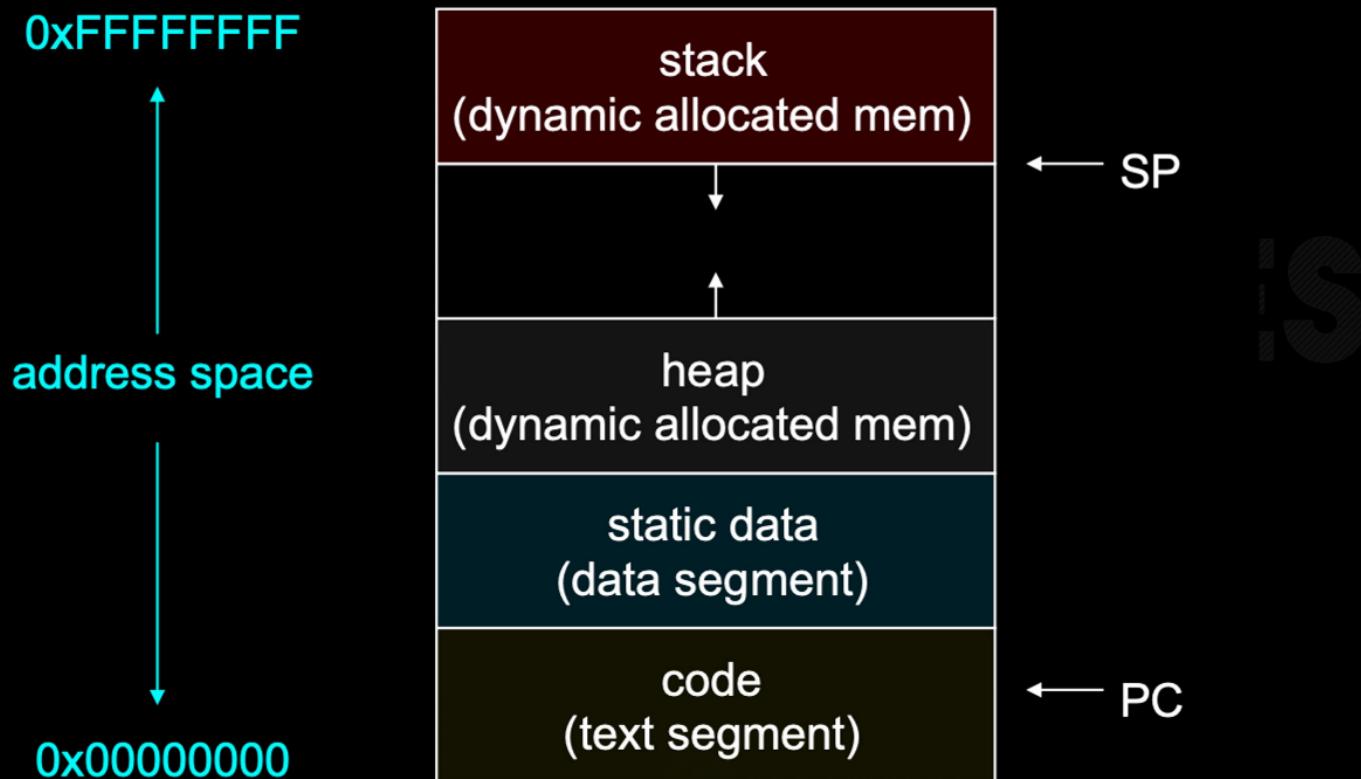
However, all the threads in a process will share the heap. Some people call threads lightweight processes because they have their own stack but can access shared data.

23-Nov-2022

CLASSES

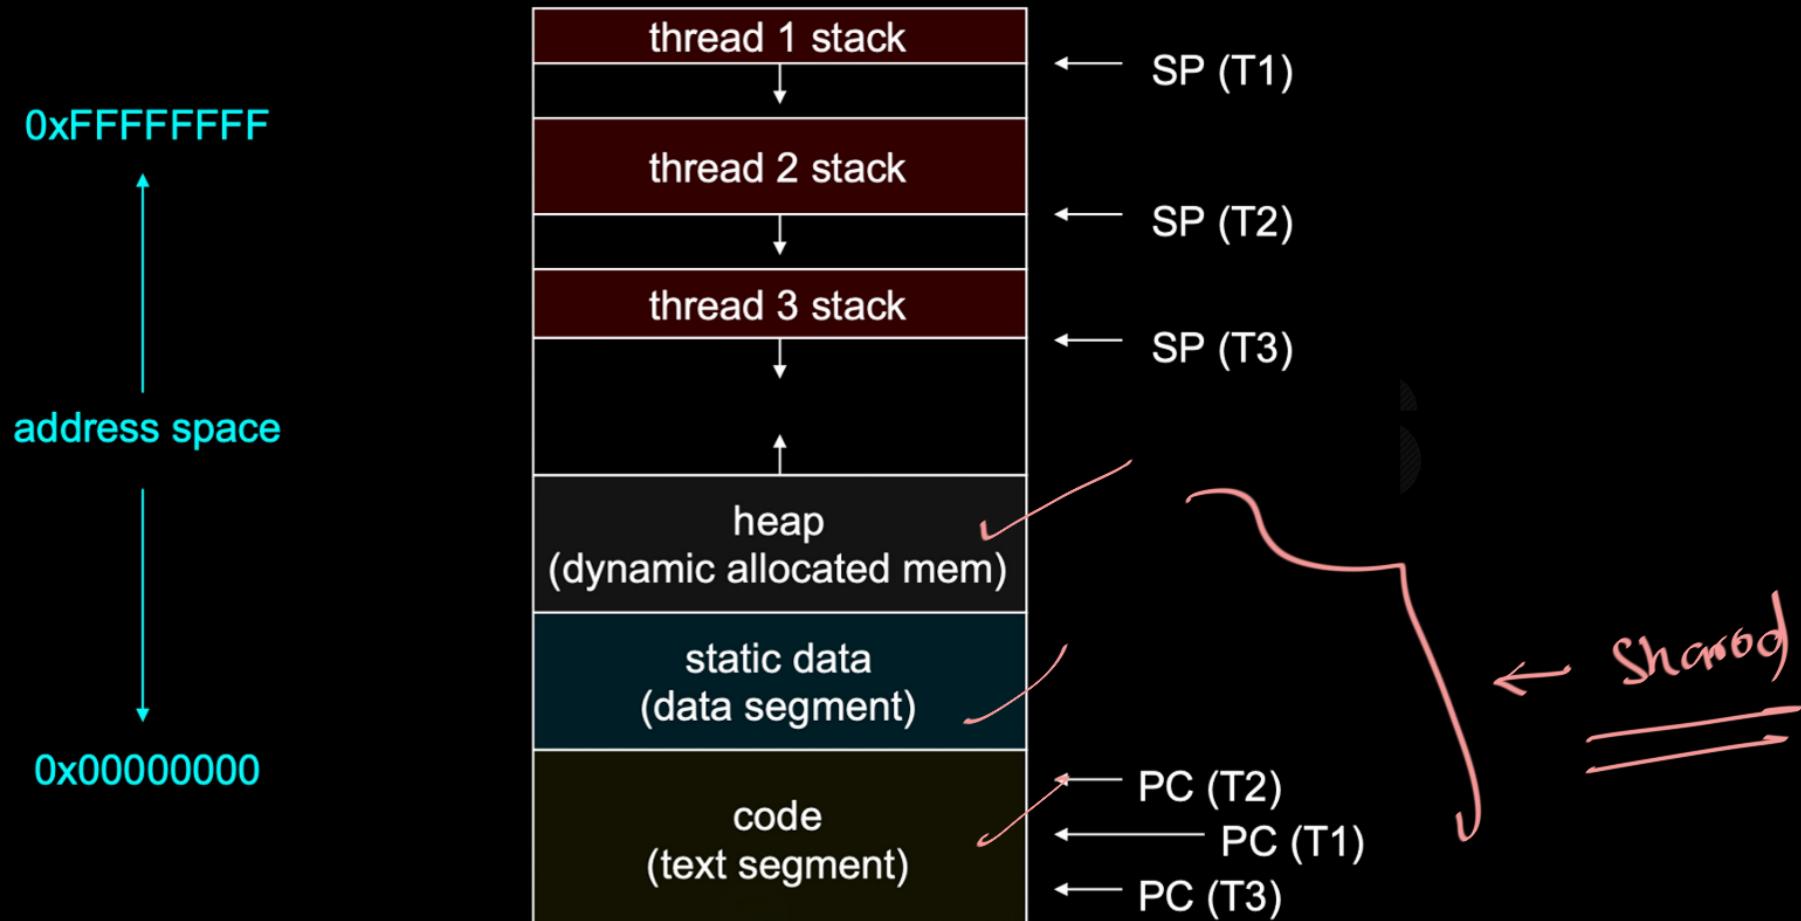


## (old) Process address space





## (new) Process address space with threads





## Types of threads

- User level
- Kernel level





## Types of threads

- User level
  - When we make threads using function calls
- Kernel level
  - When we make threads using System calls





# Types of threads

- User level
  - When we make threads using function calls
    - we have TCB
- Kernel level
  - When we make threads using System calls
  - For each thread there is new PCB (TCB)
  - Since only Kernel can create PCB hence System call is needed



## How to initiate threads ?

- User level
  - Some library help us
- Kernel level
  - Some library help us





# How to initiate threads ?

## 4.4 Thread Libraries

A **thread library** provides the programmer with an API for creating and managing threads. There are two primary ways of implementing a thread library. The first approach is to provide a library entirely in user space with no kernel support. All code and data structures for the library exist in user space. This means that invoking a function in the library results in a local function call in user space and not a system call.

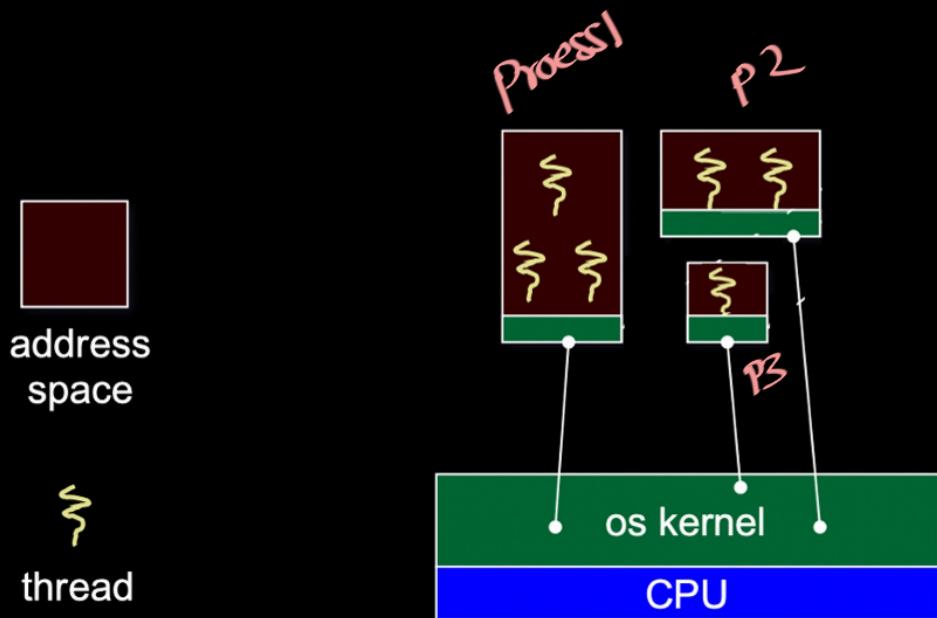
The second approach is to implement a kernel-level library supported directly by the operating system. In this case, code and data structures for the library exist in kernel space. Invoking a function in the API for the library typically results in a system call to the kernel.

(Optional read from Galvin)

ES  
user level

Kernel level

## User-level threads

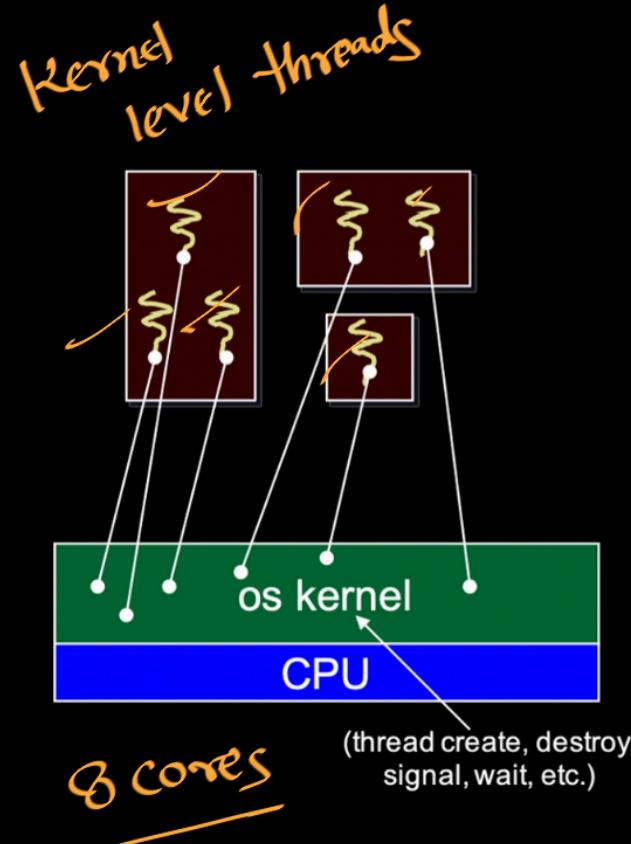
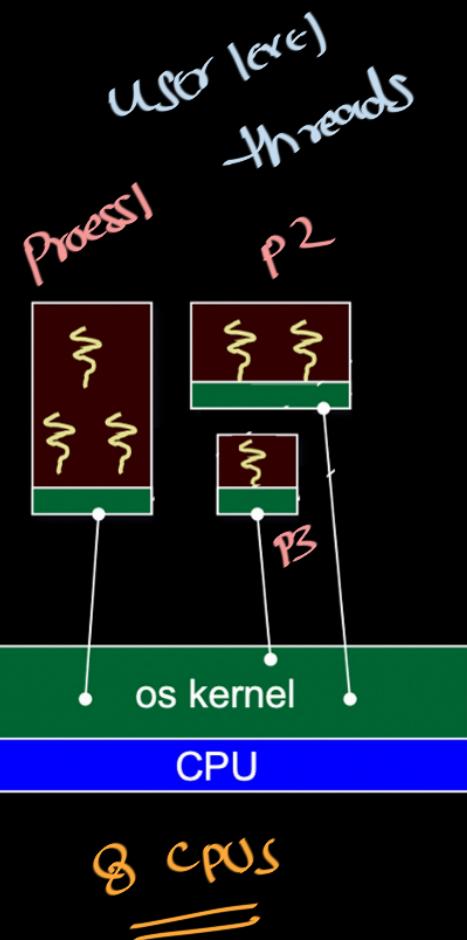


Now thread id is unique within the context of a process, not unique system-wide

## User-level threads

No free lunch

The diagram illustrates user-level threads. At the bottom is a blue 'CPU' block, above which is a green 'os kernel' block. Two processes, P1 and P2, are shown above the kernel. Process P1 contains two threads, each represented by a brown rectangle with a wavy line inside. Process P2 contains three threads, with the top one being a brown rectangle with a wavy line and the bottom two being green rectangles with wavy lines. Lines connect the threads in each process to the CPU and os kernel. A legend indicates that a brown square represents 'address space' and a wavy line represents a 'thread'.



Now thread id is unique within the context of a process, not unique system-wide



## User-level threads

- User-level threads are small and fast
  - managed entirely by user-level library
    - E.g., `pthreads` (`libpthreads.a`)
  - each thread is represented simply by a PC, registers, a stack, and a small **thread control block** (TCB)
  - creating a thread, switching between threads, and synchronizing threads are done via procedure calls
    - no kernel involvement is necessary!
- User-level thread operations can be 10-100x faster than kernel threads as a result



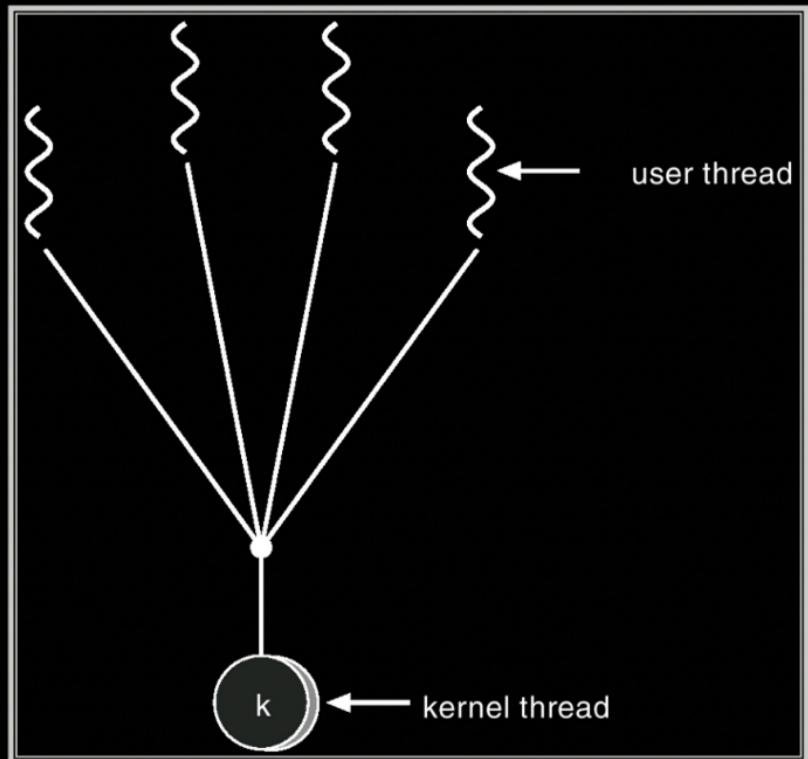


## User-level threads



# THREADS

- Many user-level threads mapped to single kernel thread.
- Used on systems that do not support kernel threads.
- Examples:  
Solaris Green Threads  
GNU Portable Threads



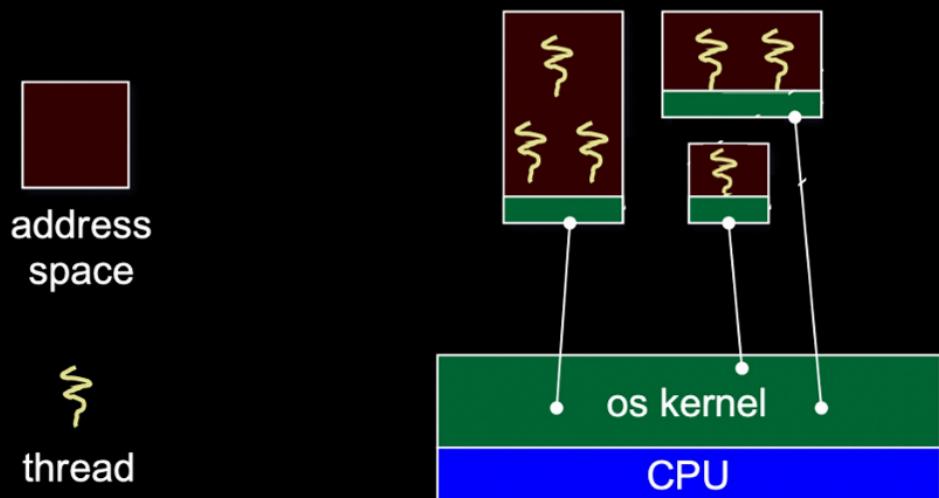
4: Threads



## User-Level Threads

- A **user-level thread** is a thread that the OS does *not* know about.
- The OS only knows about the process containing the threads.
- The OS only schedules the process, not the threads within the process.
- The programmer uses a *thread library* to manage threads (create and delete them, synchronize them, and schedule them).

## User-level threads

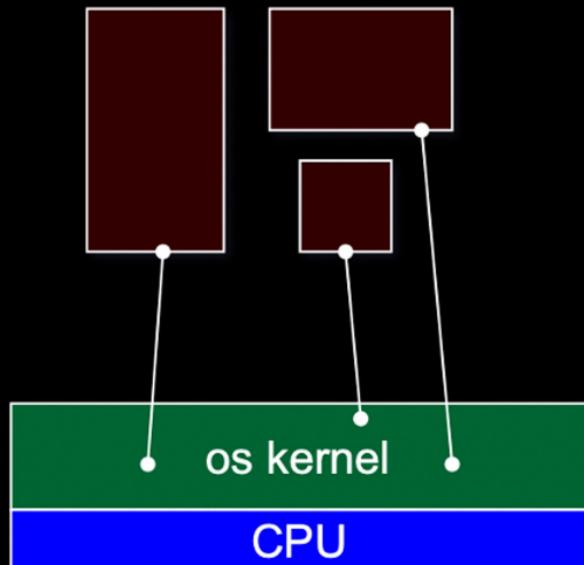


Now thread id is unique within the context of a process, not unique system-wide

## User-level threads: What the kernel sees

  
address  
space

  
thread



ES



## Question:

What is one major advantage and major disadvantage of User level threads ?





## Question:

What is one major advantage and major disadvantage of User level threads ?

Advantage: They are fast since no OS involvement (no context switch and all)

Disadvantage: OS doesn't know about threads so if one threads block because of I/O then all threads get block



## Question:

3. What is the biggest advantage of implementing threads in user space? What is the biggest disadvantage?

*The biggest advantage is efficiency. No traps to the kernel are needed to switch threads. The ability of having their own scheduler can also be an important advantage for certain applications. The biggest disadvantage is that if one thread blocks, the entire process blocks.*



## User-Level Threads: Disadvantages

- Since the OS does not know about the existence of the user-level threads, it may make poor scheduling decisions:
  - It might run a process that only has idle threads.
  - If a user-level thread is waiting for I/O, the entire process will wait.
  - Solving this problem requires communication between the kernel and the user-level thread manager.
- Since the OS just knows about the process, it schedules the process the same way as other processes, regardless of the number of user threads.
- For kernel threads, the more threads a process creates, the more time slices the OS will dedicate to it.

## Kernel threads



address  
space



thread

