

file system

implementation

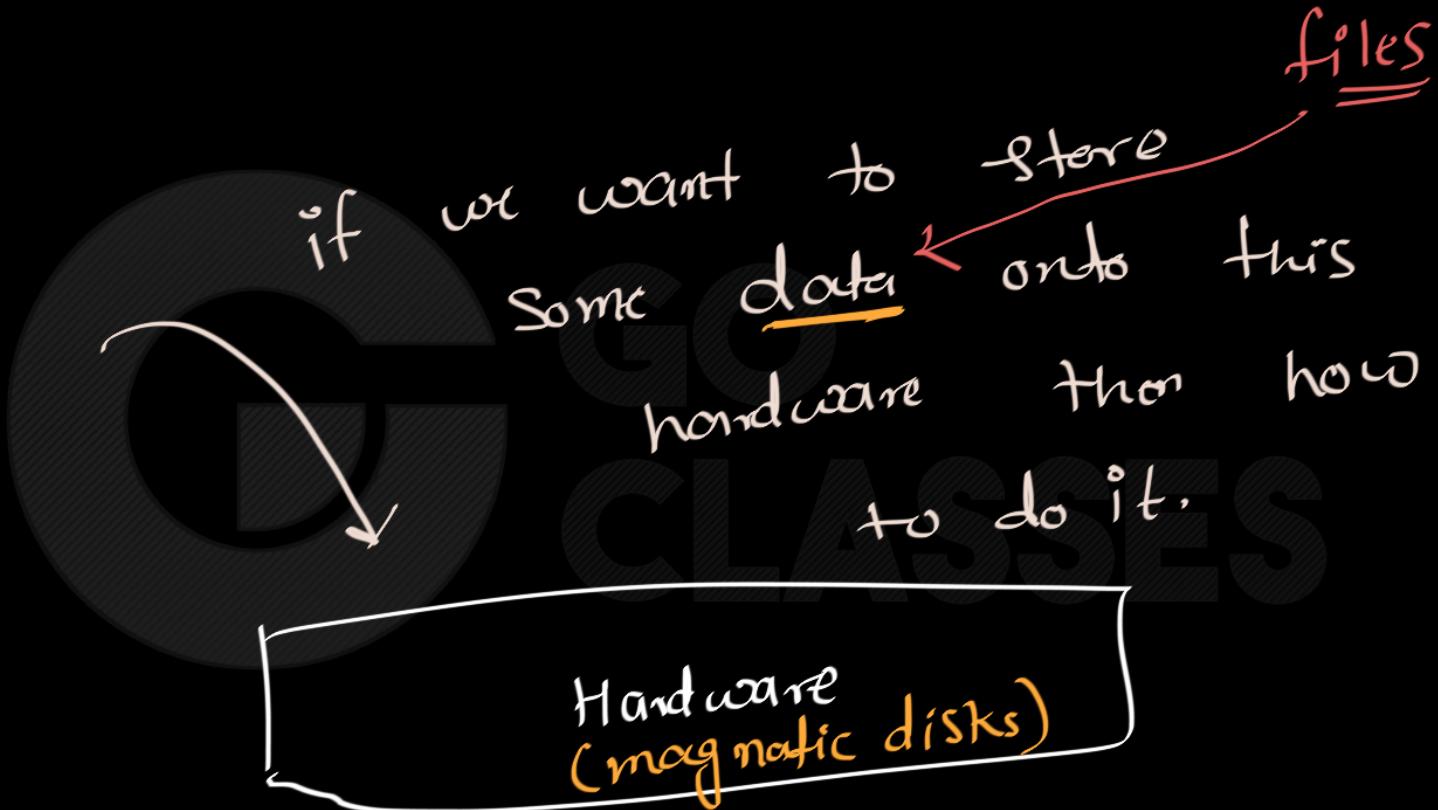
C

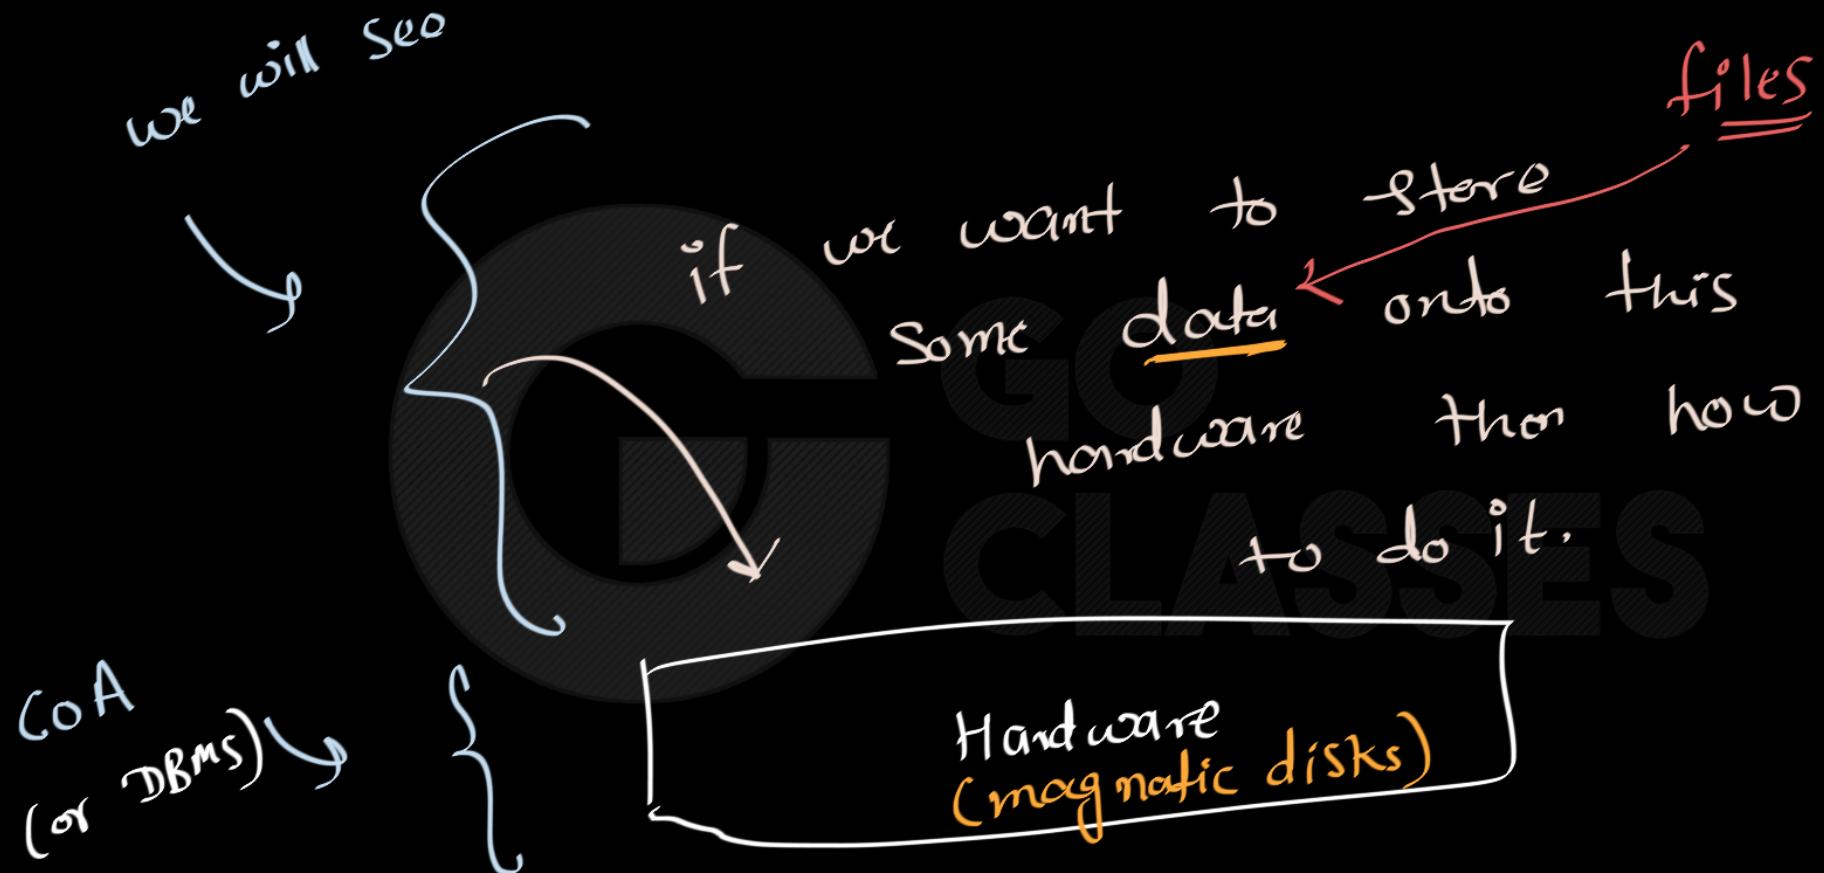
Hardware

how does secondary memory  
looks like

↓  
magnetic  
tapes

implementation







File System *(implementation)* CLASSES

# What is File ?

Data in some format



GO  
CLASSES  
• pdf  
• mp3  
• mp4

# What is File System ?

- Provides a way to store files
  - Provides a way to name files
  - Provides a way to store metadata of file
- ....

How to manage the  
files on hardware.

# File types – name, extension

Extension	Meaning
.bak	Backup file
.c	C source program
.gif	Compuserve Graphical Interchange Format image
.hlp	Help file
.html	World Wide Web HyperText Markup Language document
.jpg	Still picture encoded with the JPEG standard
.mp3	Music encoded in MPEG layer 3 audio format
.mpg	Movie encoded with the MPEG standard
.o	Object file (compiler output, not yet linked)
.pdf	Portable Document Format file
.ps	PostScript file
.tex	Input for the TEX formatting program
.txt	General text file
.zip	Compressed archive

Figure 4-1. Some typical file extensions.

Source: Tanenbaum

# File Operations

- Create
  - Rename
  - Open
  - Read
  - Write
- ..etc



Store files on disk

what should be key features for file system

→ easy to fetch (efficient fetching)

you should not take too much time to just calculate where is your file.

→ file should not get mixed up.

→

## Plan

{ we will see few common ways to implement file system.

- method 1 : this is how you should store contiguous / extent based files
- method 2 : "
- method 3 : "

How modern file systems are implemented

- windows
- Mac
- linux
- unix

(90% method 2 + 10% method 3)



Major component of any file system;

"How to store files"

Sector or Block = unit of atomicity

— —  
modify one byte in a file

- Disk reads/writes in terms of sectors, not bytes

- Read/write single sector



ES

Size of block : 4 KB

- How to write a single byte? “Read-modify-write”

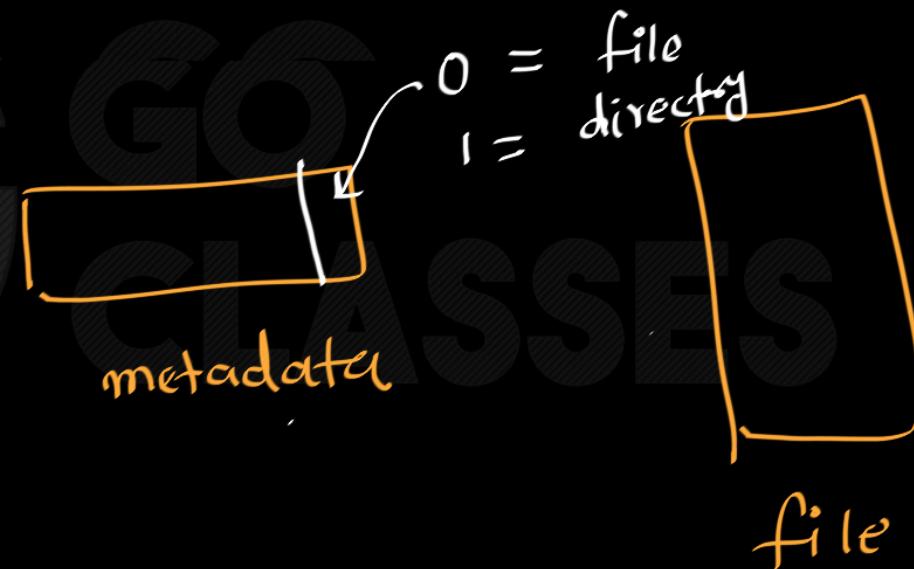
- Read in sector containing the byte
- Modify that byte
- Write entire sector back to disk



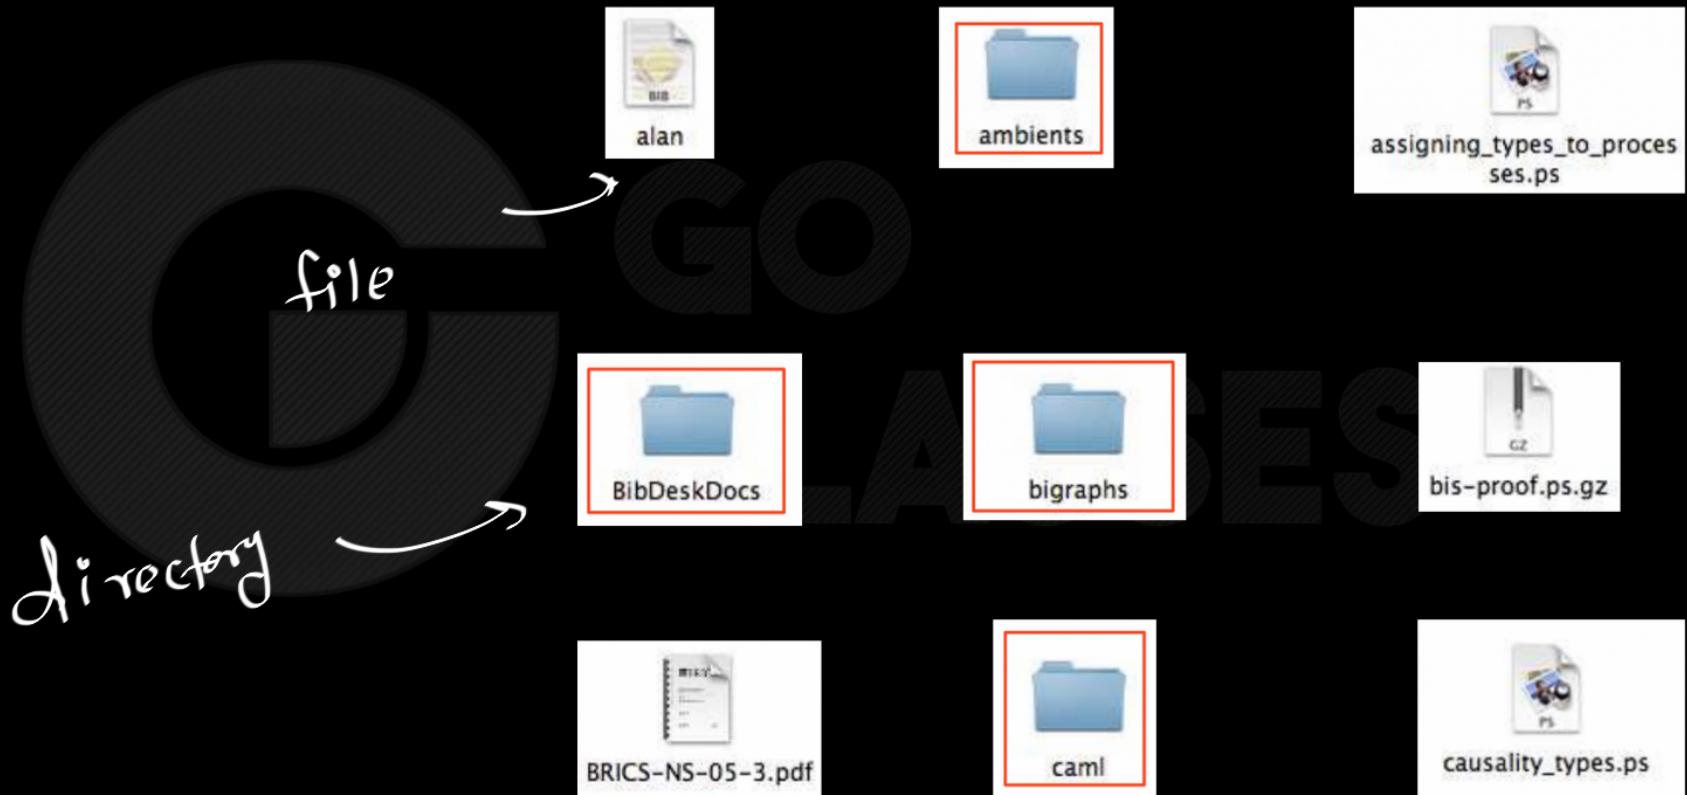
# Directory

---

A File (in UNIX, **everything is file**) with special bit set in metadata to indicate directory.



# Directories



# File Metadata

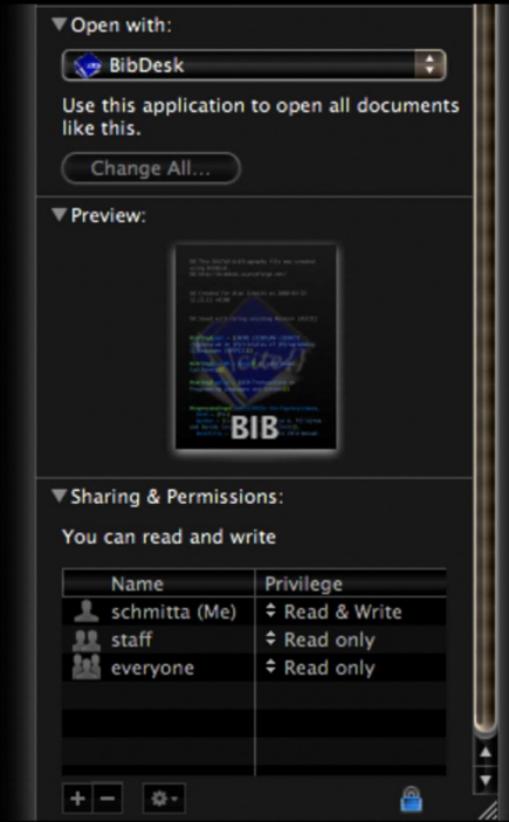
---

- **Name** - only information kept in human-readable form
  - **Identifier** - unique tag (number) identifies file within file system (**inode number** in UNIX)
  - **Location** - pointer to file location on device
  - **Size** - current file size
  - **Protection** - controls who can do reading, writing, executing
  - **Time, date, and user identification** - data for protection, security, and usage monitoring
  
  - How is metadata stored? (**inode** in UNIX)
- 

ES

Stored in  
some data structure  
= Inside

# File Metadata



# Inodes

## Inodes

- Store the metadata for a file (which data blocks belong to the file, file size, owner, access rights, ...)
  - Inode stands for index node
- 



- **Inodes** - An inode is the data structure that describes the meta data of files (regular files and directories). One inode represents one file or one directory. An inode is composed of several fields such as ownership, size, modification time, mode (permissions), reference count, and data block pointers (which point to the data blocks that store the content of the file). Note that the inode does not include a file/directory name. A file's inode number can be found using the `ls -i` command.

- for every file or directory  $\Leftrightarrow$  One inode
- Every inode has one number      inode number

File(or Directory)  $\equiv$  inode (header) + data



metadata



What does inode of file contain ?

What does data of file contain ?



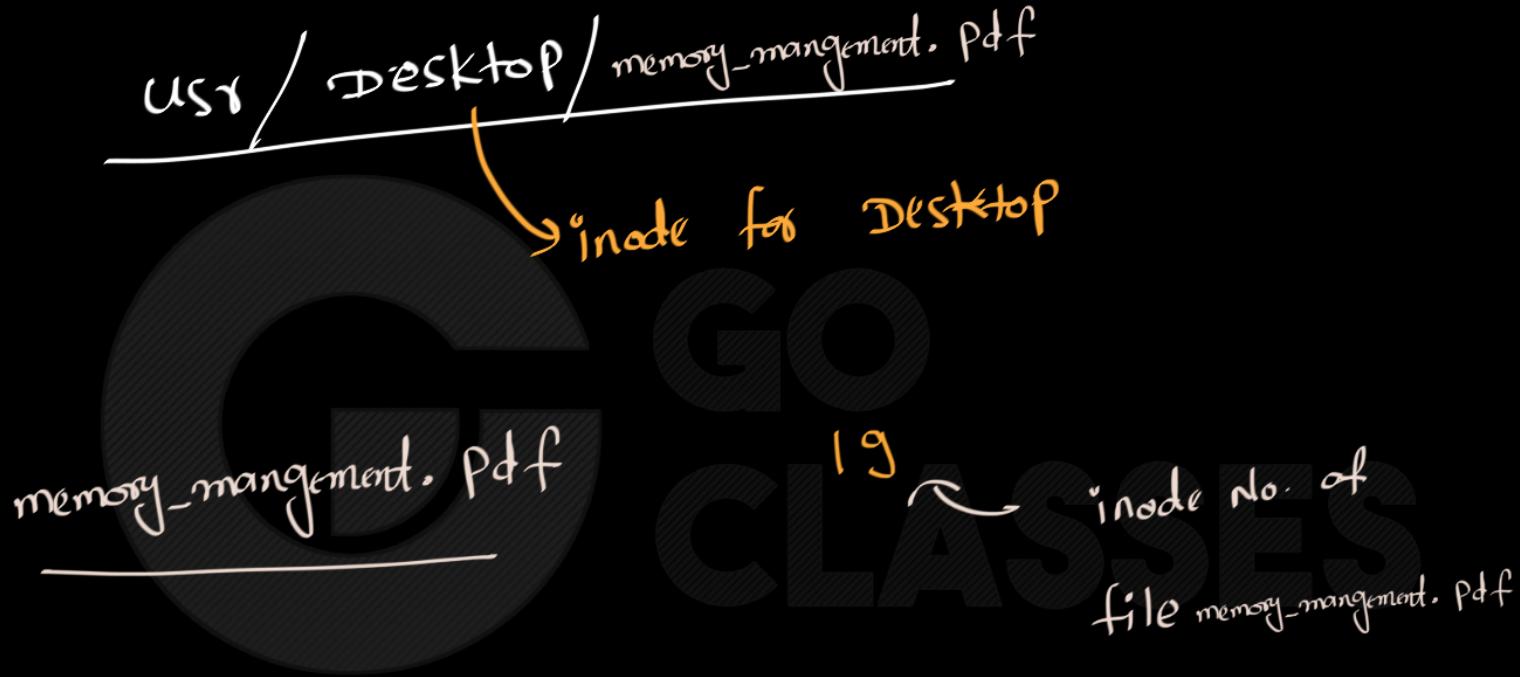
What does inode of file contain ?

→ metadata

What does data of file contain ?

→ data

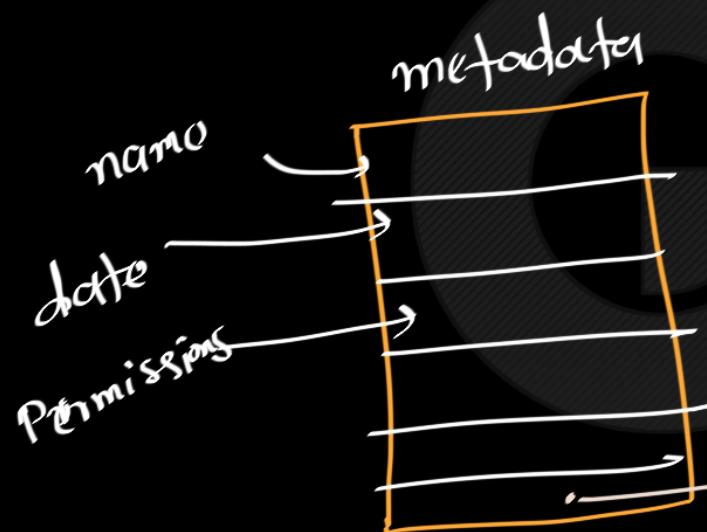
GO  
CLASSES



How to get that inode number?

---

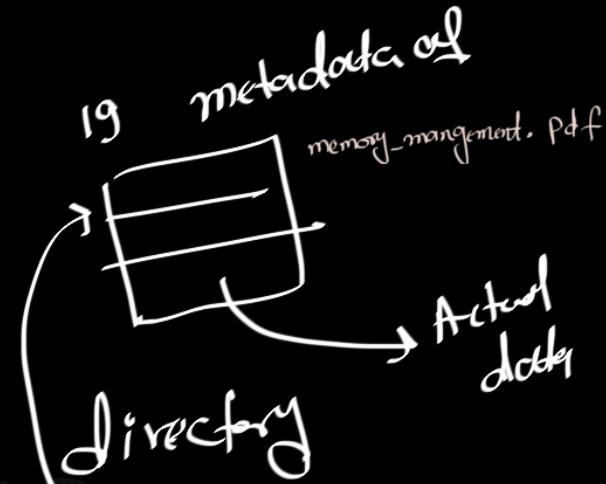
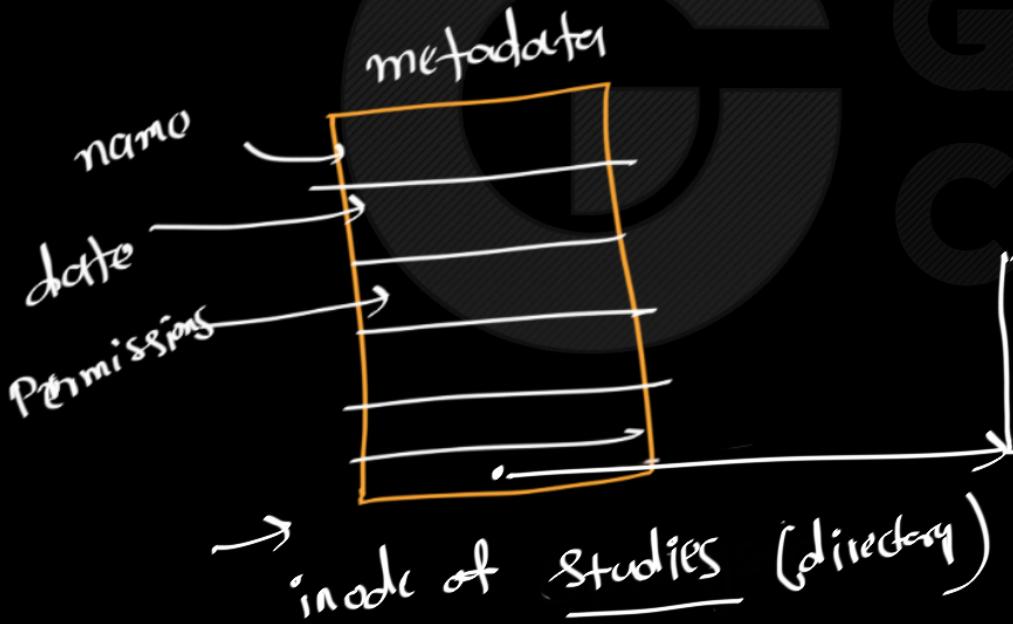
'inode' for a file



data block (actual data)

where is this file stored  
(data block pointer)

inode for a directory



19 metadata of

memory-management.pdf

Actual  
data

directory

19

inode table

# Inodes: How to index the content of a file?

19

## Indexing inodes

- An inode is identified by an **inumber**
- Corresponds to its index in the inode table
- Computing in which sector an inode is stored is easy (inputs: inode table start address, inumber, size of inode, size of block, size of sector)

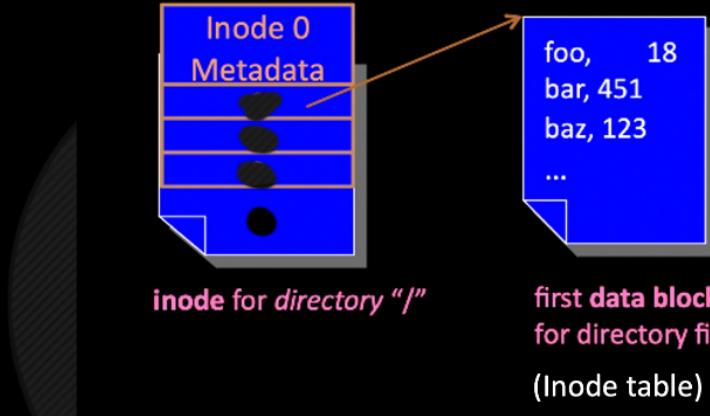
What does inode of directory contain ?

What does data of directory contain ?

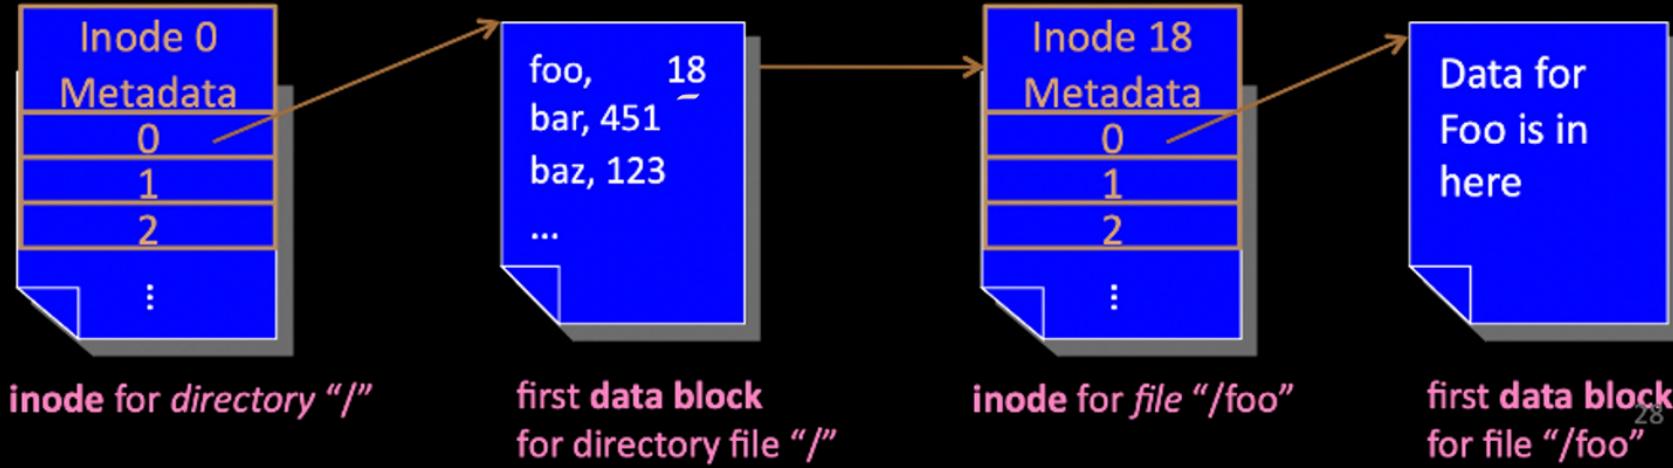
What does inode of directory contain ?

What does data of directory contain ?

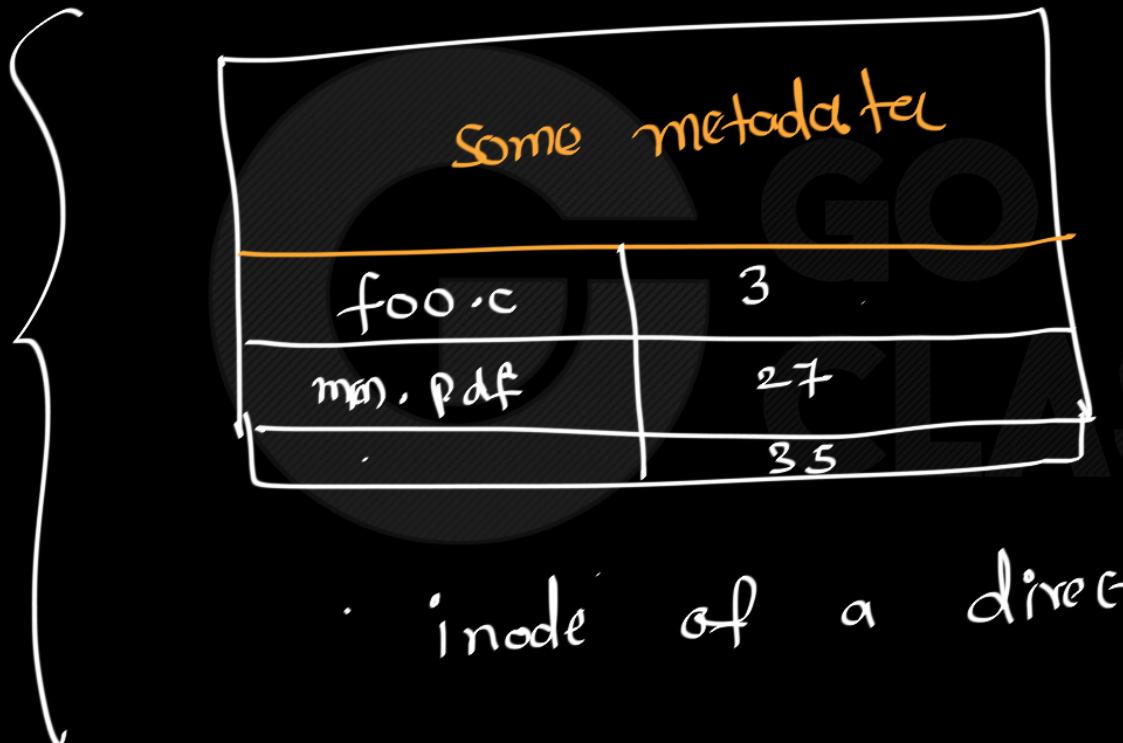
Inode table



CLASSES



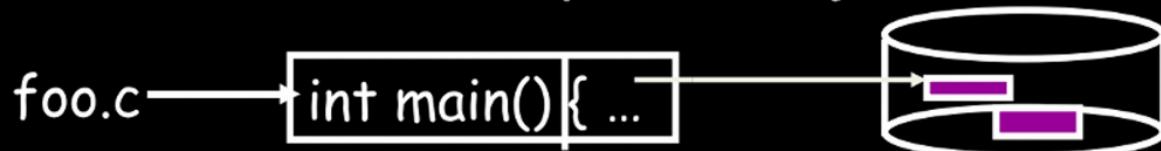
windows 95



is this a  
possible implementation  
Yes

- **File abstraction:**

- User's view: named sequence of bytes



- FS's view: collection of disk blocks
- File system's job: translate name & offset to disk blocks:



(foo.c, S<sub>12</sub>)

- Like page tables, file system metadata are simply data structures used to construct mappings

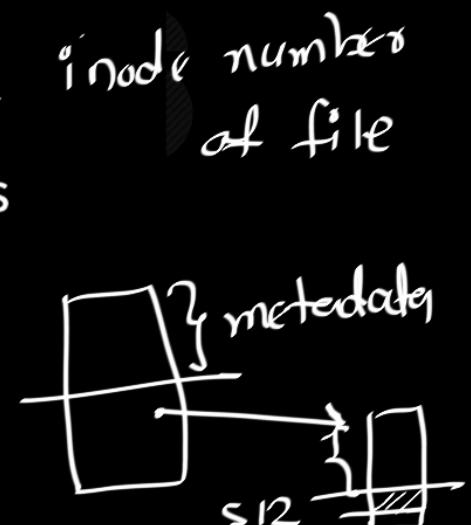
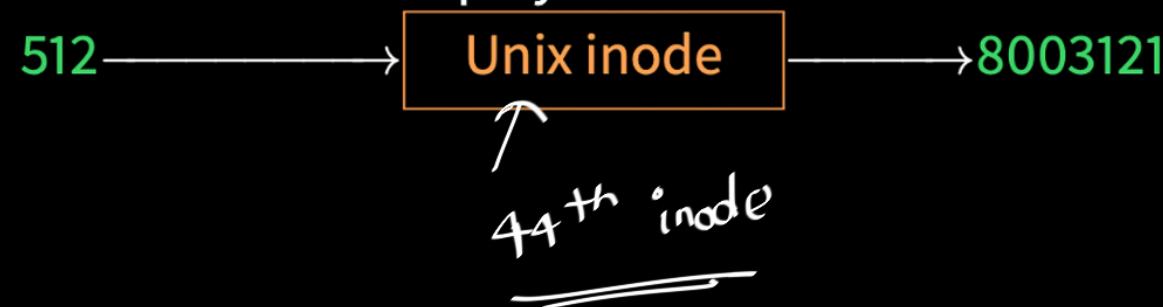
- Page table: map virtual page # to physical page #



- Directory: map name to disk address or file #



- File metadata: map byte offset to disk block address



# A short history of directories

- **Approach 1: Single directory for entire system**

- Put directory at known location on disk
- Directory contains  $\langle \text{name}, \text{inumber} \rangle$  pairs
- If one user uses a name, no one else can
- Many ancient personal computers work this way

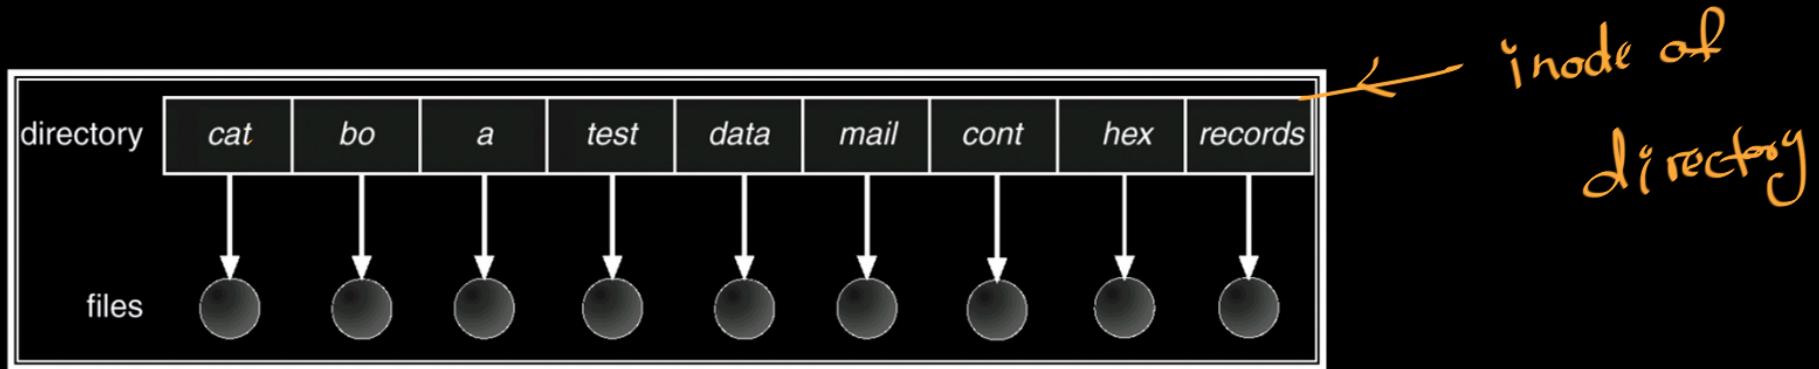
- **Approach 2: Single directory for each user**

- Still clumsy, and `ls` on 10,000 files is a real pain

- **Approach 3: Hierarchical name spaces**

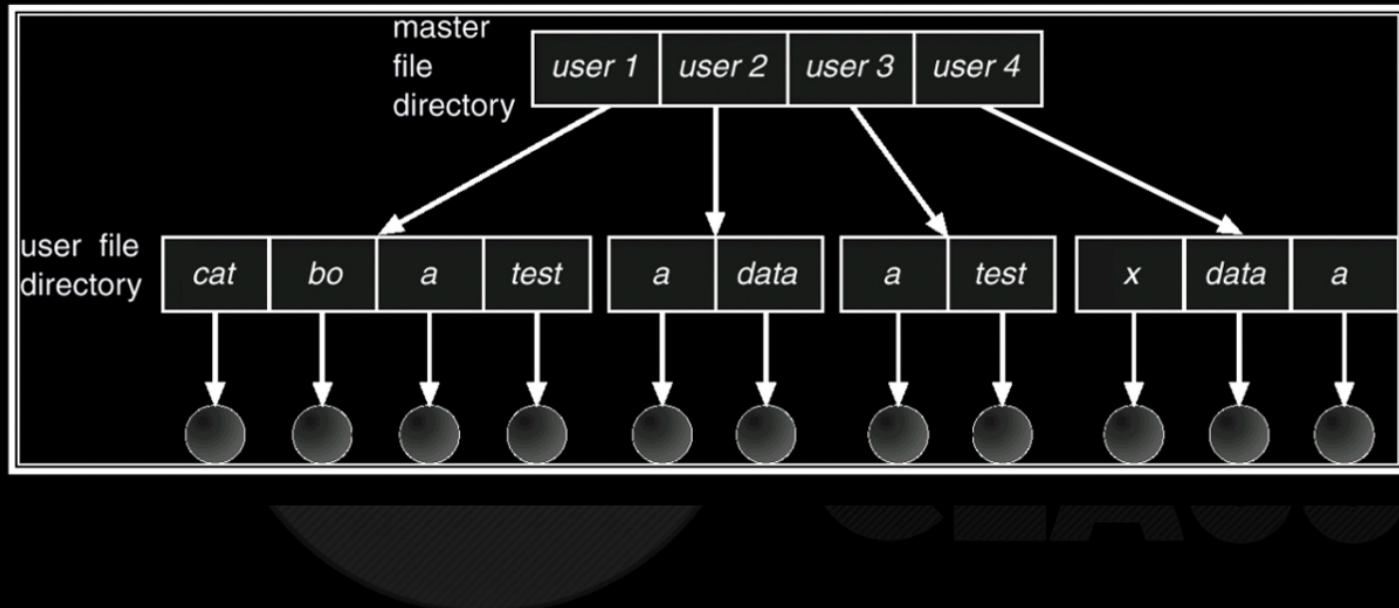
- Allow directory to map names to files *or other dirs*
- File system forms a tree (or graph, if links allowed)
- Large name spaces tend to be hierarchical (ip addresses, domain names, scoping in programming languages, etc.)

- Single-Level Directory

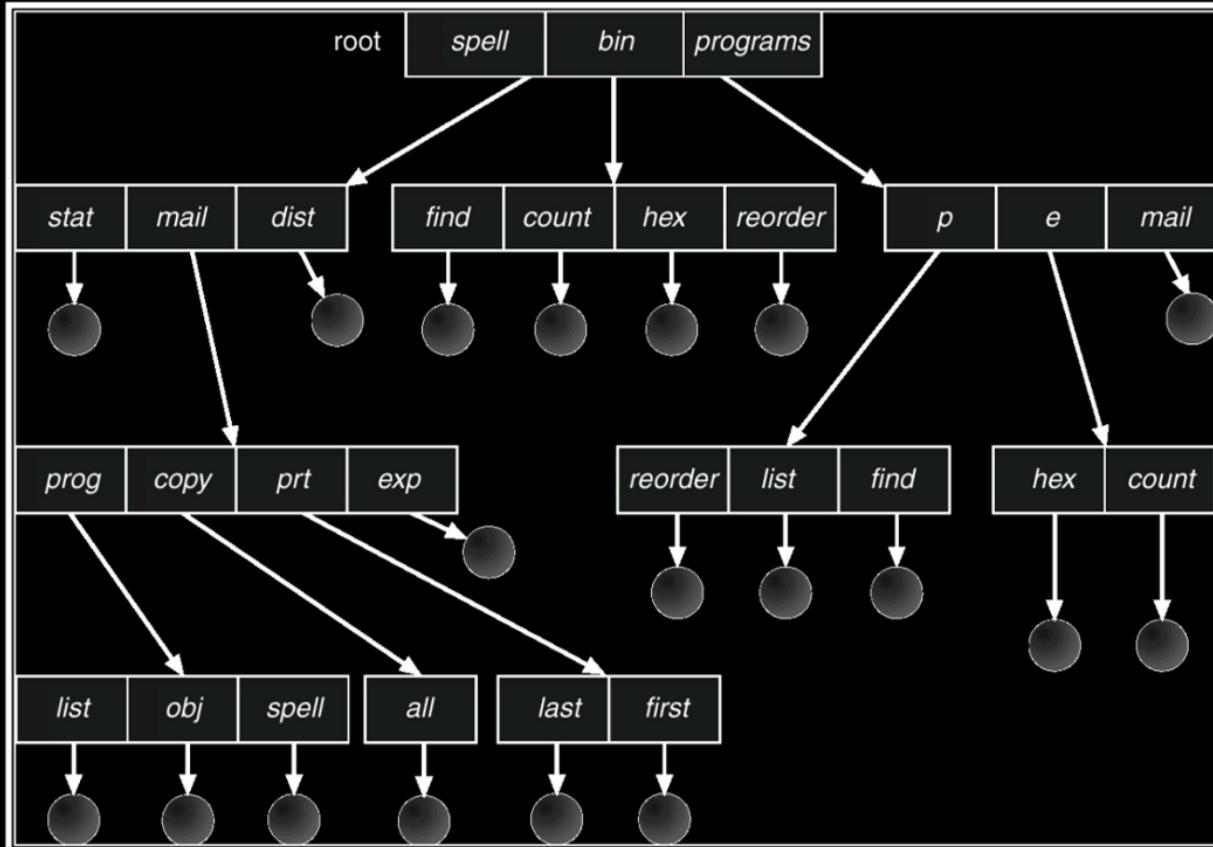


CLASSES

- Two-Level Directory

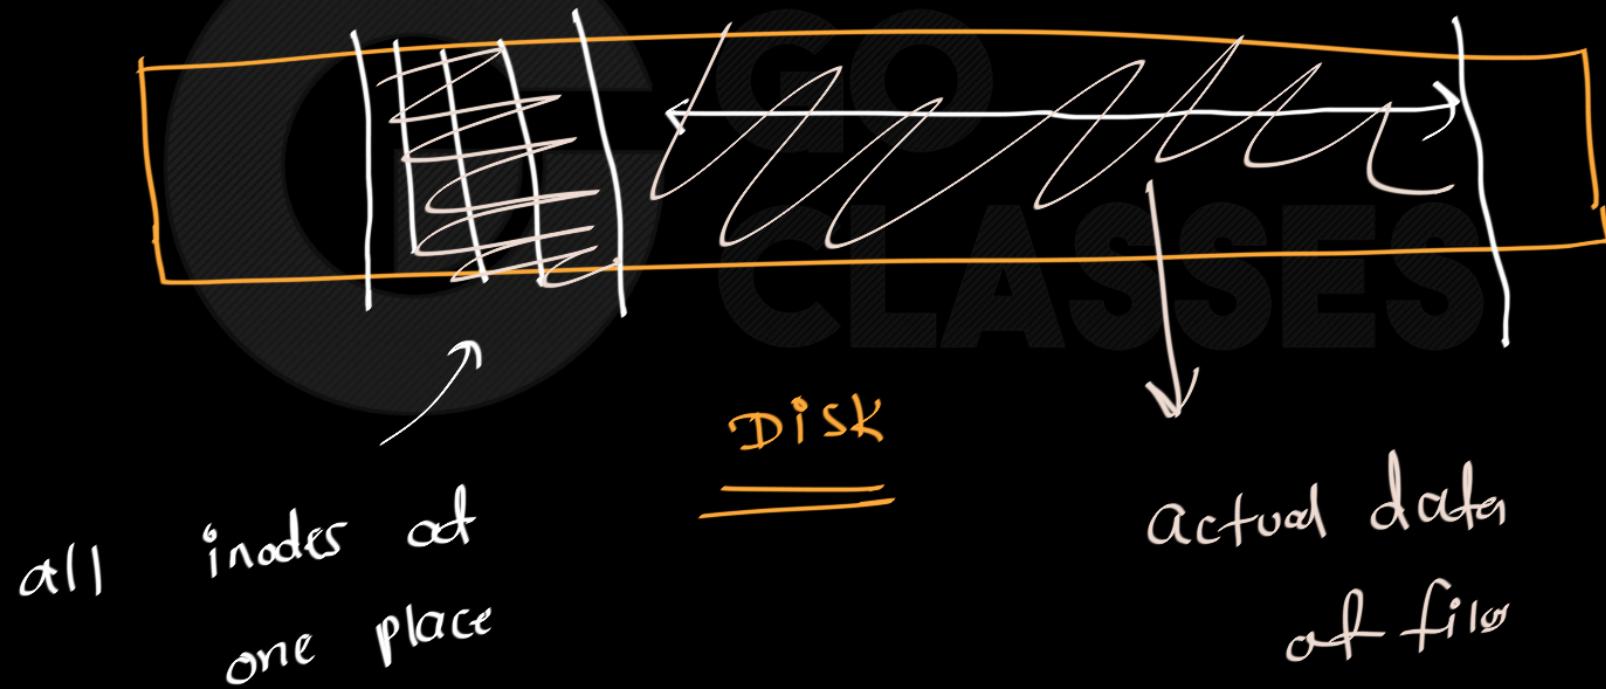


- Tree-Structured Directories



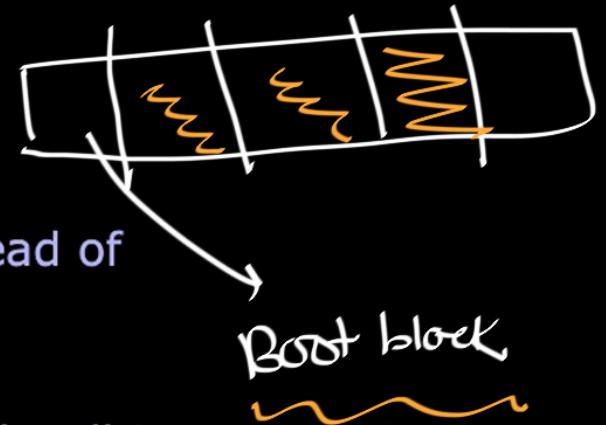
modern system

SSES



## (Old) Unix disks are divided into five parts ...

- Boot block
  - can boot the system by loading from this block
- Superblock
  - specifies boundaries of next 3 areas, and contains head of freelists of inodes and file blocks
- i-node area
  - contains descriptors (i-nodes) for each file on the disk; all i-nodes are the same size; head of freelist is in the superblock
- File contents area ← file
  - fixed-size blocks; head of freelist is in the superblock
- Swap area
  - holds processes that have been swapped out of memory





- Boot block: usually contains bootstrap code
- Super block: size, # files, free blocks, etc.
- Inode list
  - size fixed when configuring file system
  - contains all inodes
- Data blocks: file data (huge area)
- Files, directories organized into tree-like structure



"inode



contenus      metadate

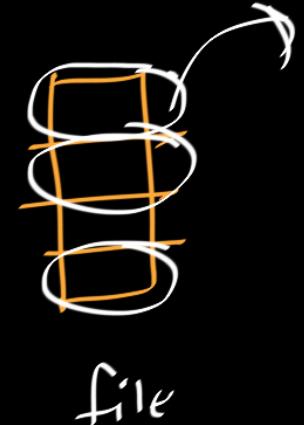
**GO  
CLASSES**

1. (16 points total) Short answer questions:

a. (2 points) What UNIX structure is used to keep track of the sectors allocated to a given file?

b. (2 points) What is the smallest addressable piece of data on a disk drive?

c. (2 points) What is a persistent, named collection of data?

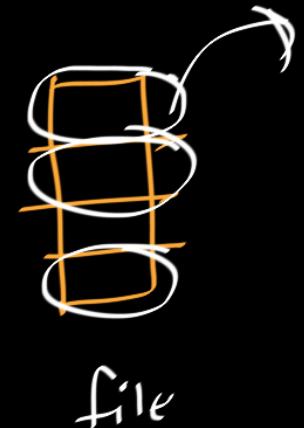


ES

1. (16 points total) Short answer questions:

- a. (2 points) What UNIX structure is used to keep track of the sectors allocated to a given file?

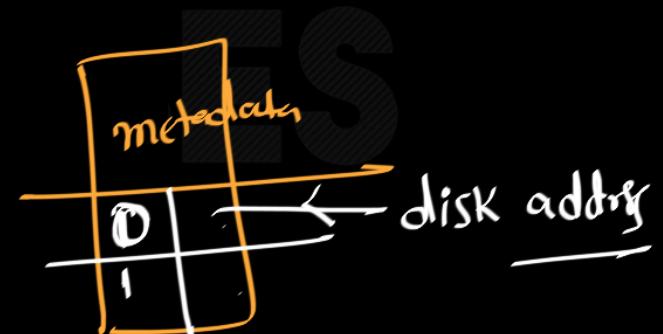
inode



- b. (2 points) What is the smallest addressable piece of data on a disk drive?

Sector

- c. (2 points) What is a persistent, named collection of data?



1. (16 points total) Short answer questions:

a. (2 points) What UNIX structure is used to keep track of the sectors allocated to a given file?

*i-node, file header, or index block*

---

b. (2 points) What is the smallest addressable piece of data on a disk drive?

*Disk sector or block*

---

c. (2 points) What is a persistent, named collection of data?

*File or directory*

---





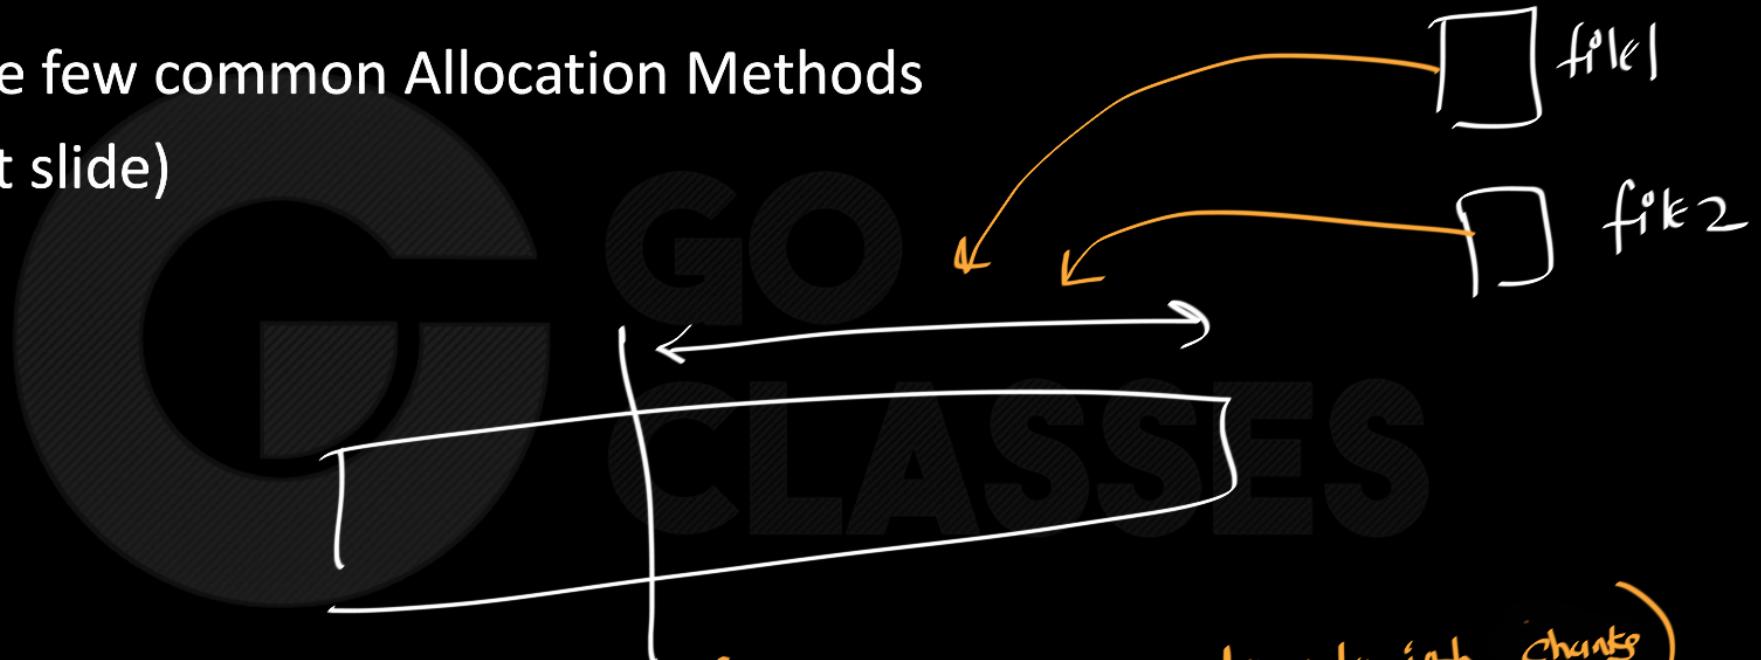
# File Implementation

How to allocate disk space to files ?

# How to allocate disk space to files ?

There are few common Allocation Methods

(see next slide)



- How to store file ( contiguous or divide into chunks )
- how inodes are organised.

# Allocation strategies

## □ Various approaches (similar to memory allocation)

- Contiguous
  - Extent-based (Variation to Contiguous)
- Linked
  - FAT tables (Variation to Linked)
- Indexed
  - Multi-Level Indexed (Variation to Indexed)

## □ Key metrics

- Fragmentation (internal & external)?
- Grow file over time after initial creation?
- Fast to find data for sequential and random access?
- Easy to implement?
- Storage overhead?

SES

# Allocation strategies

## □ Various approaches (similar to memory allocation)

### → Contiguous

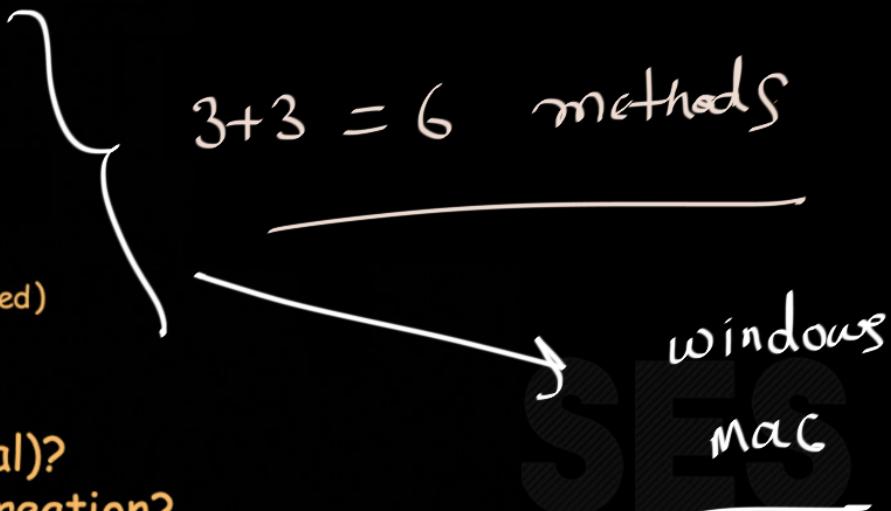
- Extent-based (Variation to Contiguous)

### → Linked

- FAT tables (Variation to Linked)

### → Indexed

- Multi-Level Indexed (Variation to Indexed)

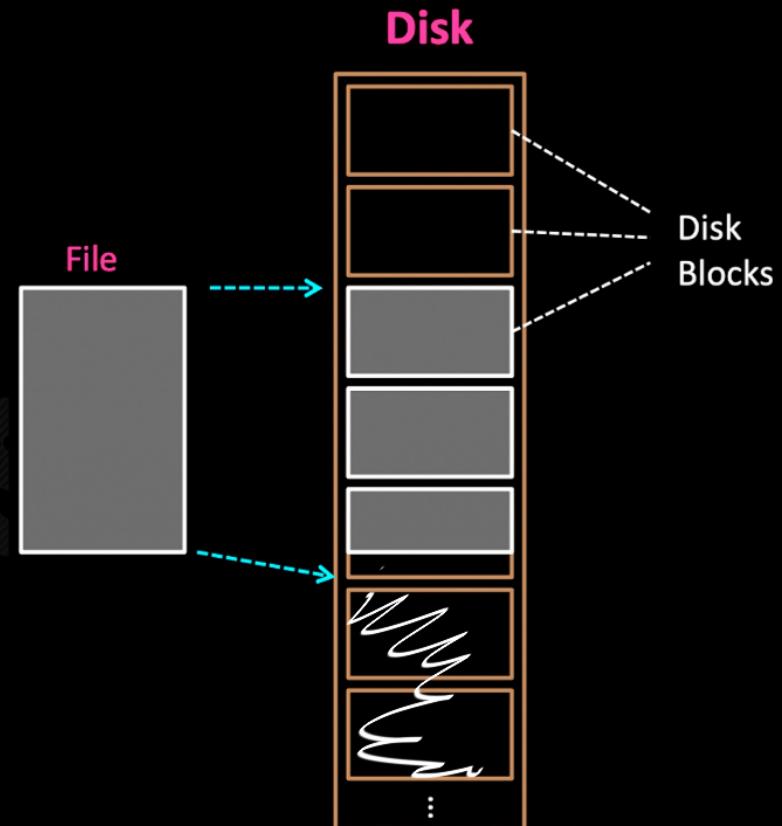
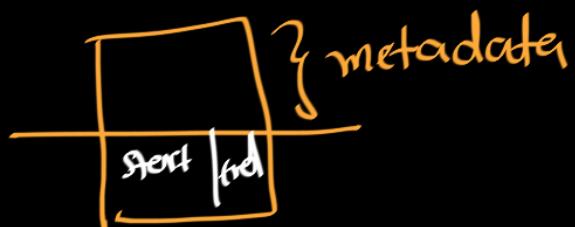


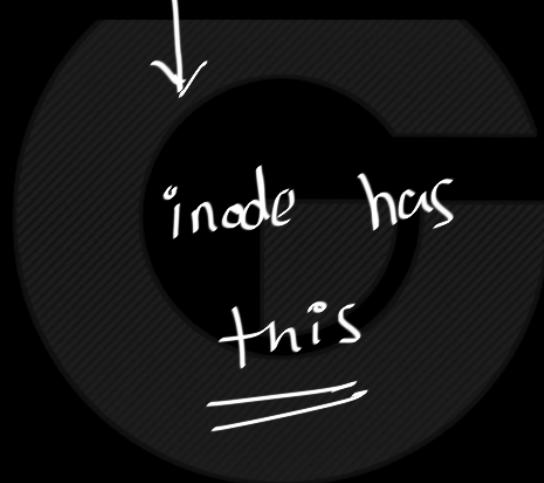
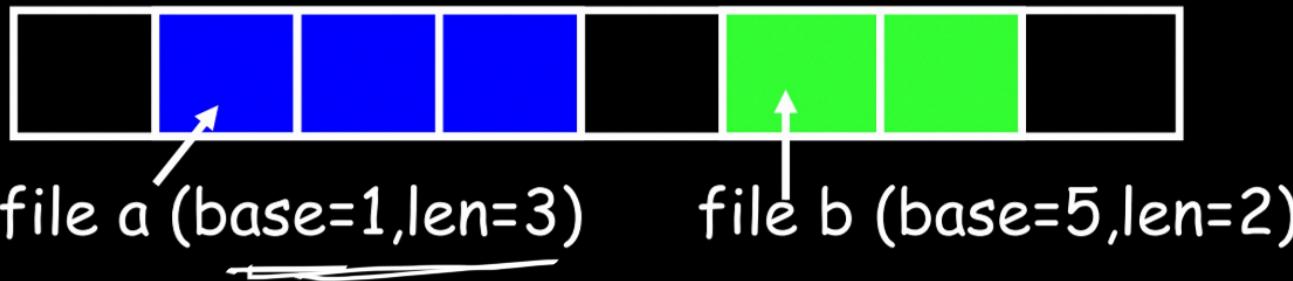
## □ Key metrics

- Fragmentation (internal & external)?
- Grow file over time after initial creation?
- Fast to find data for sequential and random access?
- Easy to implement?
- Storage overhead?

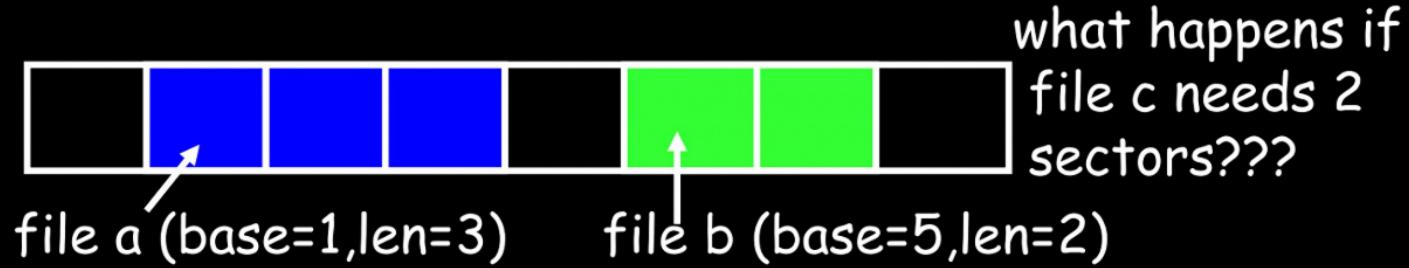
# Contiguous allocation

- User specifies length, file system allocates space all at once
- Metadata (inode):  
Contains starting location and size of file





GO  
CLASSES



GO  
CLASSES

# Pros and cons

## □ Pros

- Easy to implement
- Low storage overhead (two variables to specify disk area for file)
- Fast sequential access since data stored in continuous blocks
- Fast to compute data location for random addresses. Just an array index

Start location and  
length

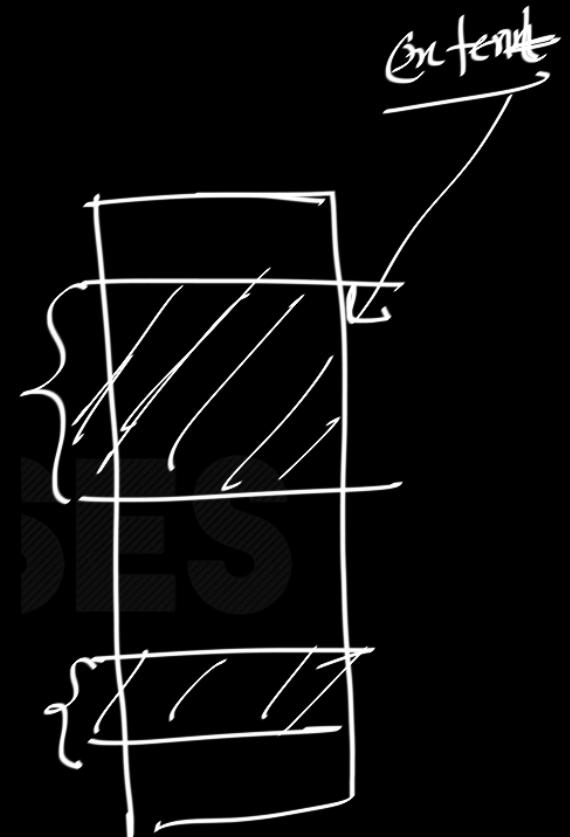
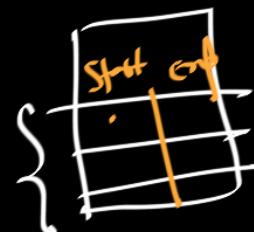
## □ Cons

- Small internal fragmentation
- Large external fragmentation
- Difficult to grow file

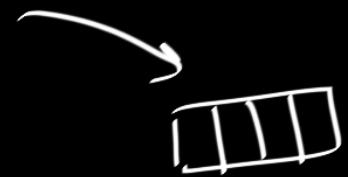
# Extent-based allocation

- Multiple contiguous regions per file (like segmentation)
  - Each region is an extent
  - Metadata: contains small array of entries designating extents
    - Each entry: start and size of extent

all extents  
info



# Pros and cons



## □ Pros

- Easy to implement
- Low storage overhead (a few entries to specify file blocks)
- ✓ ▪ File can grow overtime (until run out of extents)
- Fast sequential access
- Simple to calculate random addresses

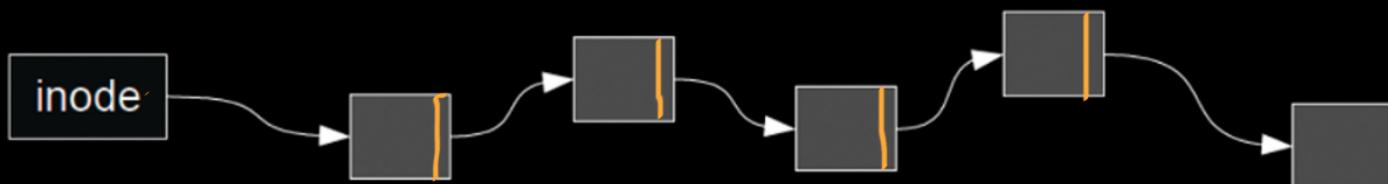
## □ Cons

- Help with external fragmentation, but still a problem

# Linked allocation

- All blocks of a file on linked list
  - Each block has a pointer to next
- Metadata (inode): pointer to the first block

8 bytes



# Pros and cons

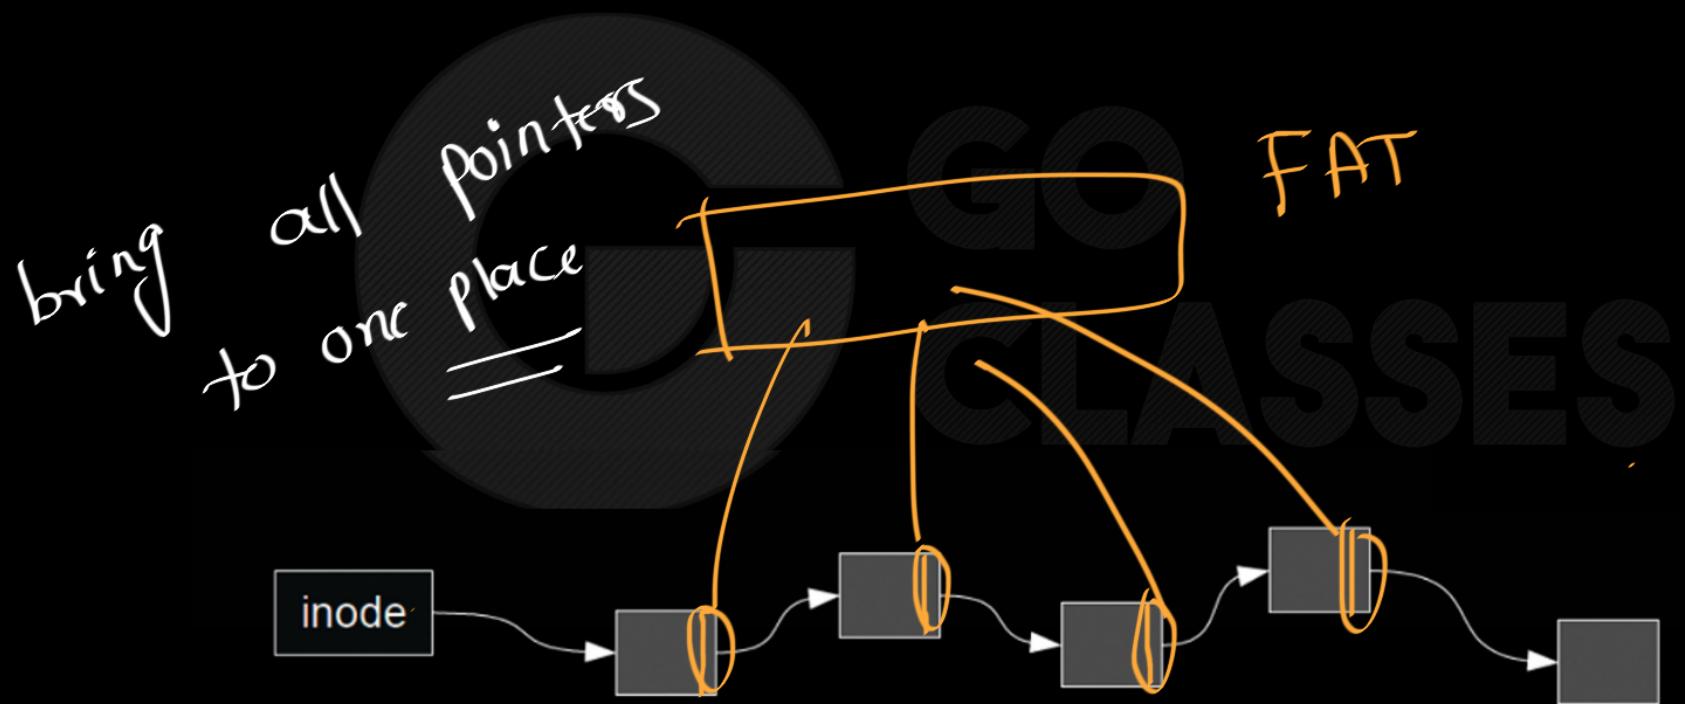
## □ Pros

- No external fragmentation
- Files can be easily grown with no limit
- Also easy to implement, though awkward to spare space for disk pointer per block

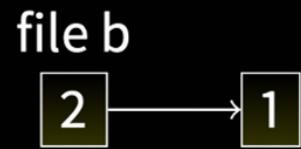
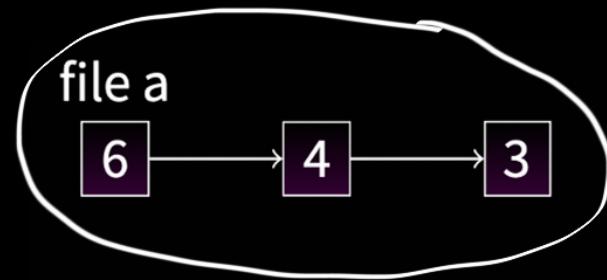
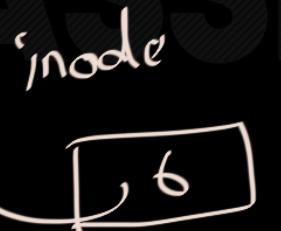
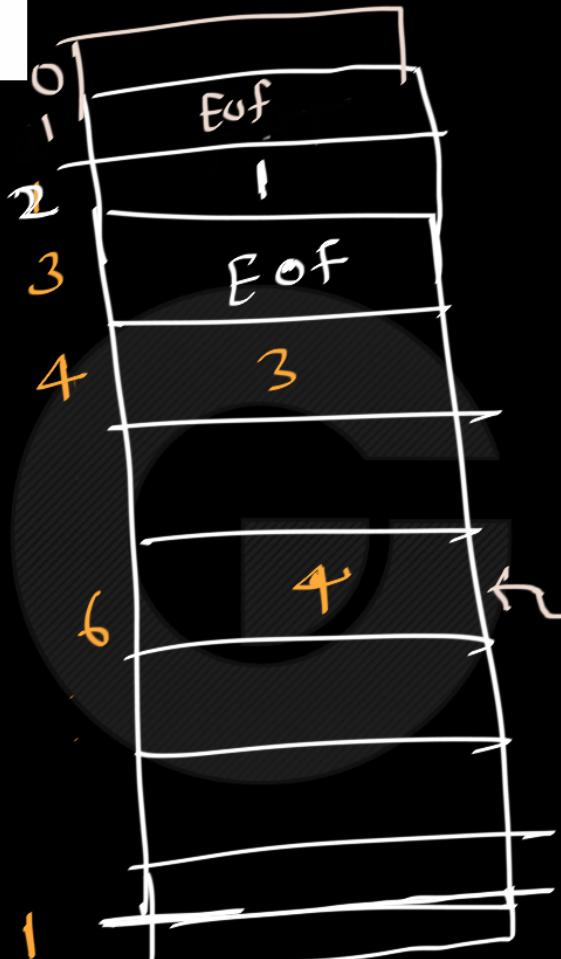
## □ Cons

- Large storage overhead (one pointer per block)
- Potentially slow sequential access
- Difficult to compute random addresses

## Variation: FAT table $\leftarrow \underline{\underline{\text{imp}}}$



No. of Sectors - 1

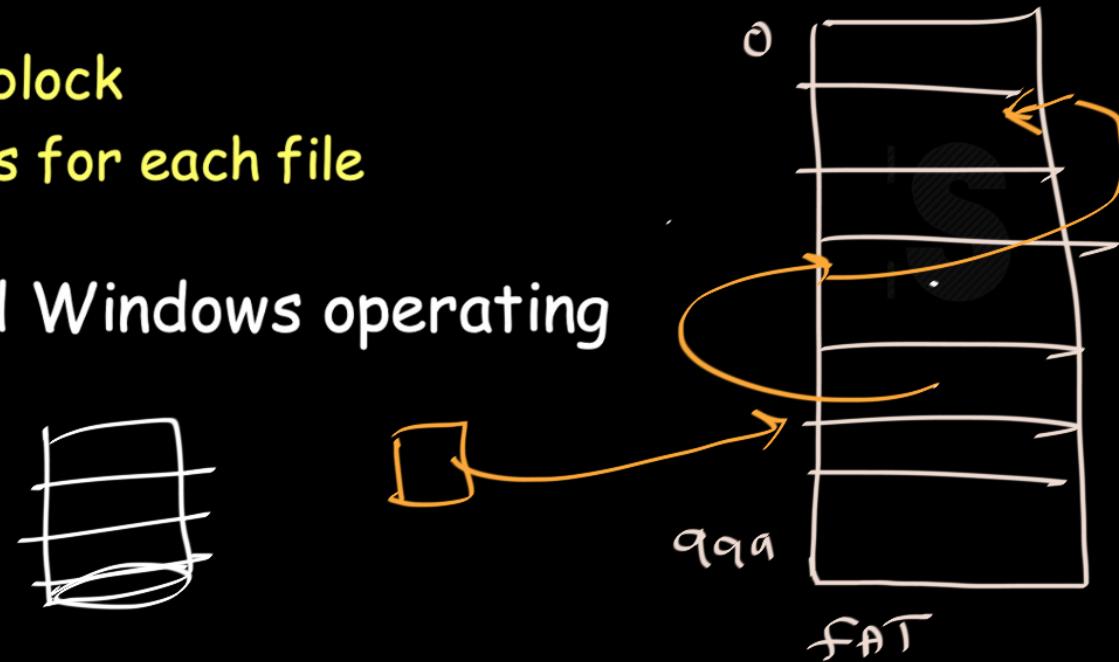


3 or 1 IO operations

1 or 2 IO operations

## Variation: FAT table $\leftarrow \underline{\text{imp}}$

- Store linked-list pointers outside block in **File-Allocation Table**
  - One entry for each block
  - Linked-list of entries for each file
- Used in MSDOS and Windows operating systems



*"inode of"*

Directory (5)

a: 6
b: 2

FAT (16-bit entries)

0	free
1	eof
2	1
3	eof
4	3
5	eof
6	4
	...

file a



file b



MM

*Being fat MM*

Still do pointer chasing

Directory (5)

a: 6
b: 2

FAT (16-bit entries)

0	free
1	eof
2	1
3	eof
4	3
5	eof
6	4
	...

file a

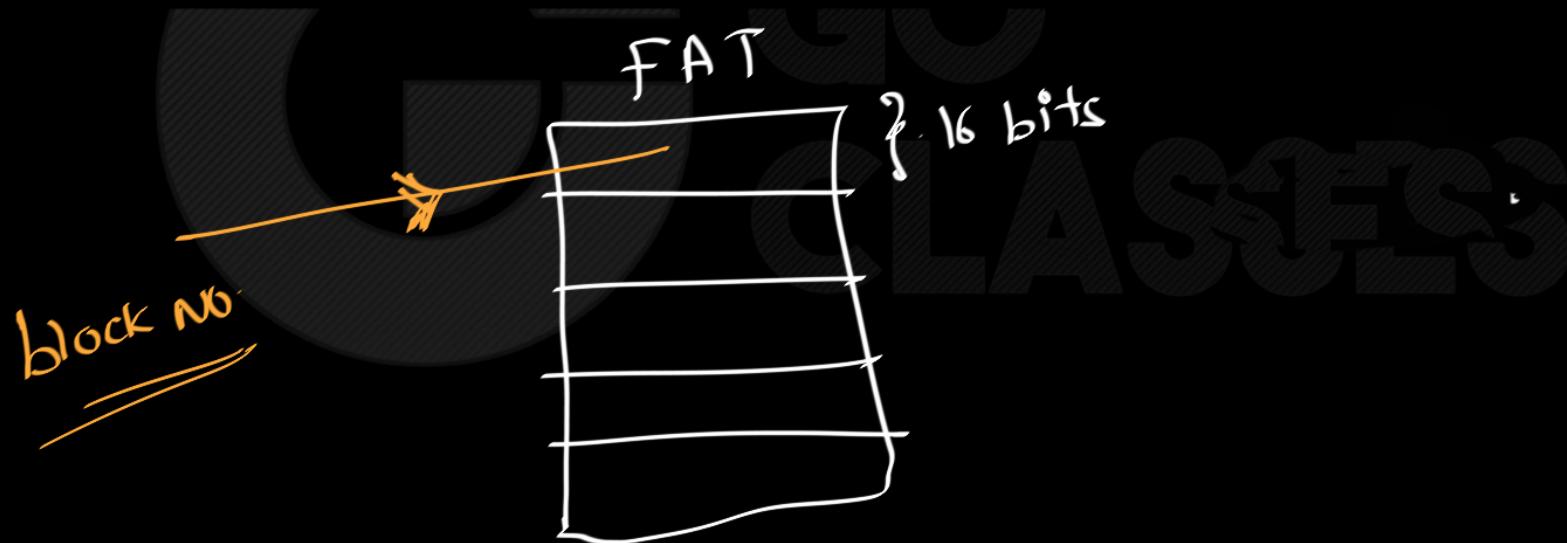


file b



Still do pointer chasing, but can bring entire FAT in MM so can be cheap compared to disk access.

- Entry size = 16 bits
  - What's the maximum size of the FAT?
  - Given a 512 byte block, what's the maximum size of FS?

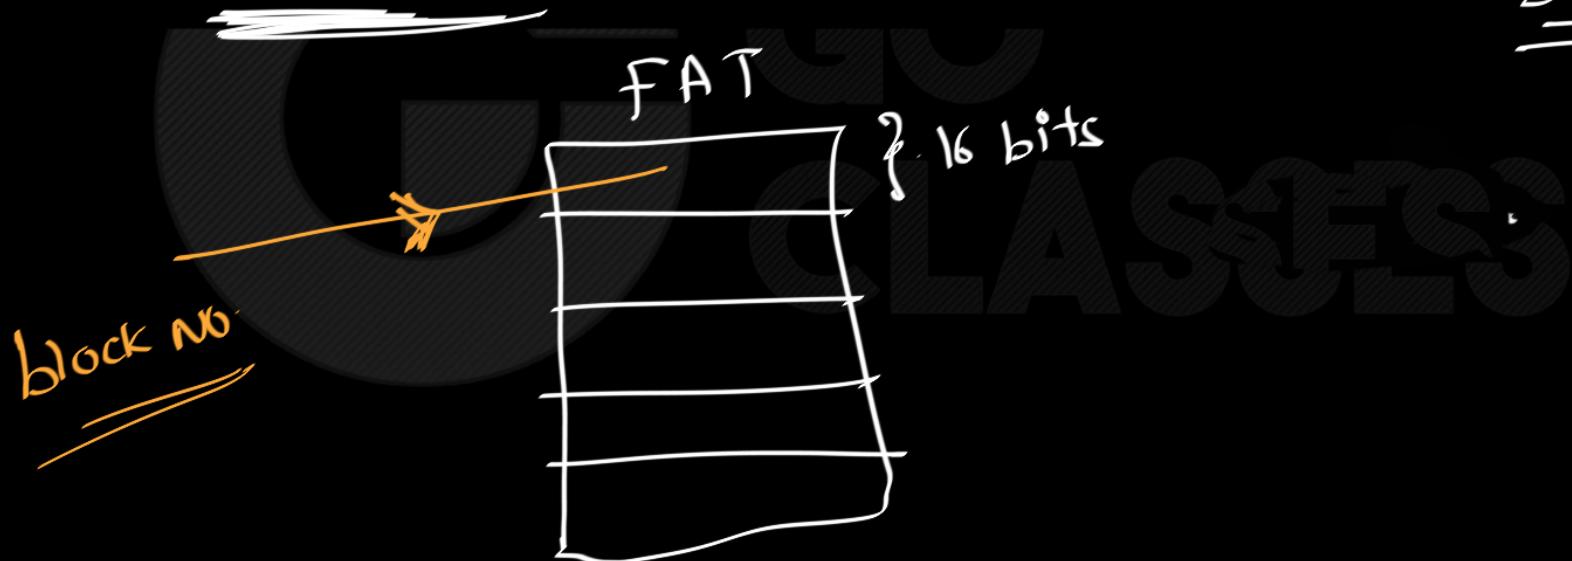


- Entry size = 16 bits

- What's the maximum size of the FAT?
- Given a 512 byte block, what's the maximum size of FS?

$2^{16}$  blocks

$2^{16} \times 512$



- **Entry size = 16 bits**

- What's the maximum size of the FAT? 65,536 entries

- Given a 512 byte block, what's the maximum size of FS? 32 MB

$2^{16}$  blocks  $\Rightarrow$

One entry for each block

$2^{16}$  entries

$$\leftarrow 2^{16} \times 512$$

$$= 2^9 \times 2^{16}$$

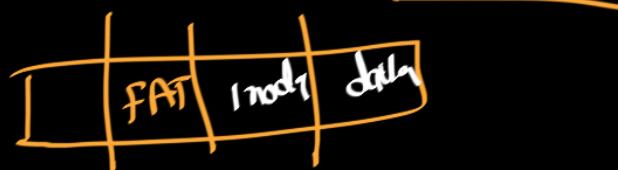
$$= 2^5 \times 2^{20}$$

$$= 32 \text{ MB}$$

An important variation on linked allocation is the use of a **file-allocation table (FAT)**. This simple but efficient method of disk-space allocation was used by the MS-DOS operating system. A section of disk at the beginning of each volume is set aside to contain the table. The table has one entry for each disk block and is indexed by block number. The FAT is used in much the same way as a linked list. The directory entry contains the block number of the first block of the file. The table entry indexed by that block number contains the block number of the next block in the file. This chain continues until it reaches the last block, which has a special end-of-file value as the table entry. An unused block is indicated by a table value of 0. Allocating a new block to a file is a simple matter of finding the first 0-valued table entry and replacing the previous end-of-file value with the address of the new block. The 0 is then replaced with the end-of-file value. An illustrative example is the FAT structure shown in Figure 12.7 for a file consisting of disk blocks 217, 618, and 339.

The FAT allocation scheme can result in a significant number of disk head seeks, unless the FAT is cached. The disk head must move to the start of the volume to read the FAT and find the location of the block in question, then move to the location of the block itself. In the worst case, both moves occur for each of the blocks. A benefit is that random-access time is improved, because the disk head can find the location of any block by reading the information in the FAT.

ALVIN



# Pros and cons

## □ Pros

- Fast random access. Only search cached FAT
- 

## □ Cons

- Large storage overhead for FAT table }
- Potentially slow sequential access

ES

12. Consider a FAT file system where disk is divided into M byte blocks, and every FAT entry can store an N bit block number. What is the maximum size of a disk partition that can be managed by such a FAT design?

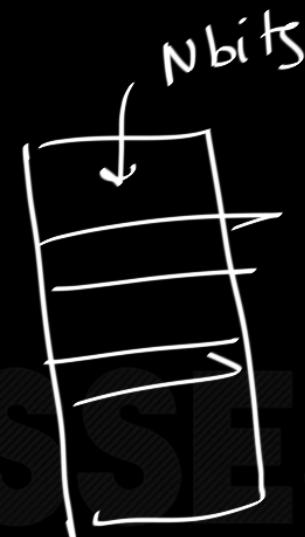
block



12. Consider a FAT file system where disk is divided into M byte blocks, and every FAT entry can store an N bit block number. What is the maximum size of a disk partition that can be managed by such a FAT design?

block size = M bytes

No. of blocks =  $2^N$



12. Consider a FAT file system where disk is divided into M byte blocks, and every FAT entry can store an N bit block number. What is the maximum size of a disk partition that can be managed by such a FAT design?

**Ans:**  $2^N * M$  bytes



The logo features a stylized circular emblem on the left composed of concentric arcs. A horizontal arrow points from the text "GO CLASSES" towards the emblem. The text "GO" is on top and "CLASSES" is below it, both in large, bold, sans-serif capital letters.

GO  
CLASSES

## GATE CSE 2014 Set 2 | Question: 20

asked in **Operating System** Sep 28, 2014 • edited Jun 26, 2018 by **kenzou**

15,059 views



47



A FAT (file allocation table) based file system is being used and the total overhead of each entry in the FAT is 4 bytes in size. Given a  $100 \times 10^6$  bytes disk on which the file system is stored and data block size is  $10^3$  bytes, the maximum size of a file that can be stored on this disk in units of  $10^6$  bytes is \_\_\_\_\_.

gatecse-2014-set2

operating-system

disk

numerical-answers

normal

file-system



## GATE CSE 2014 Set 2 | Question: 20

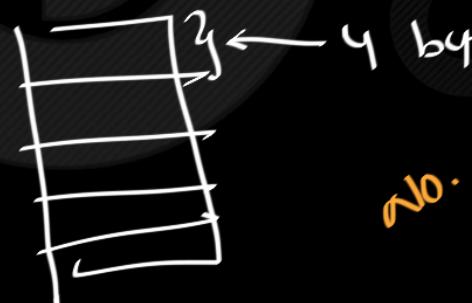
$$\frac{x}{10^6} \text{ } \left. \begin{array}{c} \text{in bytes} \\ \downarrow \\ x \end{array} \right\} \leftarrow \text{final answer}$$

asked in Operating System Sep 28, 2014 • edited Jun 26, 2018 by kenzou

15,059 views

- 47 A FAT (file allocation table) based file system is being used and the total overhead of each entry in the FAT is 4 bytes in size. Given a  $100 \times 10^6$  bytes disk on which the file system is stored and data block size is  $10^3$  bytes, the maximum size of a file that can be stored on this disk in units of  $10^6$  bytes is \_\_\_\_\_.

gatecse-2014-set2 operating-system disk numerical-answers normal file-system



$$\begin{aligned} & \text{No. of entries} \rightarrow \{ \text{No. of blocks} = \\ & \text{block size} = \underline{\underline{10^3}} \text{ bytes} \\ & 4 \times 10^5 = \underline{\underline{0.4 \times 10^6}} \end{aligned}$$

No. of entries =  $10^5$

one entry size = 4 bytes

Size of FAT =

Total disk size =

Space left to store file =

$4 \times 10^5$  bytes

$100 \times 10^6$  bytes

$= 10^3 \times 10^5$  bytes

$10^3 \times 10^5 - 4 \times 10^5$

=  $996 \times 10^5$  bytes

in units of  $10^6 \rightarrow$

$\frac{1}{10^6}$

= 99.6 MB



Each datablock will have its entry.

93

$$\text{So, Total Number of entries in the FAT} = \frac{\text{Disk Capacity}}{\text{Block size}} = \frac{100MB}{1KB} = 100K$$



Each entry takes up  $4B$  as overhead



$$\text{So, space occupied by overhead} = 100K \times 4B = 400KB = 0.4MB$$

Best answer

We have to give space to Overheads on the same file system and at the rest available space we can store data.

So, assuming that we use all available storage space to store a single file = Maximum file size =

$$\text{Total File System size} - \text{Overhead} = 100MB - 0.4MB = 99.6MB$$

 answered Jan 18, 2015 •  edited Jun 26, 2018 by Milicevic3306

 comment  Follow  share this



kalpish

# Indexed Allocation

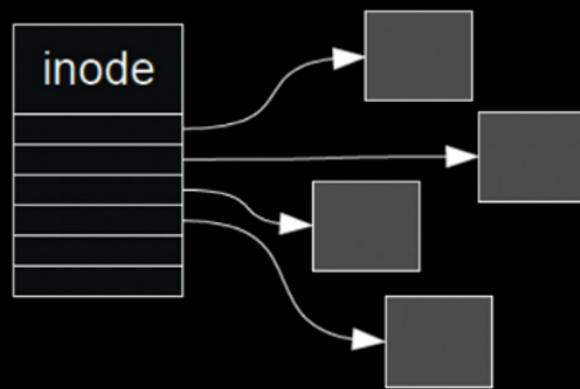
Linked allocation solves the external-fragmentation and size-declaration problems of contiguous allocation.

However, linked allocation cannot support efficient direct access, since the pointers to the blocks are scattered with the blocks themselves all over the disk and must be retrieved in order. **Indexed allocation** solves this problem by bringing all the pointers together into one location: the **index block**.

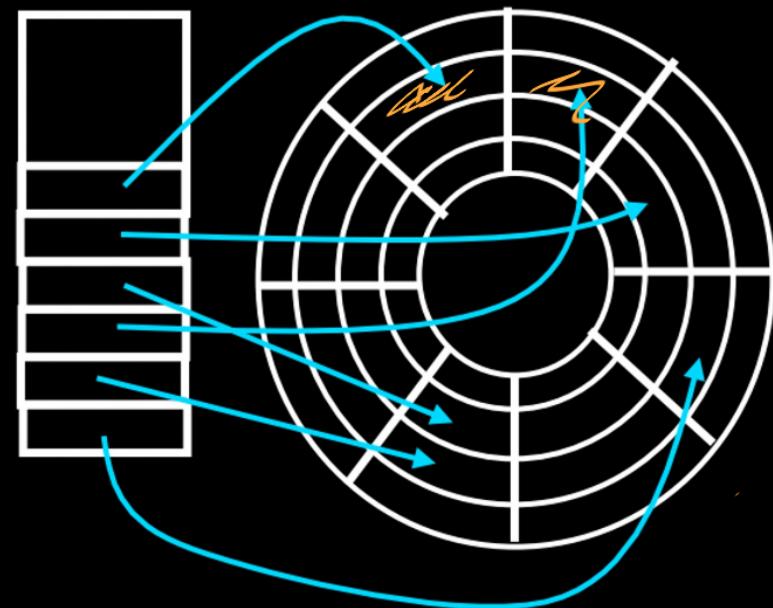
↑

GRAVIN

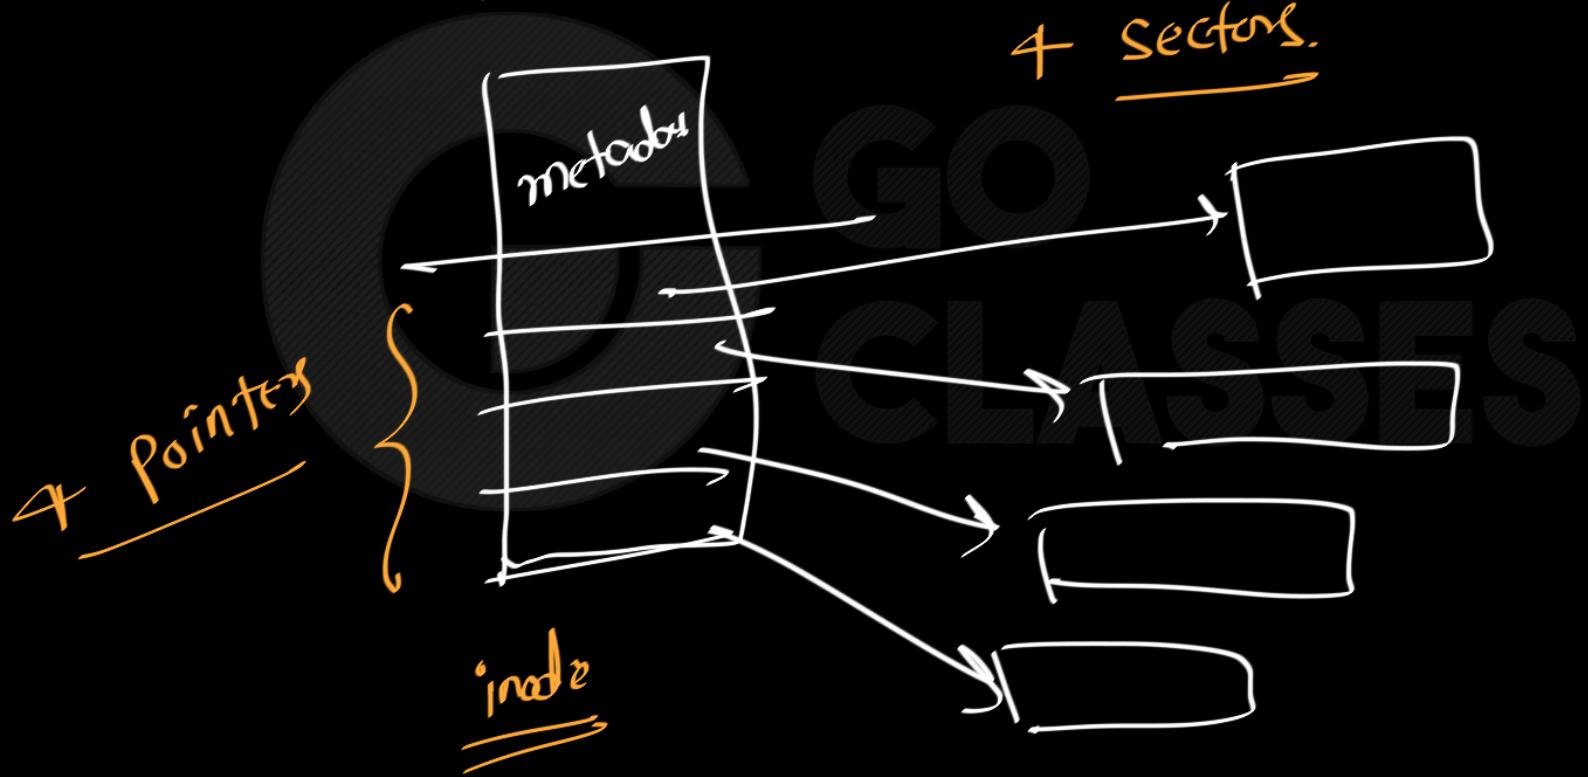
# Indexed Allocation

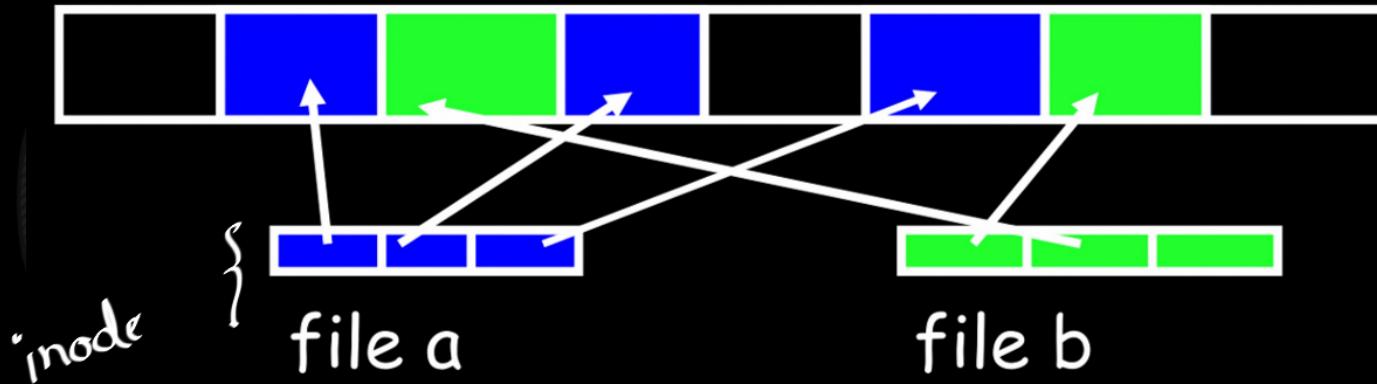


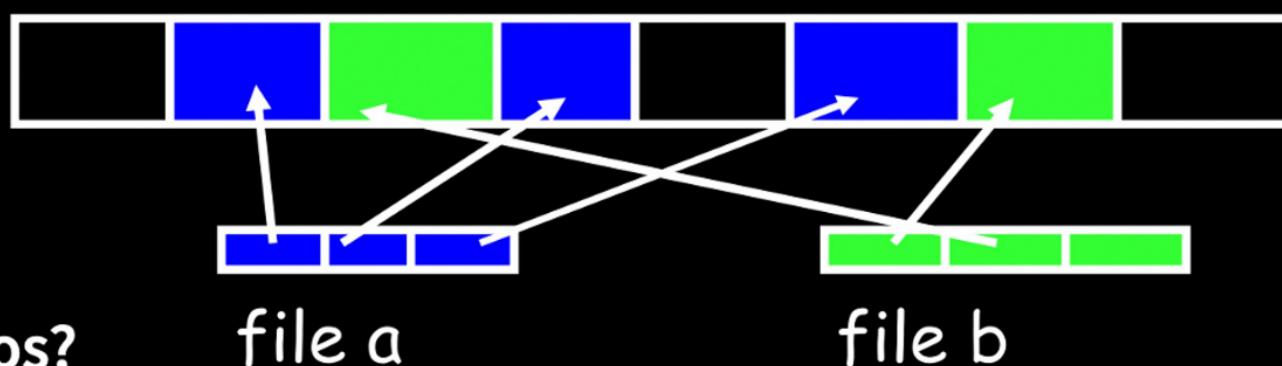
inode



if you have space for 4 pointers  
in inode then you can use at most  
4 sectors.







- Pros?

**file a**

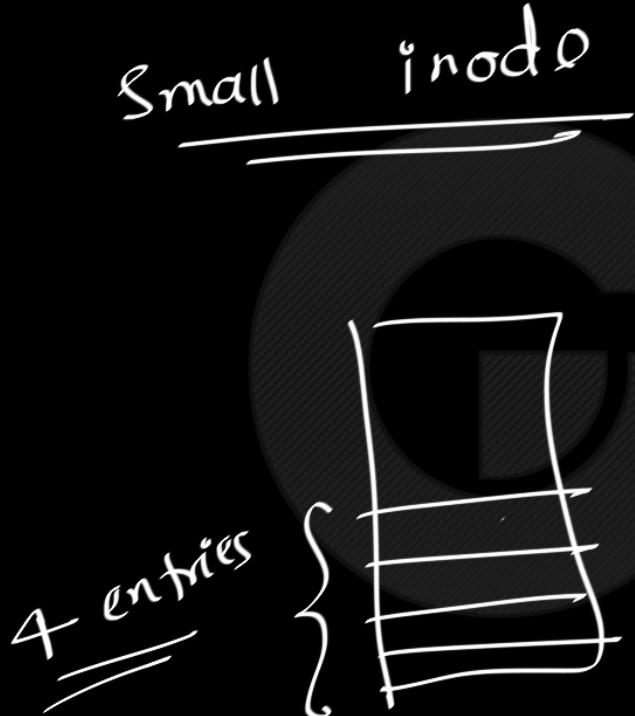
**file b**

- Both sequential and random access easy

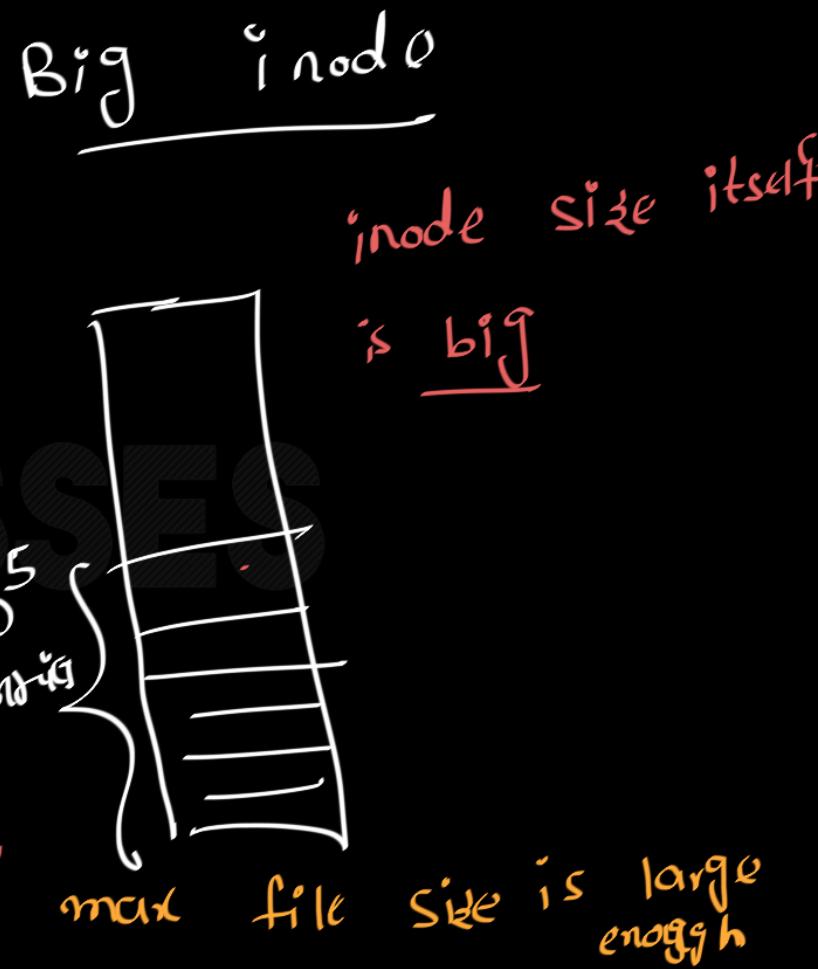
- Cons?

- Mapping table requires large chunk of contiguous space  
...Same problem we were trying to solve initially

↖ *inode size  
is Big.*



max. file size is less



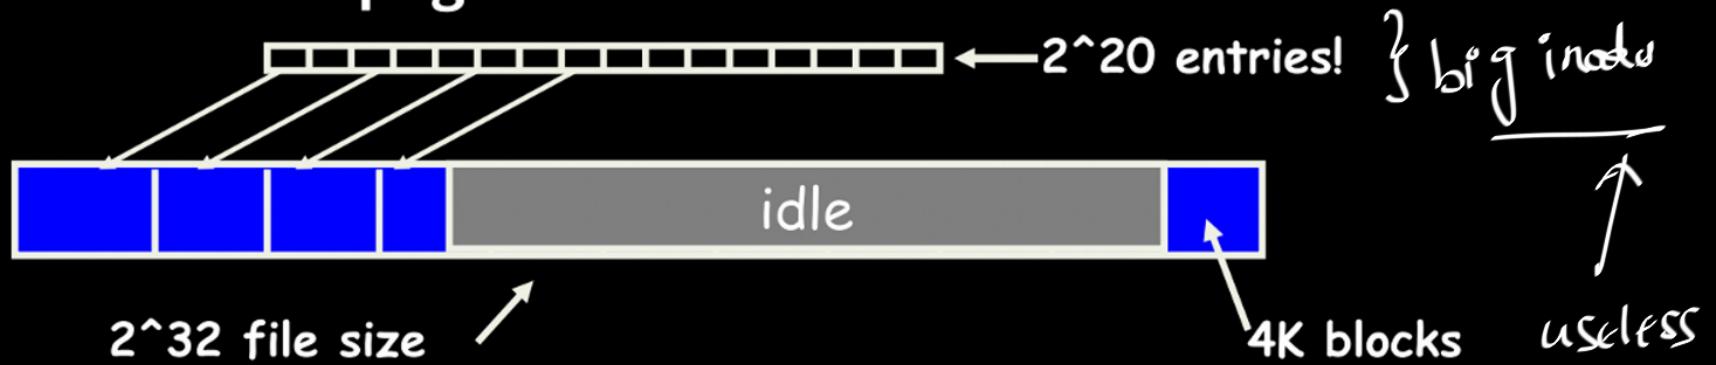
max. file size is large  
enough



Can we have a method by which  
we can have small inode size  
still have large max file size?

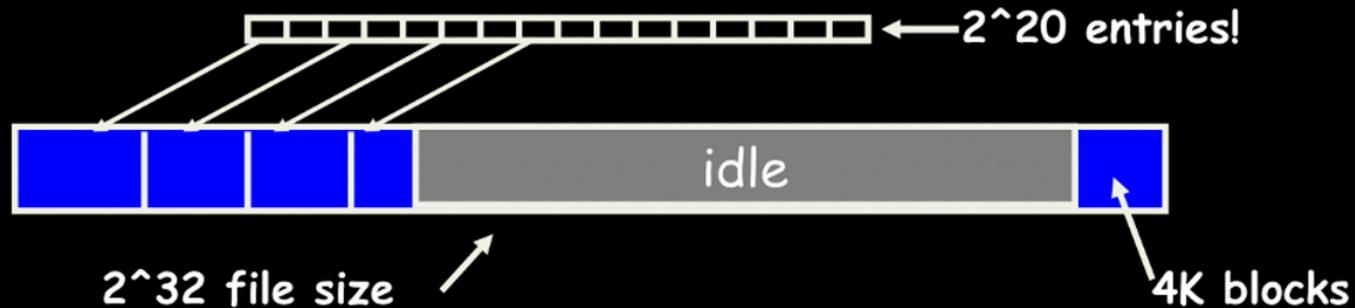
→ Yes

- Issues same as in page tables



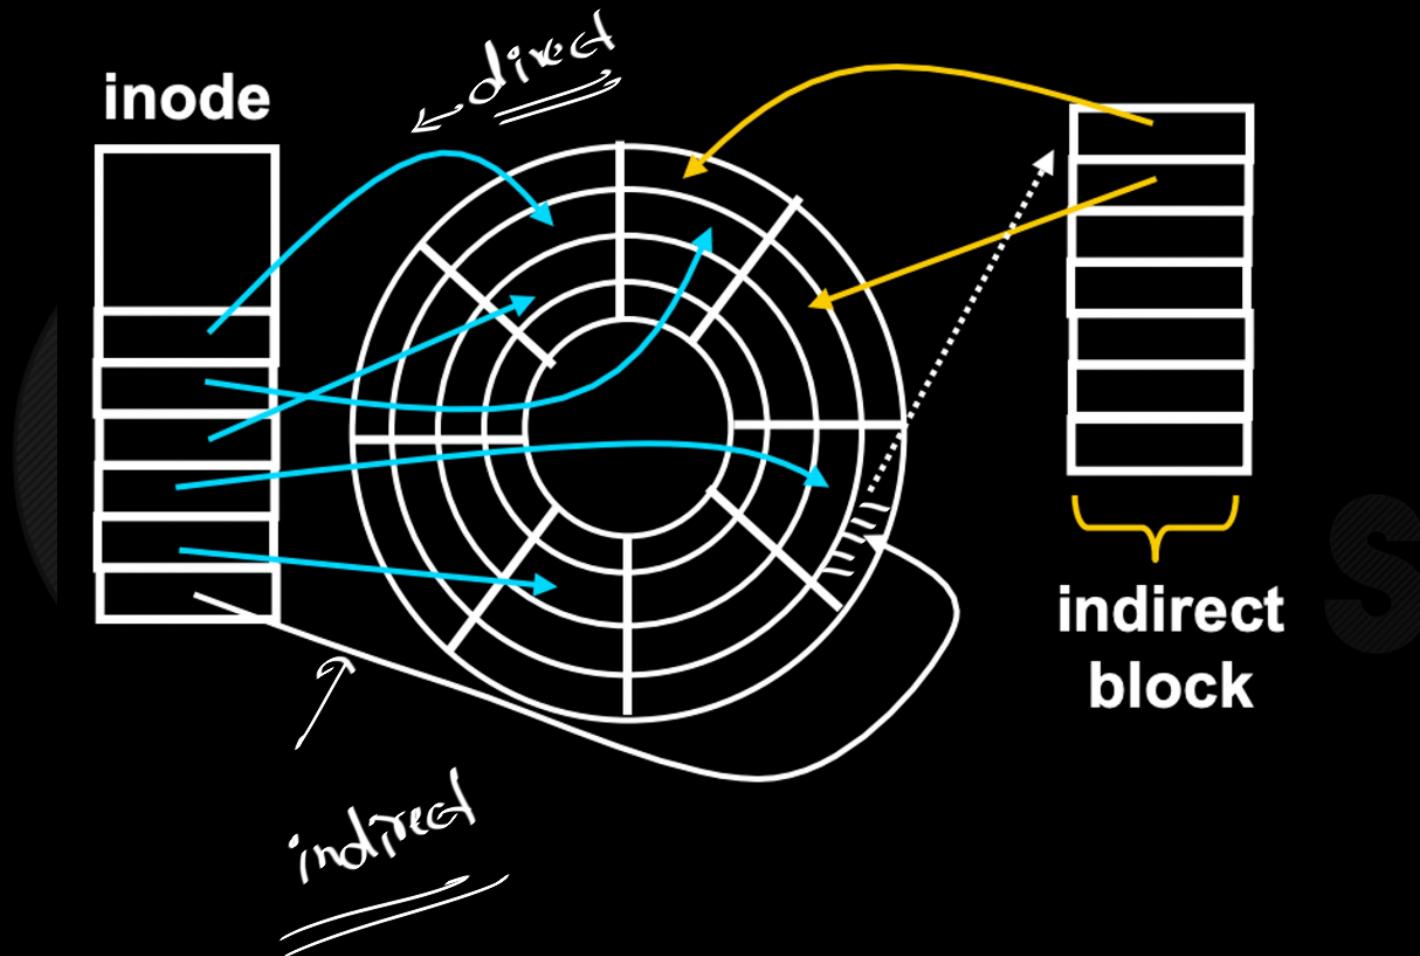
- Large possible file size = lots of unused entries ↗
  - Large actual size? table needs large contiguous disk chunk
- most of the files are  
small
- for  
most  
of the  
files

- Issues same as in page tables

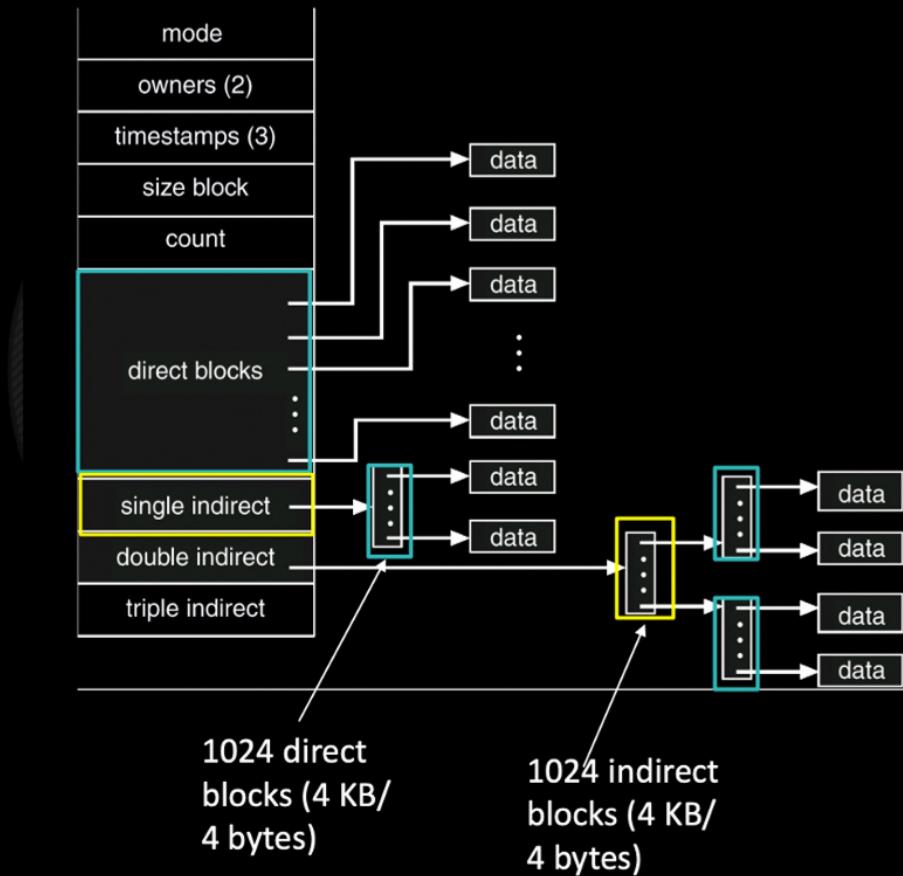


- Large possible file size = lots of unused entries
- Large actual size? table needs large contiguous disk chunk
- Solve identically: small regions with index array, this array with another array, ... Downside?

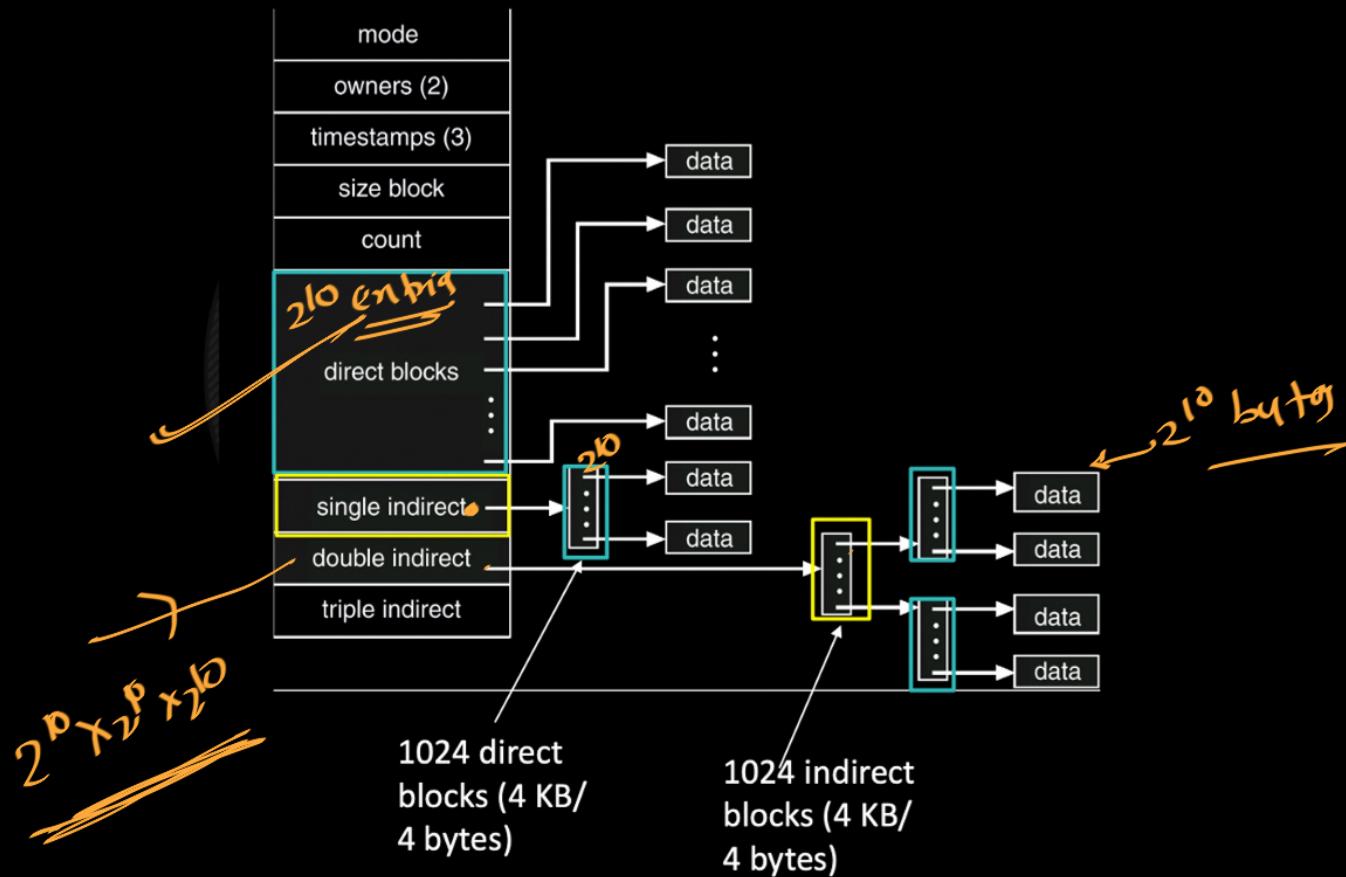




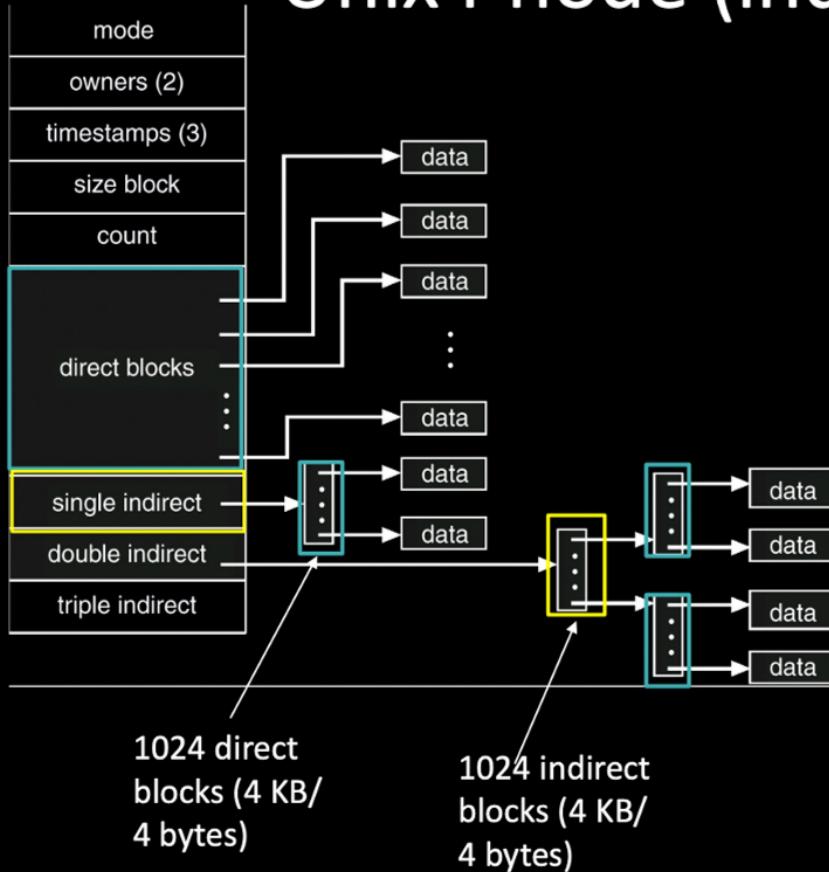
# Unix i-node (index node)



# Unix i-node (index node)



# Unix i-node (index node)



- Small files can be accessed quickly
- If each block is 4KB, each block address is 4 bytes
  - First 48KB of file reachable from 12 direct blocks
  - Next 4MB available from single-indirect blocks
  - Next 4GB available from double-indirect blocks
  - Next 4TB available through the triple-indirect blocks
- To access any one block of a 4TB file, how many disk accesses at most do we need?

4. Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.
- (a) How large of a disk can this file system support?

(b) What is the maximum file size?

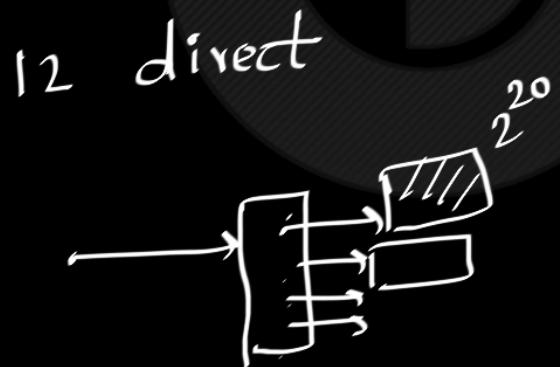
4. Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.

(a) How large of a disk can this file system support?

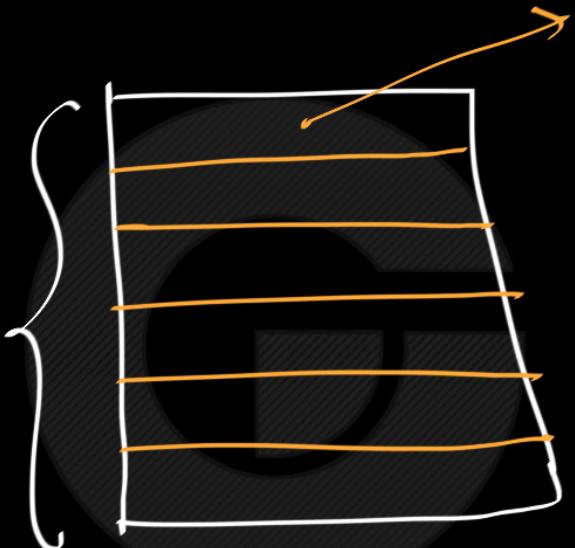
$$\text{block size} = 2^{11} \text{ B}$$

(b) What is the maximum file size?

$$2^{32} \text{ blocks disk}$$



# CLASSES

$2^9$ 


$2^{11}$  Bytes  
block

block address = 32 bits  
 -  
 [ what are no. of bits we  
 require to address one block?  
 $= 32 \text{ bits} = 4 \text{ B}$

each block can  
 hold =  $\frac{2^{11}}{2^2} = \underline{\underline{2^9 \text{ entries}}}$

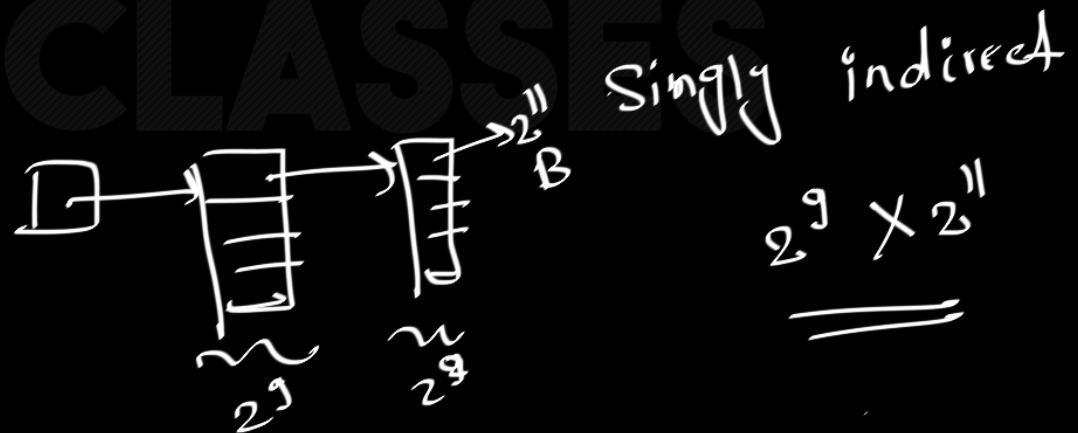
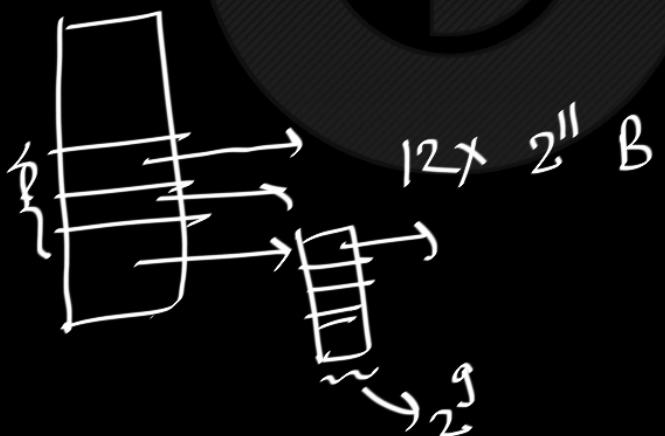


4. Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.

(a) How large of a disk can this file system support?

(b) What is the maximum file size?

$$\begin{aligned} \text{block size} &= 2^{11} \text{ B} \\ &= 2^9 \text{ entries} \end{aligned}$$





$$\left[ 12x^{2^{11}} + 1x^{2^9}x^{2^{11}} + 1x^{2^9}x^{2^9}x^{2^{11}} + \dots \right]$$

A mathematical expression enclosed in large curly braces. The expression consists of three terms separated by plus signs. The first term is  $12x^{2^{11}}$ . The second term is  $1x^{2^9}x^{2^{11}}$ . The third term is  $1x^{2^9}x^{2^9}x^{2^{11}}$ . Ellipses at the end of the braces indicate that the sequence continues.

GO  
CLASSES

4. Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.

- (a) How large of a disk can this file system support?

$2^{32}$  blocks

$$2^{32} \text{ blocks} \times 2^{11} \text{ bytes/block} = 2^{43} = 8 \text{ Terabytes.}$$

- (b) What is the maximum file size?

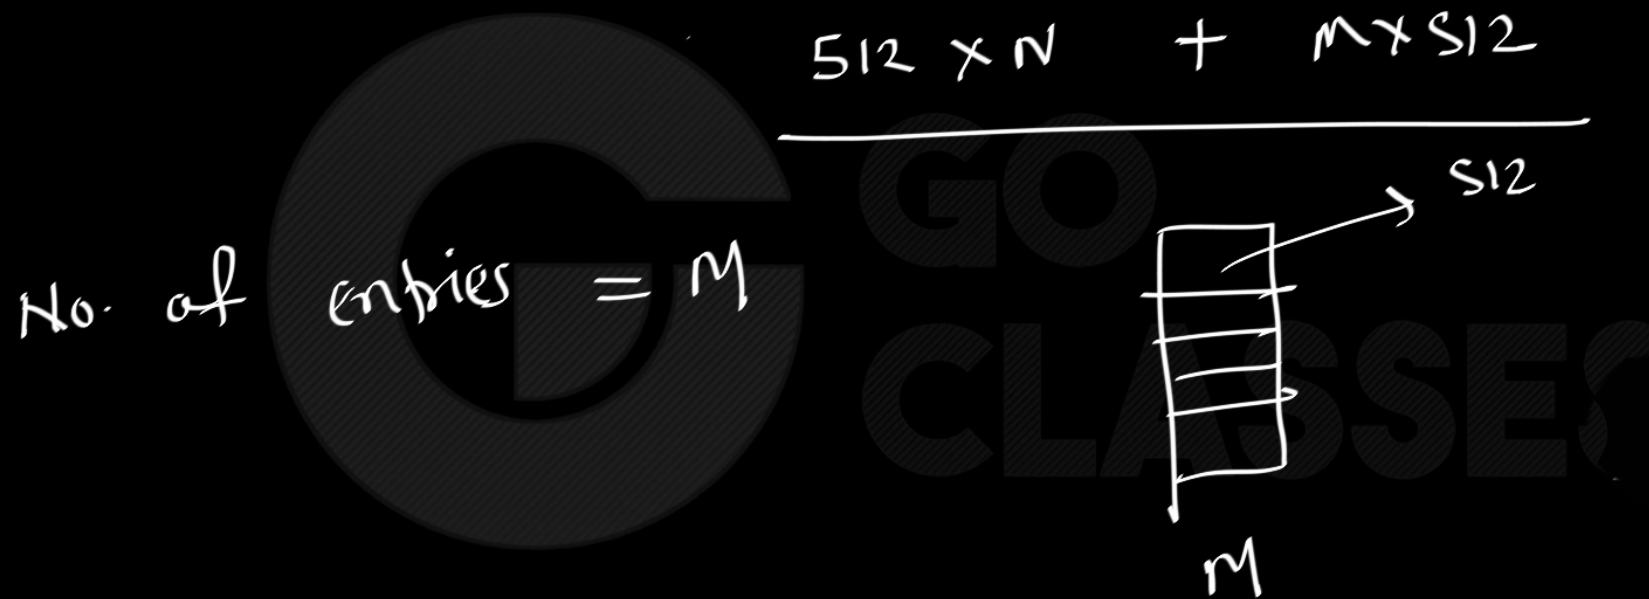
There are 512 pointers per block (i.e. 512 4-byte pointers in 2048 byte block), so:  
 $\text{blockSize} \times (\text{numDirect} + \text{numIndirect} + \text{numDoublyIndirect} + \text{numTriplyIndirect})$

$$\begin{aligned} 2048 \times (12 + 512 + 512^2 + 512^3) &= 2^{11} \times (2^2 \times 3 + 2^9 + 2^{9 \times 2} + 2^{9 \times 3}) \\ &= 2^{13} \times 3 + 2^{20} + 2^{29} + 2^{38} \\ &= 24K + 513M + 256G \end{aligned}$$

11. Consider a file system with 512-byte blocks. Assume an inode of a file holds pointers to N direct data blocks, and a pointer to a single indirect block. Further, assume that the single indirect block can hold pointers to M other data blocks. What is the maximum file size that can be supported by such an **inode** design?



11. Consider a file system with 512-byte blocks. Assume an inode of a file holds pointers to N direct data blocks, and a pointer to a single indirect block. Further, assume that the single indirect block can hold pointers to M other data blocks. What is the maximum file size that can be supported by such an **inode** design?



11. Consider a file system with 512-byte blocks. Assume an inode of a file holds pointers to N direct data blocks, and a pointer to a single indirect block. Further, assume that the single indirect block can hold pointers to M other data blocks. What is the maximum file size that can be supported by such an **inode** design?

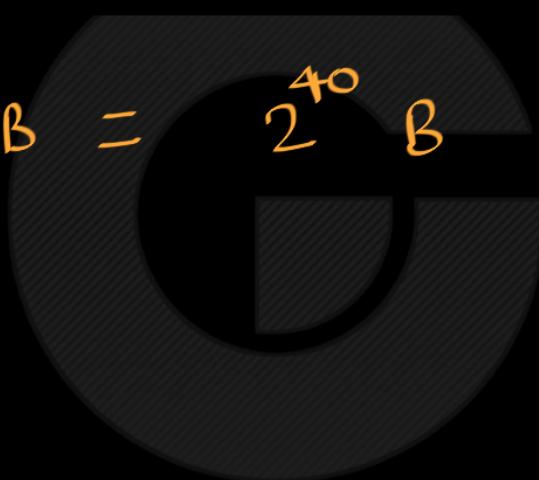
**Ans:**  $(N+M)*512$  bytes



13. Consider a secondary storage system of size 2 TB, with 512-byte sized blocks. Assume that the filesystem uses a multilevel inode datastructure to track data blocks of a file. The inode has 64 bytes of space available to store pointers to data blocks, including a single indirect block, a double indirect block, and several direct blocks. What is the maximum file size that can be stored in such a file system?

$$1TB = 2^{40} B$$

$2^{41}$  Bytes  $\leftarrow 2TB$



13. Consider a secondary storage system of size 2 TB, with 512-byte sized blocks. Assume that the filesystem uses a multilevel inode datastructure to track data blocks of a file. The inode has 64 bytes of space available to store pointers to data blocks, including a single indirect block, a double indirect block, and several direct blocks. What is the maximum file size that can be stored in such a file system?



$$\left\{ \begin{array}{l} \text{Block size} = 2^9 \text{ B} \\ \text{entry size} = 32 \text{ bits} = 2^2 \text{ bytes} \end{array} \right.$$

$$\left( \begin{array}{l} \text{No. of entries} = \end{array} \right)$$

$$2^7$$

$\equiv$   
each double can point to  
this much data

$$\text{No. of pointers}$$

$$\frac{2^7 \times 2^7 \times 2^9}{\equiv}$$

in 64 Bytes, we can store

$$\# \text{ of pointers} =$$

$$\frac{64 \text{ B}}{4 \text{ B}} = 16$$

$$\underbrace{16 \times \frac{2^7 \times 2^7 \times 2^9}{\equiv}}_{\left\{ \begin{array}{l} \text{Answer} \end{array} \right\}}$$

$$= 2^7 \text{ B}$$

13. Consider a secondary storage system of size 2 TB, with 512-byte sized blocks. Assume that the filesystem uses a multilevel inode datastructure to track data blocks of a file. The inode has 64 bytes of space available to store pointers to data blocks, including a single indirect block, a double indirect block, and several direct blocks. What is the maximum file size that can be stored in such a file system?



**Ans:** Number of data blocks =  $2^{41}/2^9 = 2^{32}$ , so 32 bits or 4 bytes are required to store the number of a data block.

Number of data block pointers in the inode =  $64/4 = 16$ , of which 14 are direct blocks. The single indirect block stores pointers to  $512/4 = 128$  data blocks. The double indirect block points to 128 single indirect blocks, which in turn point to 128 data blocks each.

3. (28 points total) File Systems.

- a. (8 points) Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.
- i) (4 points) How large of a disk can this file system support?



- ii) (4 points) What is the maximum file size?

3. (28 points total) File Systems.

a. (8 points) Consider a file system with 2048 byte blocks and 32-bit disk and file block pointers. Each file has 12 direct pointers, a singly-indirect pointer, a doubly-indirect pointer, and a triply-indirect pointer.

i) (4 points) How large of a disk can this file system support?

$$2^{32} \text{ blocks} \times 2^{11} \text{ bytes/block} = 2^{43} = 8 \text{ Terabytes.}$$

-3 missing block or disk size in calculation

ii) (4 points) What is the maximum file size?

*There are 512 pointers per block (i.e. 512 4-byte pointers in 2048 byte block), so:*

$$\begin{aligned} \text{blockSize} \times (\text{numDirect} + \text{numIndirect} + \text{numDoubly-indirect} + \text{numTriply indirect}) \\ 2048 \times (12 + 512 + 512^2 + 512^3) = 2^{11} \times (2^2 \times 3 + 2^9 + 2^{9 \times 2} + 2^{9 \times 3}) \\ = 2^{13} \times 3 + 2^{20} + 2^{29} + 2^{38} \\ = 24K + 513M + 256G \end{aligned}$$

# Question

b. (15 points) Large files. We have seen different filesystems that support fairly large files. Now let's see just how large a file various types of filesystems can support. Assume, for all of the questions in this part, that filesystem blocks are 4 KBytes.  
*Show your solutions in unsimplified form for partial credit.*

$$\beta S = 4 \text{ KB}$$

i) (3 points) Consider a really simple filesystem, **directfs**, where each inode only has 10 direct pointers, each of which can point to a single file block. Direct pointers are 32 bits in size (4 bytes). What is the maximum file size for **directfs**?



ii) (3 points) Consider a filesystem, called **extentfs**, with a construct called an extent. Extents have a pointer (base address) and a length (in blocks). Assume the length field is 8 bits (1 byte). Assuming that an inode has exactly one extent. What is the maximum file size for **extentfs**?

SES

iii) (3 points) Consider a filesystem that uses direct pointers, but also adds indirect pointers and double-indirect pointers. We call this filesystem, **indirectfs**. Specifically, an inode within **indirectfs** has 1 direct pointer, 1 indirect pointer, and 1 doubly-indirect pointer field. Pointers, as before, are 4 bytes (32 bits) in size. What is the maximum file size for **indirectfs**?

# Question

b. (15 points) Large files. We have seen different filesystems that support fairly large files. Now let's see just how large a file various types of filesystems can support. Assume, for all of the questions in this part, that filesystem blocks are 4 KBytes. *Show your solutions in unsimplified form for partial credit.*

i) (3 points) Consider a really simple filesystem, **directfs**, where each inode only has 10 direct pointers, each of which can point to a single file block. Direct pointers are 32 bits in size (4 bytes). What is the maximum file size for **directfs**?

$$4 \times 10 = 40 \text{ kB}$$

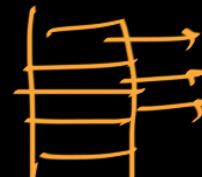
ii) (3 points) Consider a filesystem, called **extentfs**, with a construct called an extent. Extents have a pointer (base address) and a length (in blocks). Assume the length field is 8 bits (1 byte). Assuming that an inode has exactly one extent. What is the maximum file size for **extentfs**?

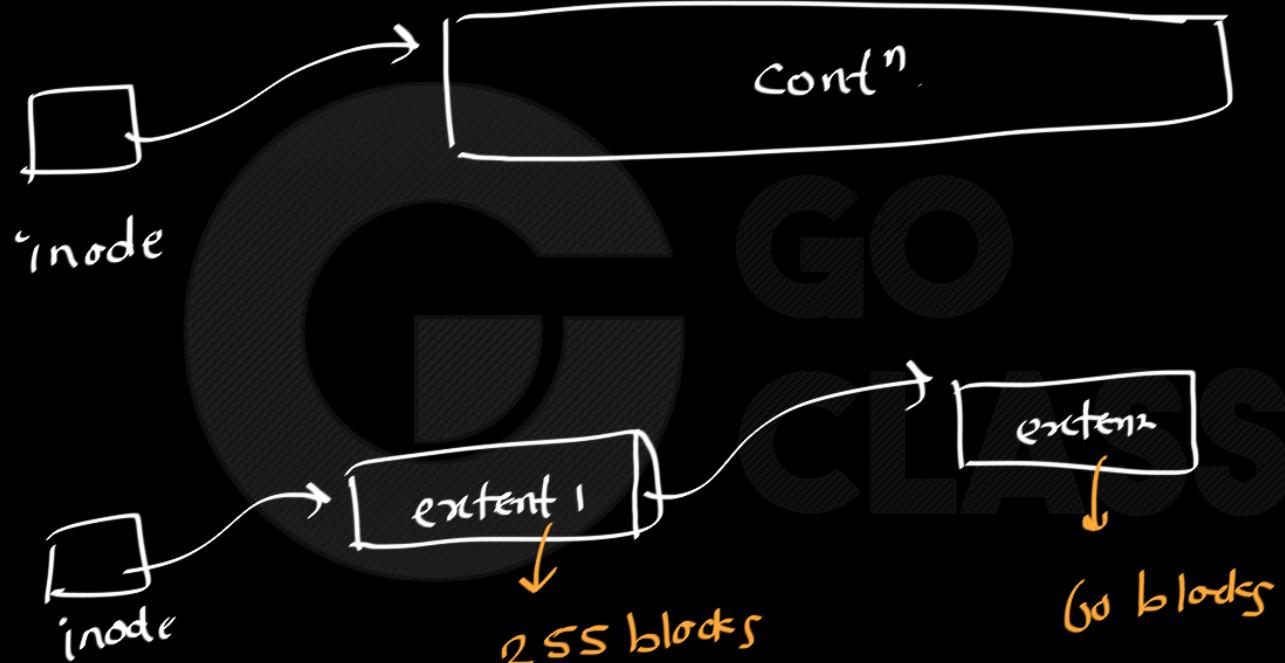
$$\underline{255} \times \underline{4 kB}$$



iii) (3 points) Consider a filesystem that uses direct pointers, but also adds indirect pointers and double-indirect pointers. We call this filesystem, **indirectfs**. Specifically, an inode within **indirectfs** has 1 direct pointer, 1 indirect pointer, and 1 doubly-indirect pointer field. Pointers, as before, are 4 bytes (32 bits) in size. What is the maximum file size for **indirectfs**?

$$\beta S = 4 \text{ kB}$$





max size of  
any extent  
= 255 blocks

$255 \times 4KB$

iii) (3 points) Consider a filesystem that uses direct pointers, but also adds indirect pointers and double-indirect pointers. We call this filesystem, **indirectfs**.

Specifically, an inode within **indirectfs** has 1 direct pointer, 1 indirect pointer, and 1 doubly-indirect pointer field. Pointers, as before, are 4 bytes (32 bits) in size. What is the maximum file size for **indirectfs**?

$$\begin{aligned}
 & 1 \text{ block} + 2^{10} \text{ blocks} + 2^{10} \times 2^{10} \text{ blocks} \\
 & = (1 + 2^{10} + 2^{10} \times 2^{10}) \text{ blocks} \\
 & = (1 + 2^{10} + 2^{20}) \times 4 \times 2^{10} \beta \\
 & = 2^{10} \underbrace{\text{entries}}
 \end{aligned}$$

iv) (3 points) Consider a compact file system, called **compactfs**, tries to save as much space as possible within the inode. Thus, to point to files, it stores only a single 32-bit pointer to the first block of the file. However, blocks within **compactfs** store 4,092 bytes of user data and a 32-bit next field (much like a linked list), and thus can point to a subsequent block (or to NULL, indicating there is no more data). First, draw a picture of an inode and a file that is 10 KBytes in size.



$$\begin{aligned} \text{No. of data blocks} &= 2^{32} \\ \text{Block size} &= 2^{12} \\ & (= 4096) \\ & \downarrow \\ & 4092 + 4 \end{aligned}$$

v) (3 points) What is the maximum file size for **compactfs** (assuming no other restrictions on file sizes)?

$$\text{Disk size} = 2^{44} \text{ B}$$

$$\frac{\text{No. of blocks}}{=} = 2^{32}$$

$$\text{No. of blocks} = 2^{32}$$

$$\text{No. of pointers} = 2^{32}$$

$$\begin{aligned}\text{No. of bytes wasted} &= 2^{32} \times 4 \text{ bytes} \\ &= \underline{\underline{2^{34} B}} \quad \xrightarrow{\text{wasted in pointers}}\end{aligned}$$

$$\text{max. file size} = 2^{44} - 2^{34}$$

# Solution

b. (15 points) Large files. We have seen different filesystems that support fairly large files. Now let's see just how large a file various types of filesystems can support. Assume, for all of the questions in this part, that filesystem blocks are 4 KBytes.

*Show your solutions in unsimplified form for partial credit.*

i) (3 points) Consider a really simple filesystem, **directfs**, where each inode only has 10 direct pointers, each of which can point to a single file block. Direct pointers are 32 bits in size (4 bytes). What is the maximum file size for **directfs**?

*The maximum **directfs** file size is  $10 \times 4 \text{ KByte} = 40 \text{ KByte}$*

ii) (3 points) Consider a filesystem, called **extentfs**, with a construct called an extent. Extents have a pointer (base address) and a length (in blocks). Assume the length field is 8 bits (1 byte). Assuming that an inode has exactly one extent. What is the maximum file size for **extentfs**?

*The maximum **extentfs** file size is  $(2^8 - 1) \times 4 \text{ KByte} = 255 \times 4 \text{ KByte} = 1 \text{ MByte}$*

SES

iii) (3 points) Consider a filesystem that uses direct pointers, but also adds indirect pointers and double-indirect pointers. We call this filesystem, **indirectfs**.

Specifically, an inode within **indirectfs** has 1 direct pointer, 1 indirect pointer, and 1 doubly-indirect pointer field. Pointers, as before, are 4 bytes (32 bits) in size. What is the maximum file size for **indirectfs**?

*The maximum **indirectfs** file size is  $(1 + 1024 + (1024 \times 1024)) \times 4 \text{ KByte} = 4 \text{ GB} + 4100 \text{ KByte}$*



- iv) (3 points) Consider a compact file system, called **compactfs**, tries to save as much space as possible within the inode. Thus, to point to files, it stores only a single 32-bit pointer to the first block of the file. However, blocks within **compactfs** store 4,092 bytes of user data and a 32-bit next field (much like a linked list), and thus can point to a subsequent block (or to NULL, indicating there is no more data). First, draw a picture of an inode and a file that is 10 KBytes in size.



SSES

- v) (3 points) What is the maximum file size for **compactfs** (assuming no other restrictions on file sizes)?

*The maximum **compactfs** file size is  $\underbrace{2^{32} \times 4 \text{ KByte}}_{2^{44}} - 2^{32} \times 4 \text{ byte} = 16 \text{ TByte}$*

**Problem 3d[4pts]:** Consider a file system with 4096 byte blocks and 32-bit disk and file block pointers. Each file has 13 direct pointers, 4 singly-indirect pointers, a doubly-indirect pointer, and a triply-indirect pointer. In the following be explicit about your work:

a) What is the maximum disk size that can be supported? Explain.

b) What is the maximum file size? Explain.

c) Give some reasonable assumptions and compute the number of **inodes** that can fit into a disk block.

**Problem 3d[4pts]:** Consider a file system with 4096 byte blocks and 32-bit disk and file block pointers. Each file has 13 direct pointers, 4 singly-indirect pointers, a doubly-indirect pointer, and a triply-indirect pointer. In the following be explicit about your work:

- a) What is the maximum disk size that can be supported? Explain.

$$BS = 4096 \text{ B}$$

$$2^{32} \times 2^{12}$$

$$\text{No. of blocks} = 2^{32}$$

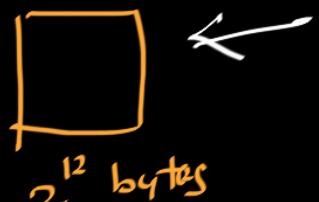
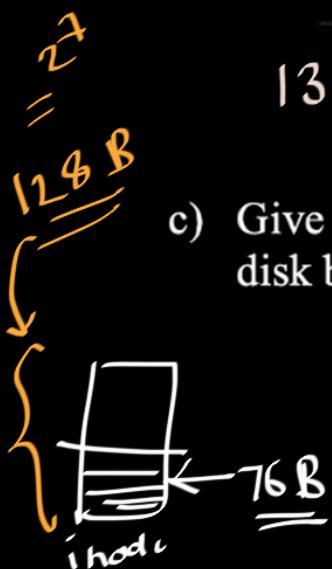
- b) What is the maximum file size? Explain.

No. of entries in one

$$13 + 4 \times 2^{10} + 1 \times 2^{10} \times 2^{10} + 2^{10} \times 2^{10} \times 2^{10}$$

$$\text{block} = \frac{4096 \text{ B}}{32 \text{ bits}} = 2^{10}$$

- c) Give some reasonable assumptions and compute the number of **inodes** that can fit into a disk block.



No. of inodes

inode has 19 pointers

$$19 \times 4 \text{ B} = 76 \text{ B}$$

Pointers  $\rightarrow$

$$\text{DISK block size} = 2^{12} \text{ B}$$

i node size = 76 B in pointers + something extra

$$= 128 \text{ B (assume)}$$

$$\text{No. of inodes in disk block} = \frac{2^{12}}{2^7} = 2^5 = 32 \text{ inodes}$$



**Problem 3d[4pts]:** Consider a file system with 4096 byte blocks and 32-bit disk and file block pointers. Each file has 13 direct pointers, 4 singly-indirect pointers, a doubly-indirect pointer, and a triply-indirect pointer. In the following be explicit about your work:

- a) What is the maximum disk size that can be supported? Explain.

*Since block pointers are 32 bits, total size =  $2^{32} \times 4096 = 2^{44}$  bytes maximum.*

- b) What is the maximum file size? Explain.

*Since pointers are 32 bits (4 bytes), one indirect block can address 1024 blocks.  
Thus, maximum file =  $4096 \times (13 + 4 \times 1024 + 1024^2 + 1024^3) = 4402358308864$  bytes*

- c) Give some reasonable assumptions and compute the number of **inodes** that can fit into a disk block.

*An inode has  $13 + 4 + 1 + 1 = 19$  pointers = 76bytes. Assuming various other file information such as ACLs, etc, we round up to 128bytes and can fit 32 inodes/block*

In lecture three file descriptor structures were discussed: (a) Indexed files. (b) Linked files. (c) Contiguous (extent-based) allocation.

Each of the structures has its advantages and disadvantages depending on the goals for the file system and the expected file access pattern. For each of the following situations, rank the three structures in order of preference. Be sure to include the justification for your rankings.

(a) You have a file system where the most important criteria is the performance of sequential access to very large files.

(b) You have a file system where the most important criteria is the performance of random access to very large files.

SES

(a) You have a file system where the most important criteria is the performance of sequential access to very large files.

1. c (extent-based)
2. b (linked)
3. a (indexed)

It is easy to see that (c) is the best structure for sequential access to very large files, since in (c) files are contiguously allocated and the next block to read is physically the next on the disk. No seek time to find the next block, and each block will be read sequentially as the disk head moves.

Both (b) and (a) require some look up operation in order to know where the next block is. However, (b) may be slightly more expensive, since for "very large files", multiple disk accesses are required to read the indirect blocks.

(b) You have a file system where the most important criteria is the performance of random access to very large files.

1. c (extent-based)
2. a (indexed)
3. b (linked)

(c) is still the best structure here: just need to use an offset.

(a) will probably need to look at some levels of indirect blocks in order to find the right block to access (we are dealing with very large files).

(b) is absolutely the worst structure. In fact, in order to find a random block, we will need to traverse a linked list of blocks, which will take a time linear in the offset size.

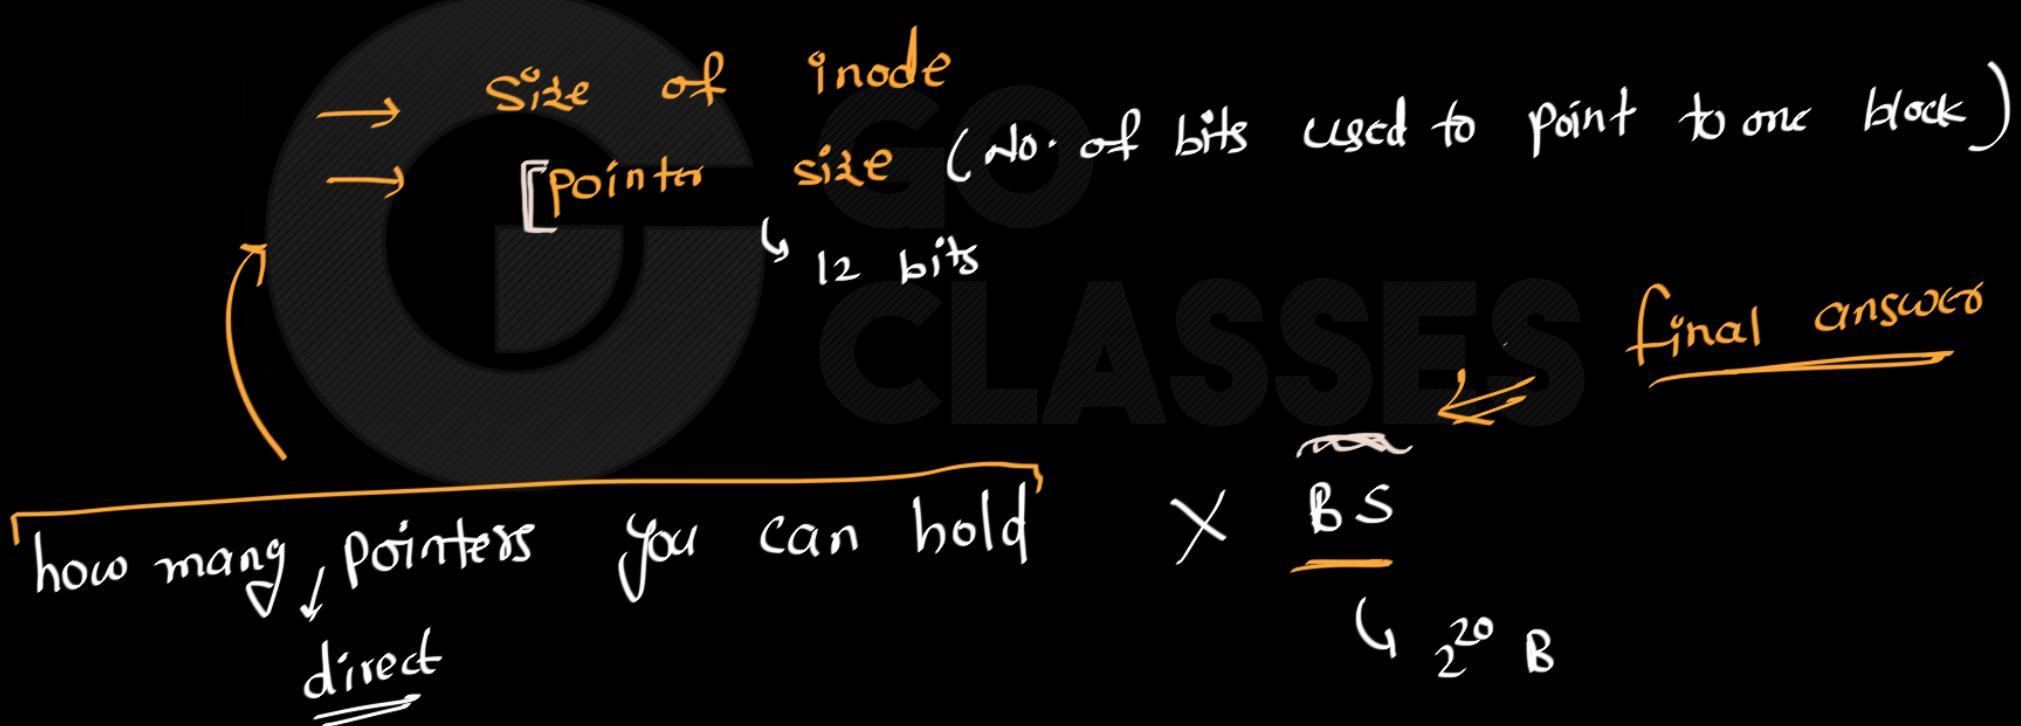


## GATE CSE 2002 | Question: 2.22

- 30 In the index allocation scheme of blocks to a file, the maximum possible size of the file depends on  
A. the size of the blocks, and the size of the address of the blocks.  
B. the number of blocks used for the index, and the size of the blocks.  
C. the size of the blocks, the number of blocks used for the index, and the size of the address of the blocks.  
D. None of the above



$$\underbrace{\text{Disk Size}}_{=} = 2^{12} \text{ blocks} \times 2^{20} B$$





file = inode + Data

$$\underline{\text{Data size}} = \underline{\text{Disk size}} - \underline{\text{inode size}}$$



In Index allocation size of maximum file can be derived like following:

52

No of addressable blocks using one Index block ( $A$ ) = Size of block / Size of block address

$\} \leftarrow$  How many  
pointers can  
inode  
can hold



No of block addresses available for addressing one file ( $B$ ) =



No of Maximum blocks we can use for the Index \* No of addressable blocks using one Index block ( $A$ )

Best answer

Size of File =  $B * \text{Size of Block}$

So, it is clear that:

Answer is (C).

**A & B** are incomplete.

## GATE CSE 2008 | Question: 20



The data blocks of a very large file in the Unix file system are allocated using

- 30  
④

- A. continuous allocation
- B. linked allocation
- C. indexed allocation
- D. an extension of indexed allocation

→ multi level indexing

## GATE CSE 2008 | Question: 20



The data blocks of a very large file in the Unix file system are allocated using

30

- A. continuous allocation
- B. linked allocation
- C. indexed allocation
- ~~D. an extension of indexed allocation~~



→ multi level indexing

## GATE IT 2005 | Question: 63

asked in Operating System Nov 4, 2014 • edited Jun 23, 2018 by kenzou

4,630 views

- 28 In a computer system, four files of size 11050 bytes, 4990 bytes, 5170 bytes and 12640 bytes need to be stored. For storing these files on disk, we can use either 100 byte disk blocks or 200 byte disk blocks (but can't mix block sizes).  
For each block used to store a file, 4 bytes of bookkeeping information also needs to be stored on the disk. Thus, the total space used to store a file is the sum of the space taken to store the file and the space taken to store the book keeping information for the blocks allocated for storing the file. A disk block can store either bookkeeping information for a file or data from a file, but not both.  
What is the total space required for storing the files using 100 byte disk blocks and 200 byte disk blocks respectively?

- A. 35400 and 35800 bytes
- B. 35800 and 35400 bytes
- C. 35600 and 35400 bytes
- D. 35400 and 35600 bytes

116 blocks ( $BS = 100$ )

11050

4990

5170

12640

$BS = \frac{100B}{\text{or}} \frac{200B}{\text{_____}}$

H.W.

{ file size = 11050 bytes

} Block size = 100 bytes

↓ No. of blocks =  $\frac{11050}{100} = \underline{\underline{111 \text{ blocks}}}$

for each block we need 4 bytes extra.



total space =  444  
 111 block +  111 × 4 bytes  
 = 111 blocks + 5 blocks  
 = 116 blocks

11050 → 116 blocks

4990 → 52 blocks

5170 →  $52 + 3 = 55$  blocks

12640 →  $127 + 6 = 133$  blocks

(ii) 4990 → 50 blocks for data  
 $50 \times 4$  bytes for book keeping  $= 200 B = 2$  blocks

52

$$\frac{127 \times 4}{100} =$$



for 100 bytes block:

53

$11050 = 111$  blocks requiring  $111 \times 4 = 444$  bytes of bookkeeping info which requires another 5 disk blocks. So,

totally  $111 + 5 = 116$  disk blocks. Similarly,

$$4990 = 50 + (50 \times 4)/100 = 52$$

$$5170 = 52 + (52 \times 4)/100 = 55$$

$$12640 = 127 + (127 \times 4/100) = 133$$

Best answer

----

$$356 \times 100 = 35600 \text{ bytes}$$

For 200 bytes block:

$$56 + (56 \times 4/200) = 58$$

$$25 + (25 \times 4/200) = 26$$

$$26 + (26 \times 4/200) = 27$$

$$64 + (64 \times 4/200) = 66$$

----

$$177 \times 200 = 35400$$

So, (C) option.

## GATE CSE 2017 Set 2 | Question: 08



37



In a file allocation system, which of the following allocation scheme(s) can be used if no external fragmentation is allowed?

1. Contiguous
  2. Linked
  3. Indexed
- 
- A. 1 and 3 only
  - B. 2 only
  - C. 3 only
  - D. 2 and 3 only





41



Best  
answer

Both Linked and Indexed allocation free from external fragmentation

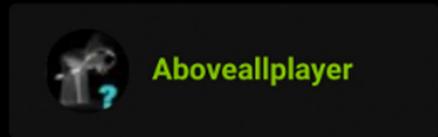
D is answer

Refer: Galvin

Reference: <https://webservices.ignou.ac.in/virtualcampus/adit/course/cst101/block4/unit4/cst101-bl4-u4-06.htm>

answered Feb 14, 2017 • edited Jun 22, 2021 by Lakshman Patel RJIT

comment Follow share this



CLASSES

## GATE CSE 2019 | Question: 42



13



- The index node (inode) of a Unix-like file system has 12 direct, one single-indirect and one double-indirect pointers. The disk block size is 4 kB, and the disk block address is 32-bits long. The maximum possible file size is (rounded off to 1 decimal place) \_\_\_\_ GB

gatecse-2019

numerical-answers

operating-system

file-system

H.W.  
      

CLASSES



Given 12 direct, 1 single indirect, 1 double indirect pointers

23

Size of Disk block =  $4kB$



Disk Block Address = 32 bit =  $4B$



Number of addresses = Size of disk block/address size =  $\frac{4kB}{4B} = 2^{10}$

Best  
answer

Maximum possible file size =  $12 * 4kB + 2^{10} * 4kB + 2^{10} * 2^{10} * 4kB$

=  $4.00395 GB \simeq 4 GB$

Hence  $4GB$  is the correct answer



6



- Consider a linear list based directory implementation in a file system. Each directory is a list of nodes, where each node contains the file name along with the file metadata, such as the list of pointers to the data blocks. Consider a given directory foo.
- Which of the following operations will necessarily require a full scan of foo for successful completion?

- A. Creation of a new file in foo
- B. Deletion of an existing file from foo
- C. Renaming of an existing file in foo
- D. Opening of an existing file in foo



Useful answer



Correct Options: A, C

15



Best  
answer

**Note:** In the question it's given "which of the following options require a full scan of foo for successful completion". Meaning the best algorithm scans the list entirely for each type of input to verify the correctness of the procedure and ,can't partially scan and complete for any particular instance...)

Each File in Directory is uniquely referenced by its **name**. So **different files** must have **different names!**

So,

A. **Creation of a New File:** For creating new file, we've to check whether the new name is same as the existing files. Hence, the linked list must be scanned in its entirety.

B. **Deletion of an Existing File:** Deletion of a file doesn't give rise to name conflicts, hence if the node representing the files is found earlier, it can be deleted without a through scan.

C. **Renaming a File:** Can give rise to name conflicts, same reason can be given as option A.

D. **Opening of existing file:** same reason as option B.

## GATE IT 2004 | Question: 67



30



In a particular Unix OS, each data block is of size 1024 bytes, each node has 10 direct data block addresses and three additional addresses: one for single indirect block, one for double indirect block and one for triple indirect block. Also, each block can contain addresses for 128 blocks. Which one of the following is approximately the maximum size of a file in the file system?

- A. 512 MB
- B. 2 GB
- C. 8 GB
- D. 16 GB





Answer: (B)

38

Maximum file size =  $10 \times 1024 \text{ Bytes} + 1 \times 128 \times 1024 \text{ Bytes} + 1 \times 128 \times 128 \times 1024 \text{ Bytes}$   
 $+ 1 \times 128 \times 128 \times 128 \times 1024 \text{ Bytes} = \text{approx } 2 \text{ GB.}$



answered Apr 6, 2015 • edited Jun 26, 2018 by **kenzou**

Best answer

comment Follow share this



Rajarshi Sarkar

# GATE CSE 2022 | Question: 53

asked in **Operating System** Feb 15 • edited Feb 19 by **Anjana5051**

2,002 views

9  
↑  
↓

Consider two file systems A and B, that use contiguous allocation and linked allocation, respectively. A file of size 100 blocks is already stored in A and also in B. Now, consider inserting a new block in the middle of the file (between 50<sup>th</sup> and 51<sup>st</sup> block), whose data is already available in the memory. Assume that there are enough free blocks at the end of the file and that the file control blocks are already in memory. Let the number of disk accesses required to insert a block in the middle of the file in A and B are  $n_A$  and  $n_B$ , respectively, then the value of  $n_A + n_B$  is\_\_\_\_\_.

gatecse-2022

numerical-answers

operating-system

file-system

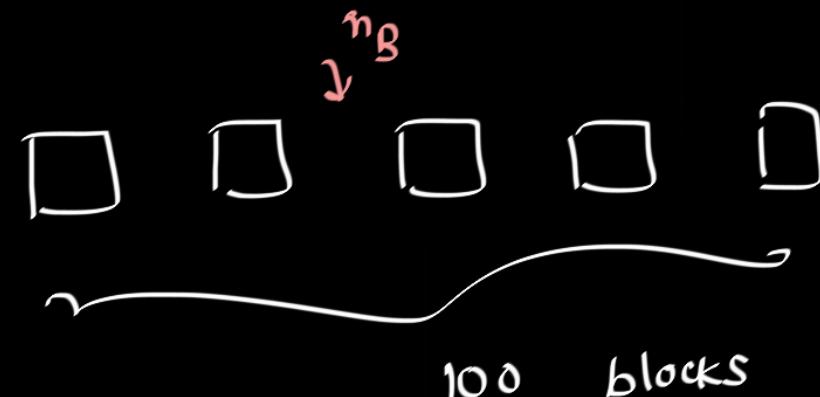
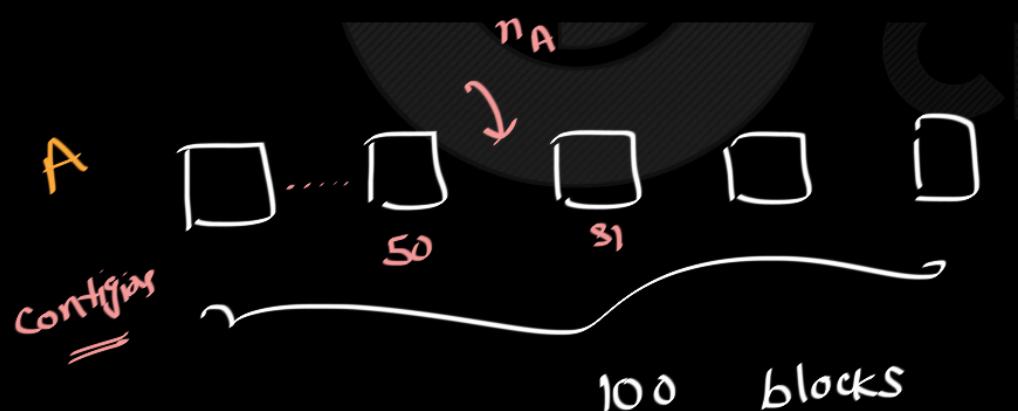
# GATE CSE 2022 | Question: 53

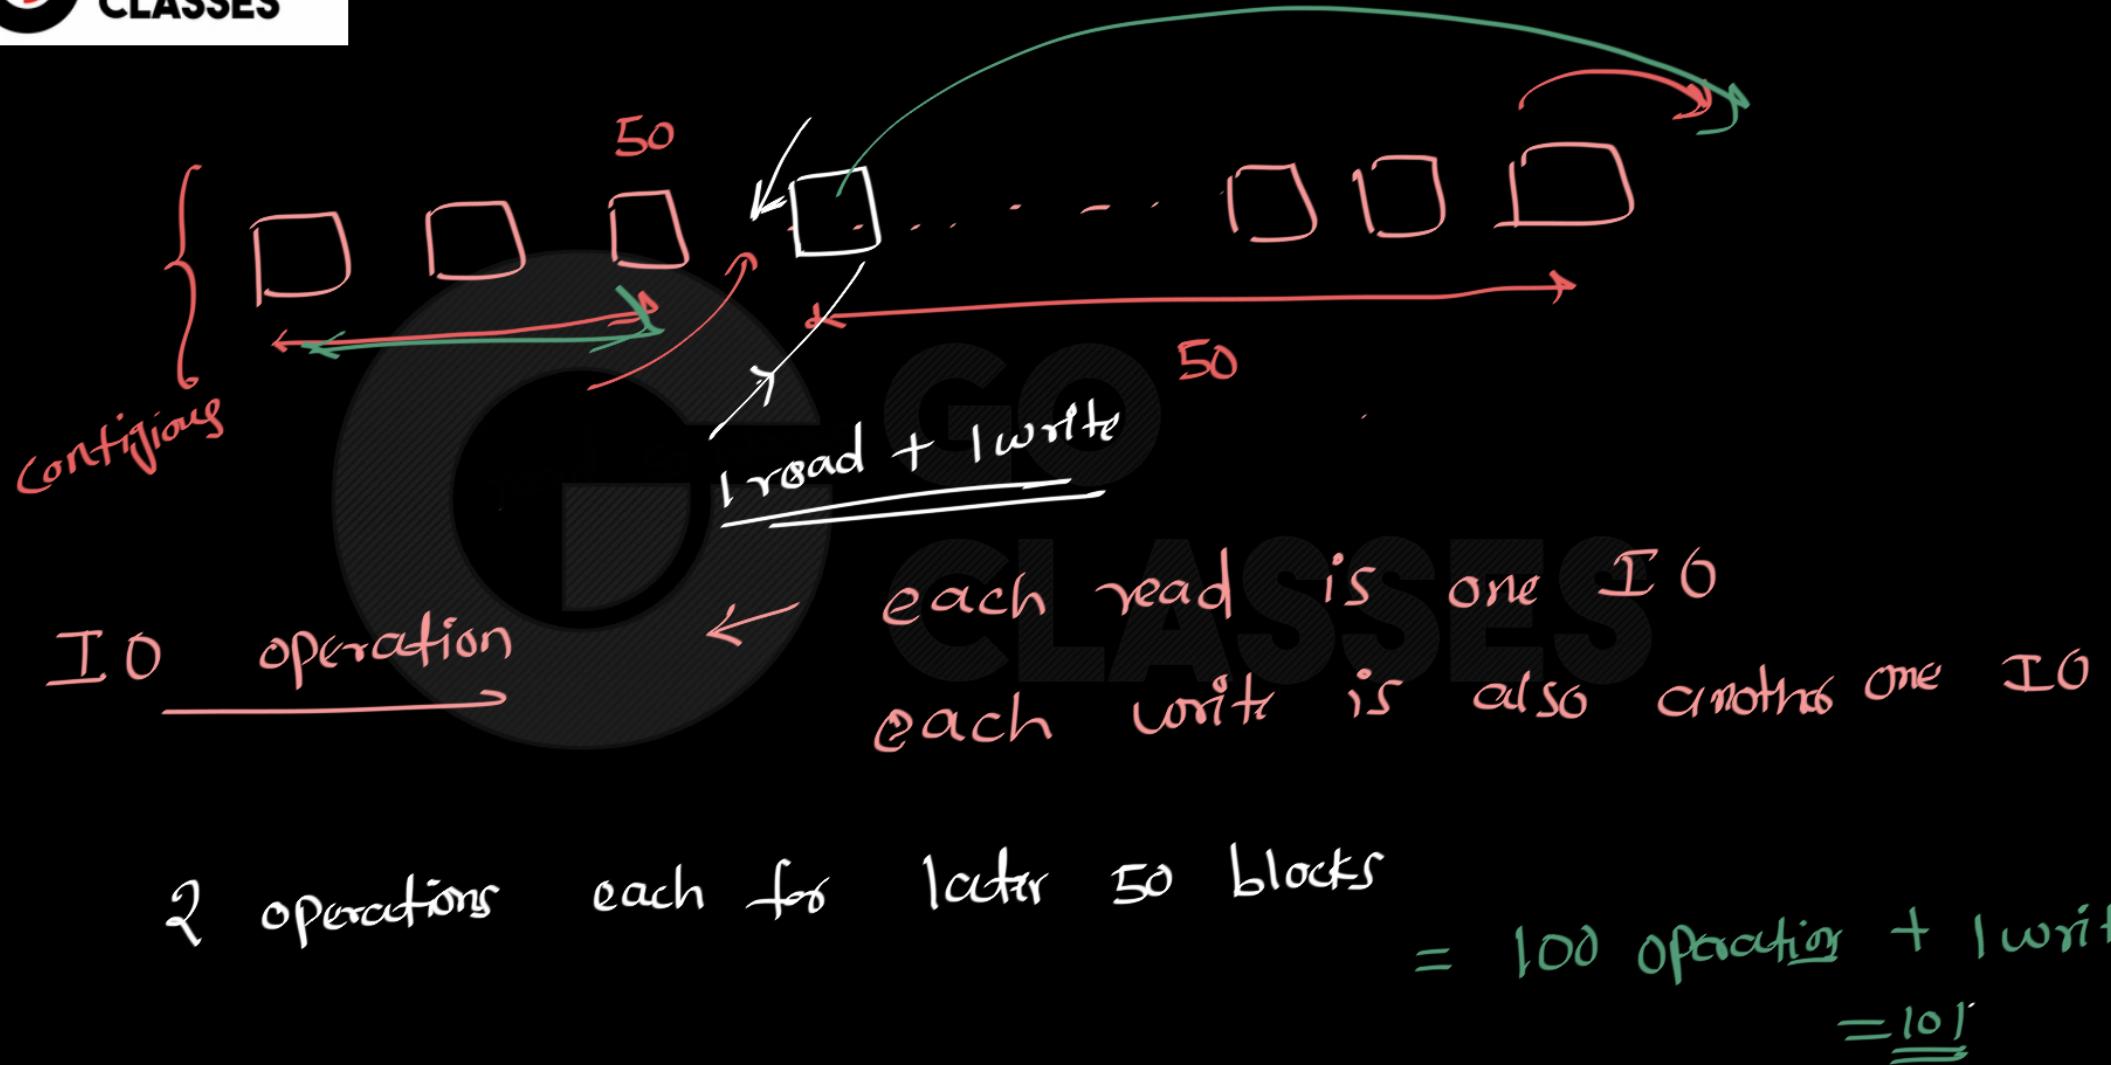
asked in Operating System Feb 15 • edited Feb 19 by Anjana5051

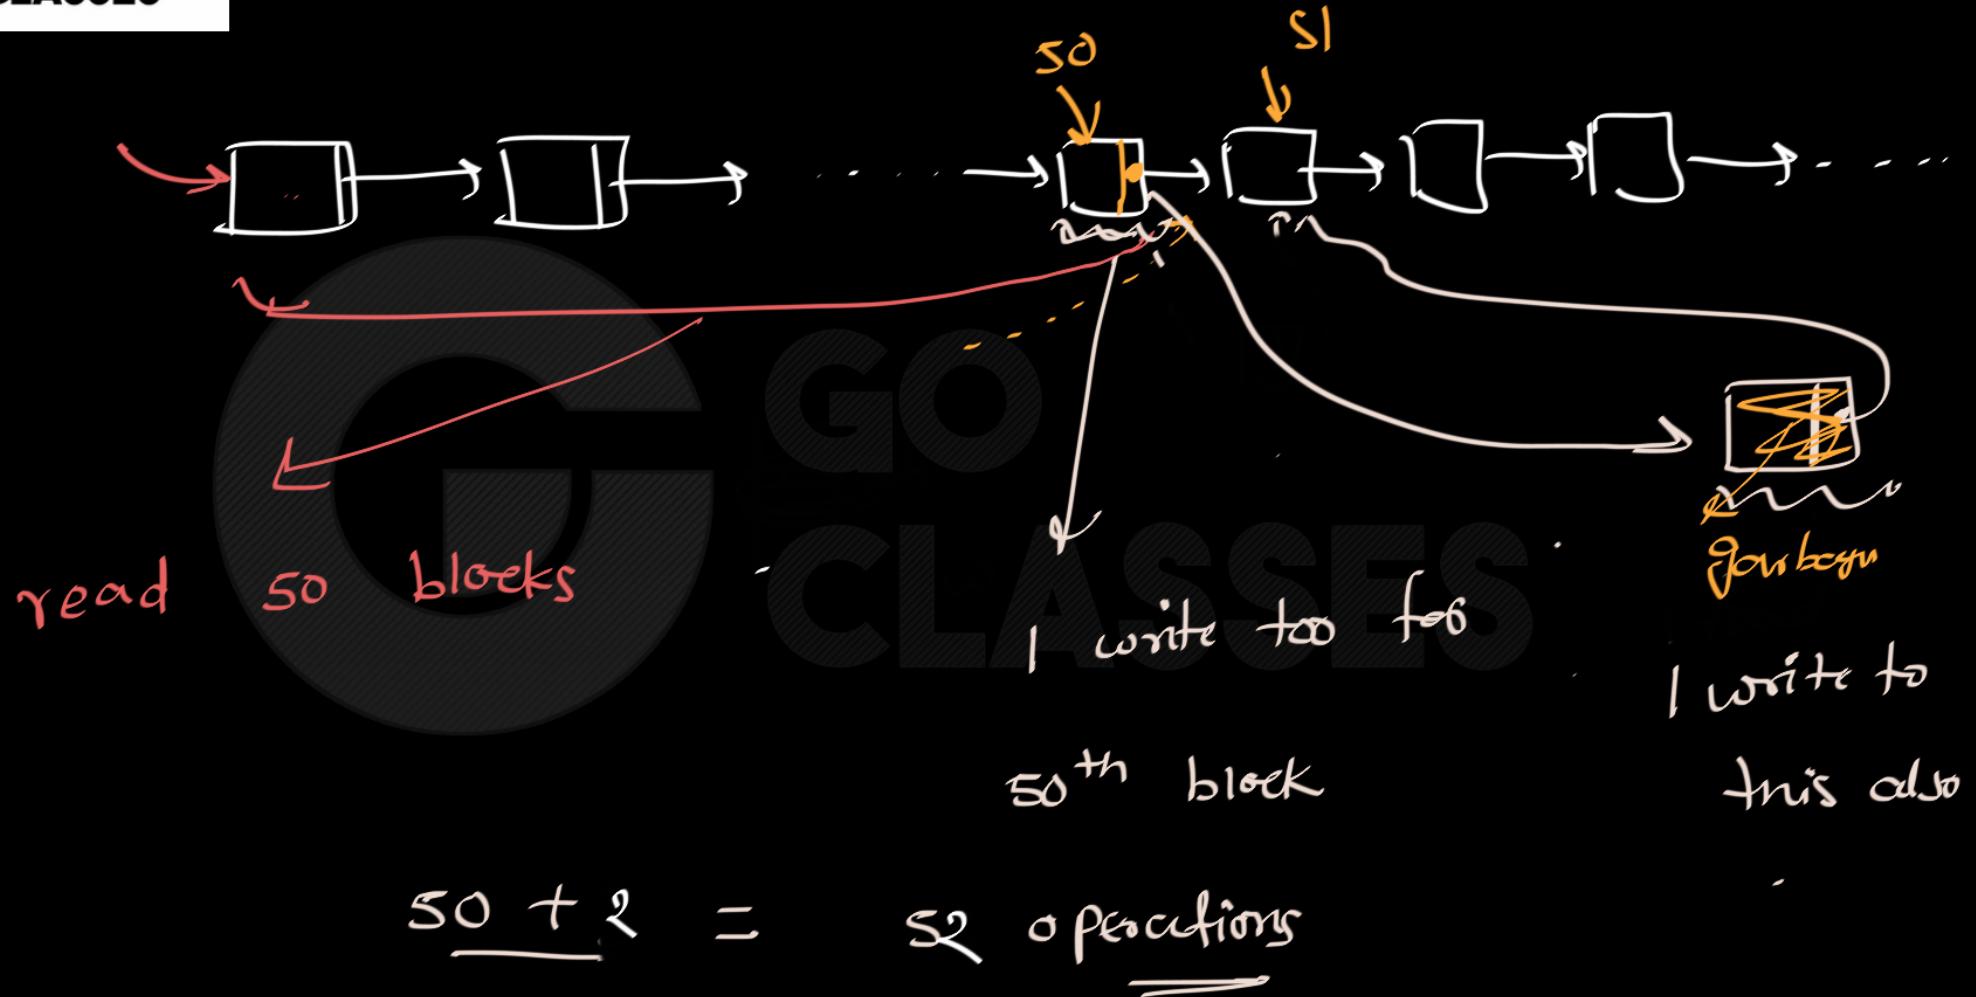
2,002 views

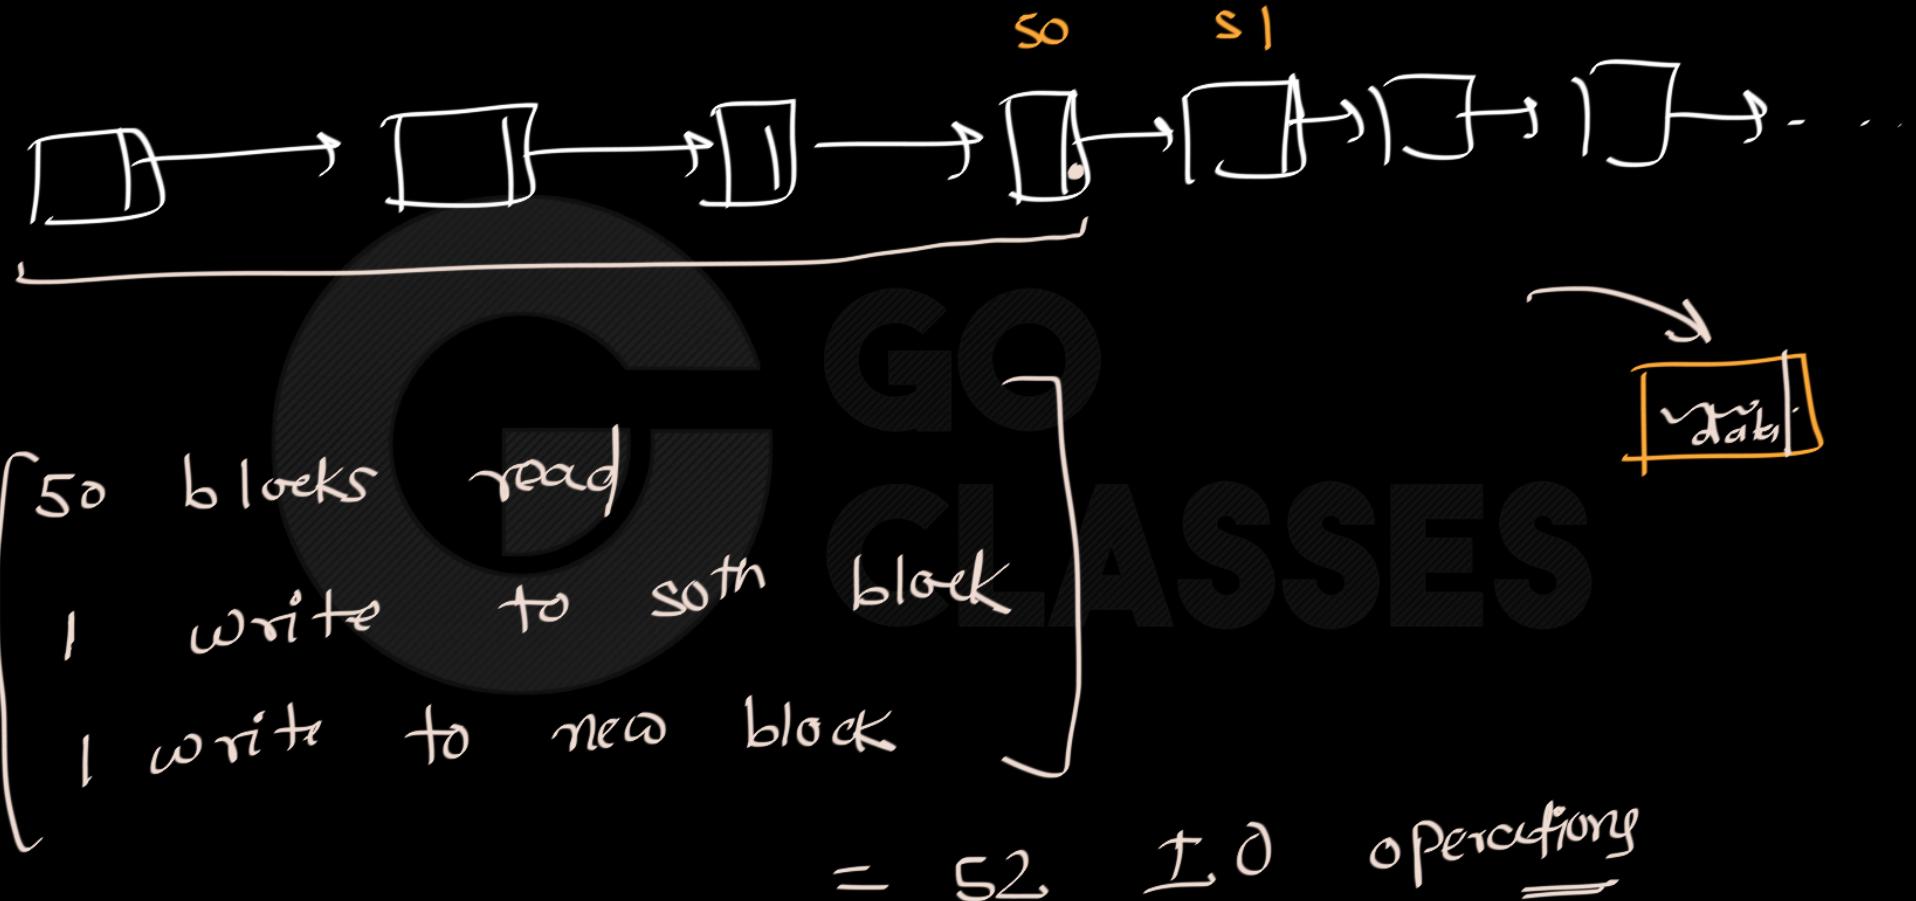
9  
Upvote  
Downvote

Consider two file systems A and B, that use contiguous allocation and linked allocation, respectively. A file of size 100 blocks is already stored in A and also in B. Now, consider inserting a new block in the middle of the file (between 50<sup>th</sup> and 51<sup>st</sup> block), whose data is already available in the memory. Assume that there are enough free blocks at the end of the file and that the file control blocks are already in memory. Let the number of disk accesses required to insert a block in the middle of the file in A and B are  $n_A$  and  $n_B$ , respectively, then the value of  $n_A + n_B$  is\_\_\_\_\_.











11

Answer: 153Explanation:

-No Free blocks- 1, 2, 3,..,49, 50, New Block, 51, 52, ...,99, 100, ... Free Blocks ...

### Contiguous Allocation:

In case of **Contiguous** allocation we can directly go to the 50<sup>th</sup> element. After this, we have to insert a block here, and since the allocation is **Contiguous**, therefore you need to shift all the remaining 50 blocks **to the right**. [As enough free blocks are available to the right and no free blocks are available in the beginning, so we can only shift the blocks to the right only].

So,

$$50 \text{ Read Operations} + 50 \text{ Write Operations} + 1 \text{ [1 operation to write a newly inserted block]} = 101 \text{ OI}$$

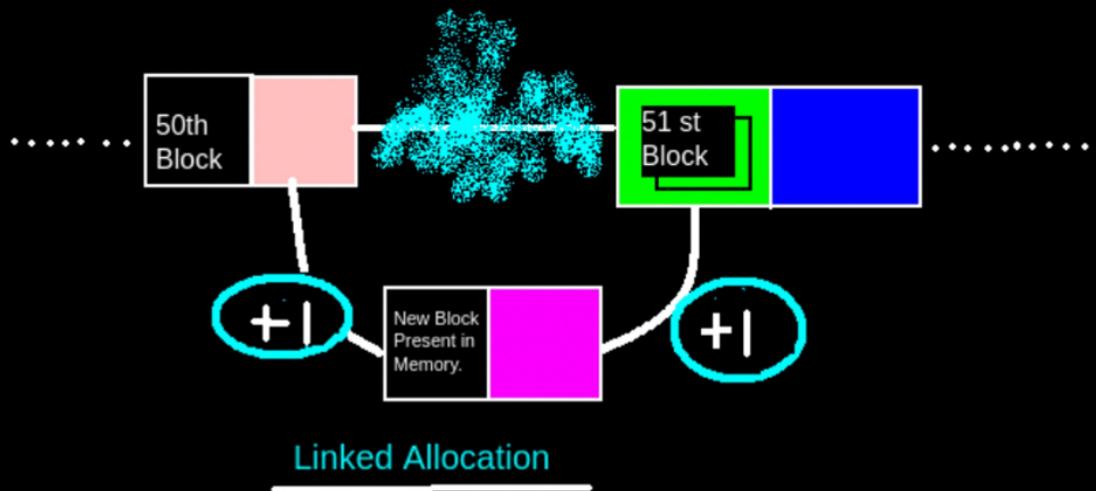
.

Also, we know from the Contiguous Memory allocation concept that overwriting an element simply means deleting it. Therefore, we don't have to worry about deleting an element specifically. We can just overwrite them, thus saving the cost of operations.

### **Linked Allocation:**

In **Linked Allocation**, 50 operations to read first 50 elements , 2 operations are needed to delete the next pointer of the 50<sup>th</sup> element, connect that link to the block which is to be inserted, and then connect the next pointer of that block to the 51<sup>st</sup> element. This takes **2 operations**.

So, Total **52** operations are needed in this case.



ES

### **Note:**

- The statement “The file Control Blocks are already there in memory ” means the information regarding the file structure is already present in the memory. (Eg: index block in the case of indexed allocation is already present in memory.)
- I/O Operations or Operations or Read & Write Operations or Disk Accesses, all of them simply represent the same thing.

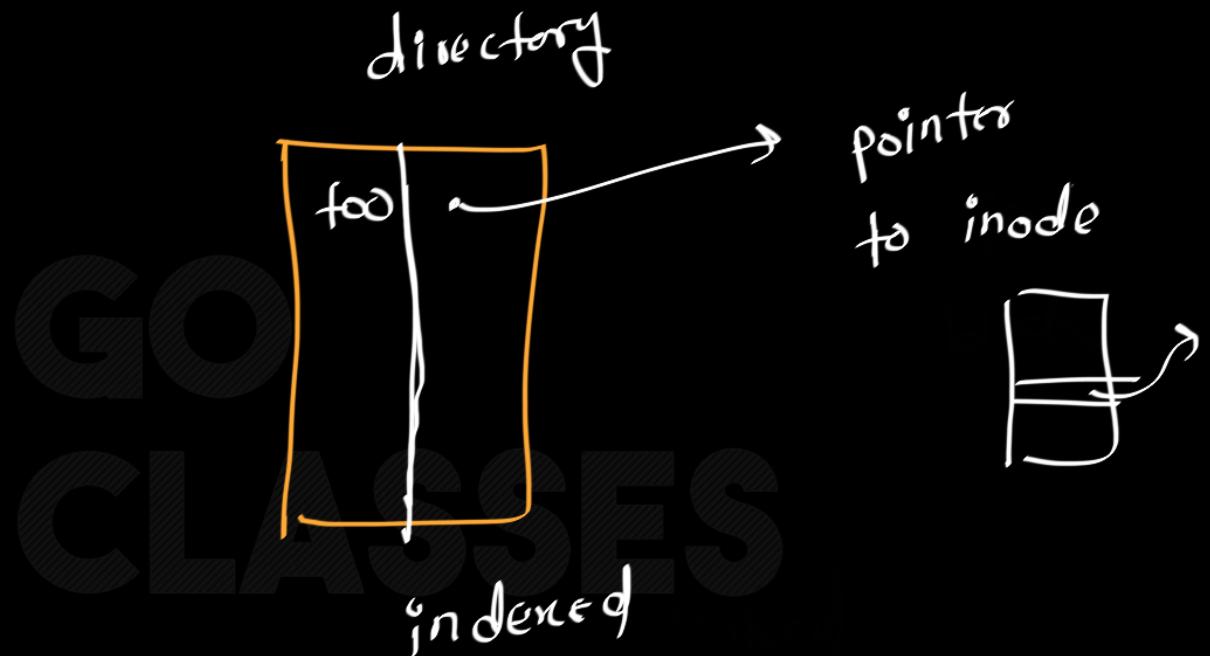
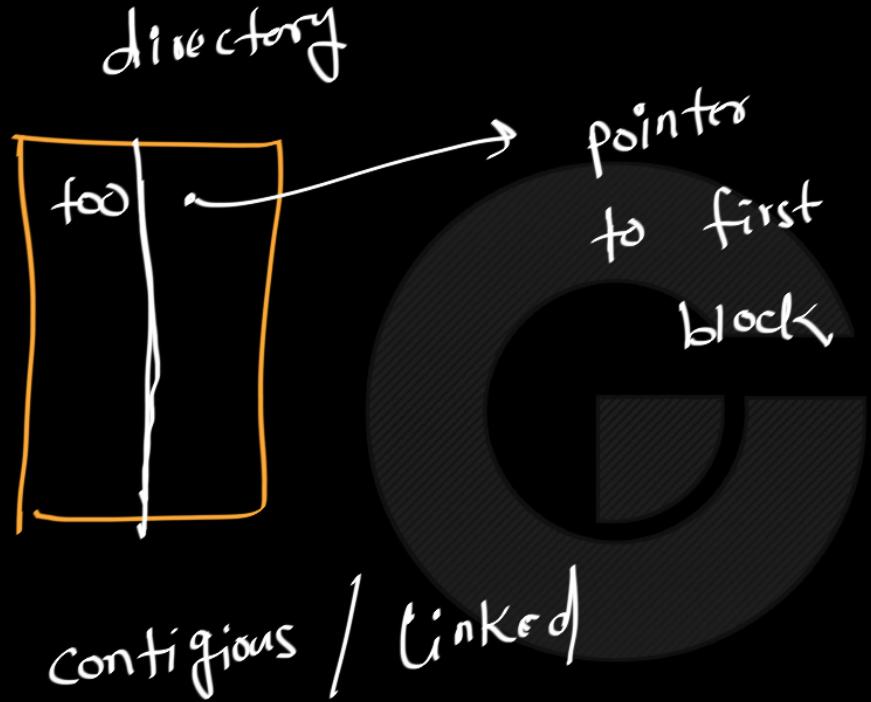
- c. (9 Points) Suppose a file system can have three disk allocation strategies, contiguous, linked, and indexed. We have just read the information for a file from its parent directory. For contiguous and linked allocation, this gives the address of the first block, and for indexed allocation this gives the address of the index block. Now we want to read the 10<sup>th</sup> data block into the memory. How many disk blocks ( $R$ ) do we have to read for each of the allocation strategies? *For partial credit, explicitly list which block(s) you have to read.*

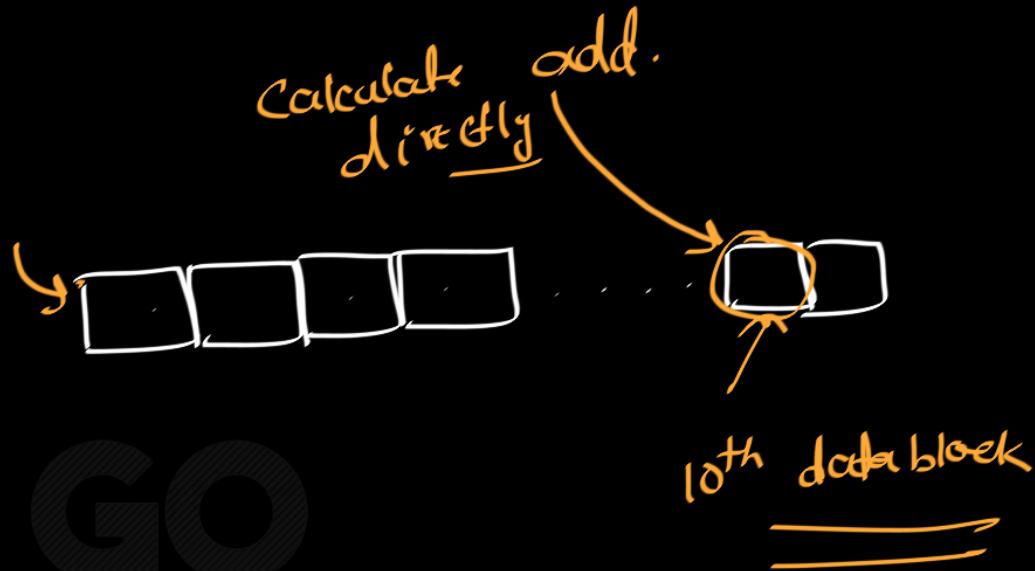
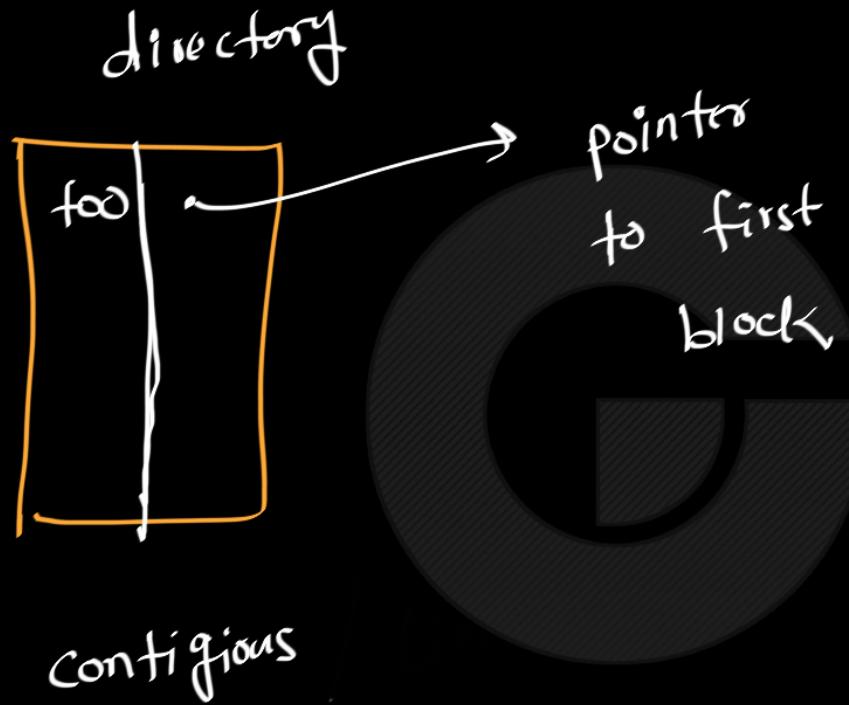
Contiguous allocation:



Linked allocation:

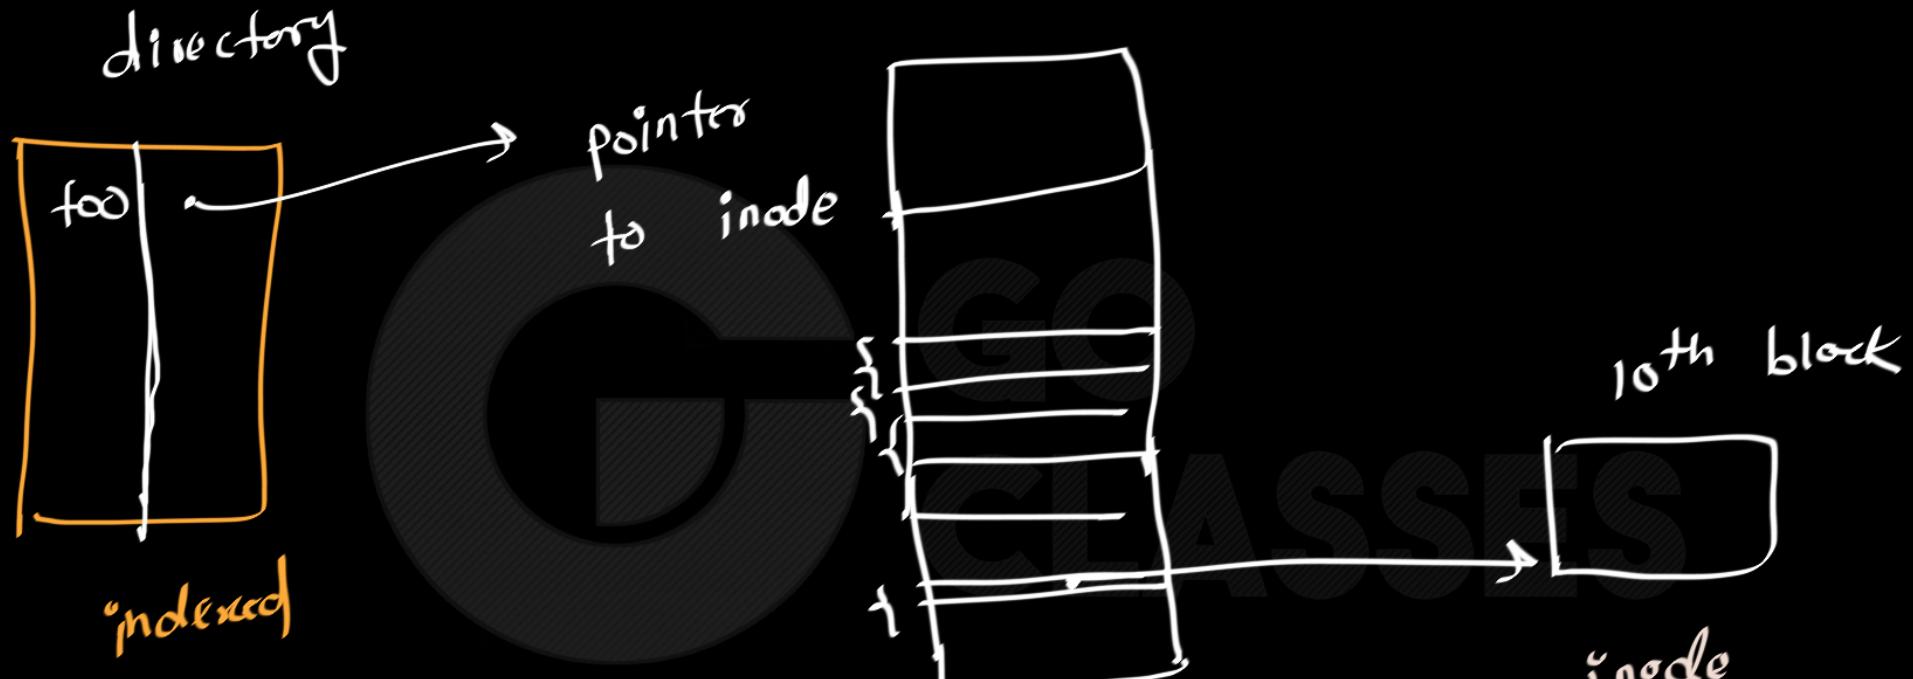
Indexed allocation:





GO  
CLASSES





No. of blocks to read      10th block =      1 + 1

- c. (9 Points) Suppose a file system can have three disk allocation strategies, contiguous, linked, and indexed. We have just read the information for a file from its parent directory. For contiguous and linked allocation, this gives the address of the first block, and for indexed allocation this gives the address of the index block. Now we want to read the 10<sup>th</sup> data block into the memory. How many disk blocks ( $R$ ) do we have to read for each of the allocation strategies? *For partial credit, explicitly list which block(s) you have to read.*

Contiguous allocation:

$$R =$$

1 block

*We took off one point if you included reading a file header.*

*For major errors, we deducted two points.*

Linked allocation:

$$R =$$

10 blocks

*We took off one point if you included reading a file header.*

Indexed allocation:

$$R =$$

2 blocks

*We took off one point if you did not read the index block. We took off two points if you assumed the index block was already in memory.*



7. [15 points] Consider a file currently consisting of 100 blocks. Assume that the file-control block (and the index block, in the case of indexed allocation) is already in memory. Calculate how many disk I/O operations are required for contiguous, linked, and indexed (single-level) allocation strategies, if, for one block, the following conditions hold. In the contiguous-allocation case, assume that there is no room to grow at the beginning but there is room to grows at the end. Also assume that the block information to be added is stored in memory.

- (a) The block is added at the beginning.
- (b) The block is added in the middle.
- (c) The block is added at the end.
- (d) The block is removed from the beginning.
- (e) The block is removed from the middle.
- (f) The block is removed from the end.

H. 10.  
\_\_\_\_\_

Fill out the following table and provide a detailed elaboration. **Without a detailed elaboration, you risk to receive a very low or even 0 points.**

*add middle* →

	<i>Contiguous</i>	<i>Linked</i>	<i>Indexed</i>
(a)			
(b)			
(c)			
(d)			
(e)			
(f)			

ES

Answer: The correct answers are shown in the table below:

(a) The block is added at the beginning.

- i. **Contiguous**: All 100 blocks have to be shifted one block. Because each shift requires one read and one write,  $200 = 100 \times 2$  disk I/O operations are needed. Then, the new block is written to the first location, which requires one write. As a result, the number of disk I/O operations is  $201 = 2 \times 100 + 1$ .
- ii. **Linked**: The address to the first block is found in the directory. Thus, this address is added to the address field of the new block, which is written to an available location. Hence, one 1 disk operation is needed.
- iii. **Indexed**: The pointers in the index block are shifted to make the first position available so that the new block could be recorded there. Therefore, only 1 disk I/O is needed to write the new block to disk.

(b) The block is added in the middle. This is very similar to adding at the beginning.

- i. **Contiguous**: The middle block can be computed from the directory because blocks are allocated in a contiguous way. Then, the 50 blocks in the second half are moved one block, which requires 50 reads and 50 writes. Finally, the new block is written to the original 50th position, and the number of disk I/O operations is  $101 = 2 \times 50 + 1$ .
- ii. **Linked**: We need 50 reads to retrieve the address of the 51th block. The address of the 51th block is added to the address of the new block, which is written to disk with 1 disk I/O operation. The address of the original 50th block has to be updated with the address of the new block. Therefore, the number of disk I/O operations is  $52 = (50 + 1) + 1$ .
- iii. **Indexed**: This is the same as the previous case, and only 1 disk I/O is needed to write the new block to disk.

(c) The block is added at the end.

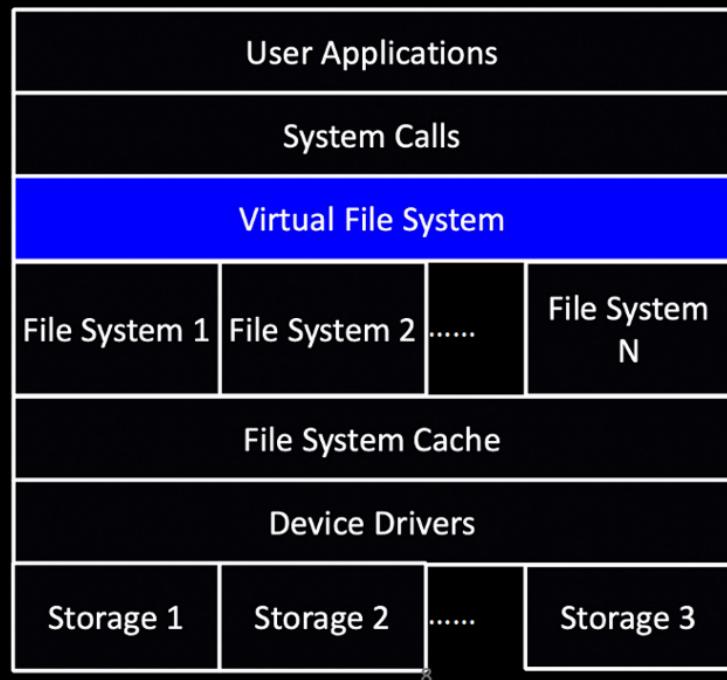
- i. **Contiguous**: One disk I/O operation is needed to write the new block to the end of the original contiguous allocation provided that there is a free block available. Otherwise, compaction is needed.
- ii. **Linked**: First we need to retrieve the last block from directory (1 read). Its address field is updated with the address of the new block, and written back to disk (1 write). One more write is needed to write the new block to disk. Therefore, the number of disk I/O operation is  $3 = 1 + 1 + 1$ .
- iii. **Indexed**: This is the same as the previous case, and only 1 disk I/O is needed to write the new block to disk.

- (d) The block is removed from the beginning.
- Contiguous:** The remaining 99 blocks are moved forward and each move requires 1 read and 1 write, The total number of disk I/O operations is  $198 = 2 \times 99$ .
  - Linked:** Read the first block to find the second one. The second block's address is used to update the directory of this file. Hence, only 1 disk I/O operation is needed.
  - Indexed:** Shifting the addresses stored in the index block one position upward does not need any disk I/O operation.
- (e) The block is removed from the middle.
- Contiguous:** The remaining 49 blocks are moved forward and each move requires 1 read and 1 write, The total number of disk I/O operations is  $98 = 2 \times 49$ .
  - Linked:** This is the same as inserting in the middle. The number of disk I/O operations is 52.
  - Indexed:** Shifting the addresses stored in the index one position upward does not need any disk I/O operation.
- (f) The block is removed from the end.
- Contiguous:** Just modify the directory to reflect the number of blocks is one less than the original. No disk I/O operation is needed.
  - Linked:** We need to read the first 99 blocks, change the address field of the 99th to NULL, and write it back. The number of disk I/O operations is  $100 = 99 + 1$ .
  - Indexed:** Shifting the addresses stored in the index one position upward does not need any disk I/O operation.

# (Not for GATE)

## Virtual File System (VFS)

- VFS provides
  1. A common system call interface to user applications to access different file systems implemented in the OS.
  2. A common interface to file systems to “plug into” the operating system and provide services to user applications.



One OS can support  
multiple FS

# Free Space

- Find the block to use when one is needed
  - Find space quickly
  - Keep storage reasonable
- Options
  - Bit vector
  - Linked List
  - Grouping
  - Counting

using

indexed

contiguous

linked list

[there are free blocks here  
and there]



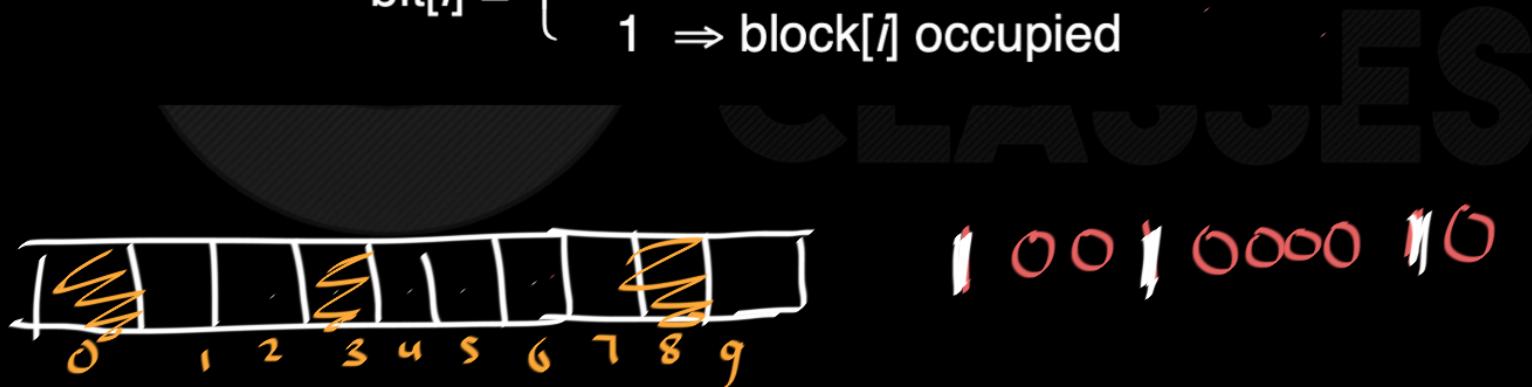
- i should get there free blocks quick
- i should not waste too much of space to store info of free blocks

- Bit vector ( $n$  blocks)

0 1 2 ... n-1

*m*

$$\text{bit}[i] = \begin{cases} 0 & \Rightarrow \text{block}[i] \text{ free} \\ 1 & \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$



- Bit vector downside

- Space

- Example:

block size =  $2^{12}$  bytes  $\leftarrow 4 \text{ kB}$

disk size =  $2^{30}$  bytes (1 gigabyte)

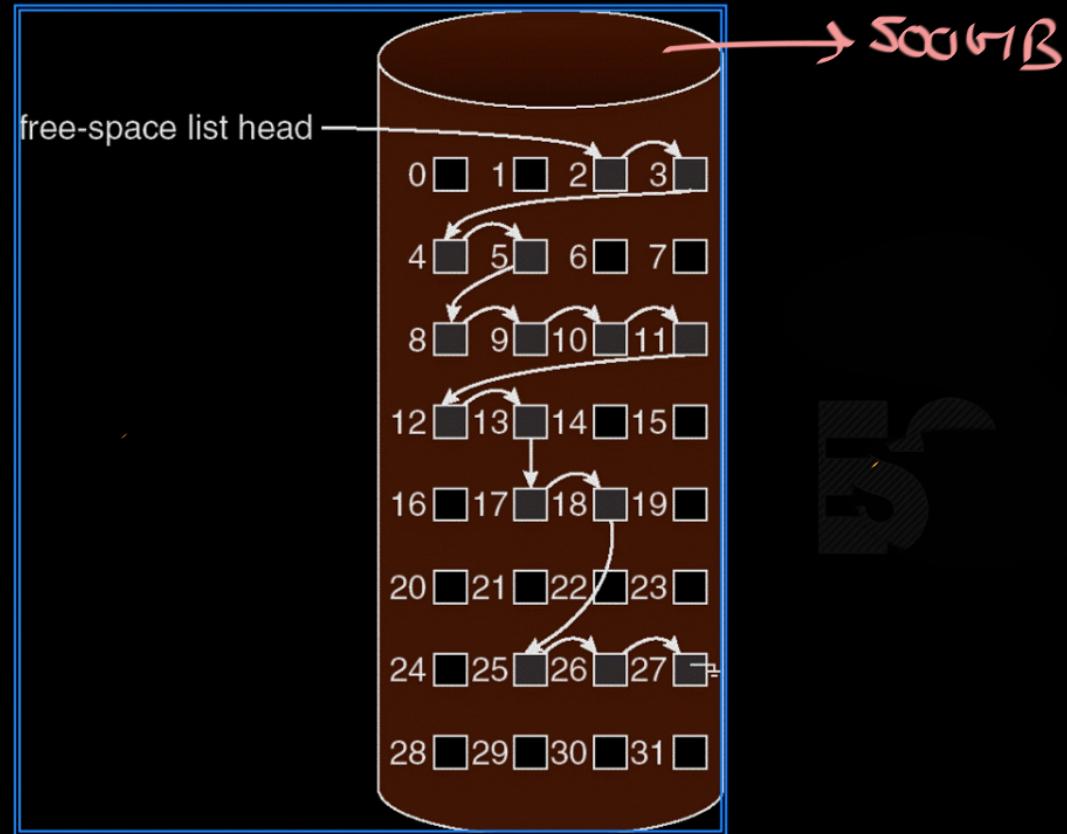
$n = 2^{30}/2^{12} = 2^{18}$  bits (or 32K bytes)

no. of blocks

$32 \text{ kB} \Rightarrow 8 \text{ blocks}$   
 $\equiv$  for bit map

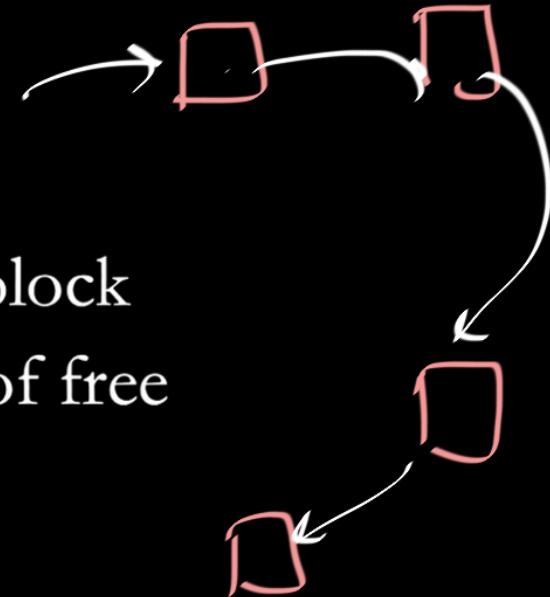
# Free-Space Linked List

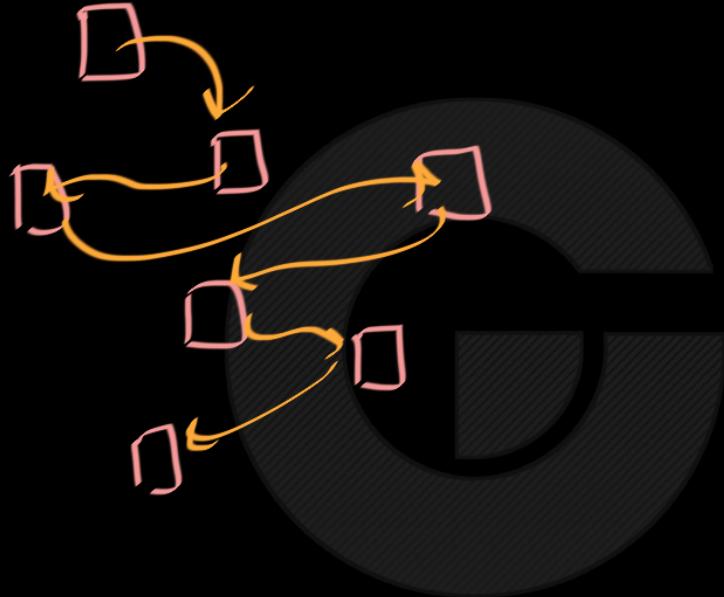
Free list  
→ counting }  
→ grouping }



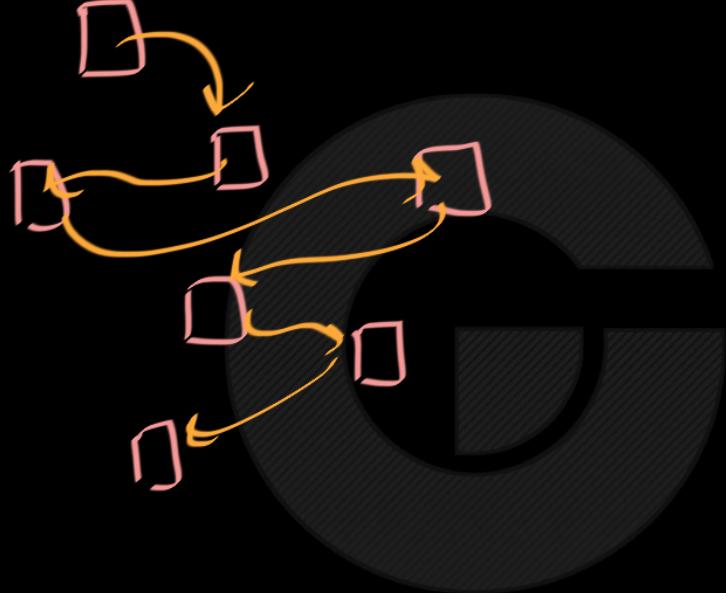
# Free-Space Linked Optimizations

- Grouping
  - Store n free blocks in first free block
  - Last entry points to next block of free blocks
- Counting
  - Specify start block and number of contiguous free blocks

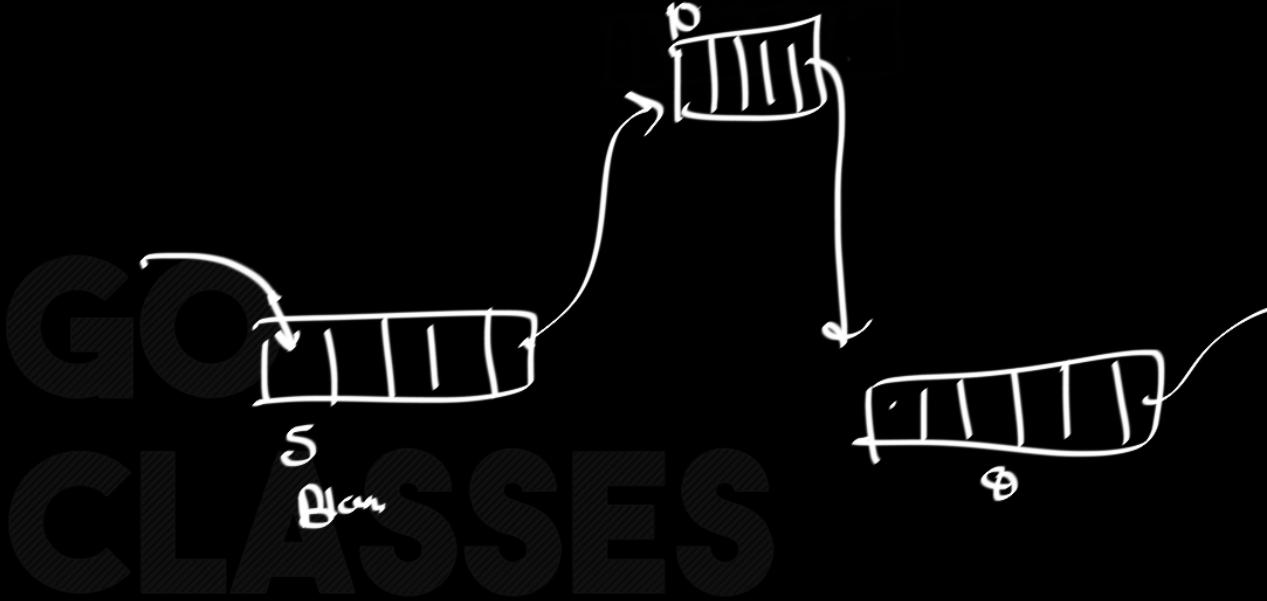




free list



free list



Counting

There are numerous file systems in use; for example, FAT32 and NTFS are Windows file systems, and HFS is used on Macs. Linux uses ext2, ext3 and FAT32. Unix systems use UFS, ext2, ext3 and ZFS.



As such, the file system is highly operating-system-dependent. In most cases you don't generally have a "choice" between different file system types. However, in some operating systems you do, and there can be a performance impact from the choice. Windows NT and 2000 typically give you your choice of file system; Windows 2000 supports FAT16, FAT32 and NTFS.

## **File System Formulas (Types)**

Below some types of file system:

**Hierarchical File System (HFS)**

**File Allocation Table (FAT)**

**New Technology File System (NTFS)**

## **Hierarchical File System (HFS)**

HFS is a file system type developed by Apple Inc. for use on computers running Mac OS. Two main variants of HFS exist: Mac OS Standard (“HFS Standard” or “HFS”) and Mac OS extended (“HFS extended” or “HFS+”). If you are running Mac OS X, your bootable drive is almost certainly using HFS+, not standard HFS. HFS+ allows for larger files with longer file names to be stored on the disk.



## **File Allocation Table (FAT)**

FAT is a brief for File Allocation Table, which dates back to the beginnings of DOS programming. The File Allocation Table (FAT) file system was the primary file system in Microsoft's older operating systems, it is a file system that was created by Microsoft in 1977. FAT was the primary file system used in all of Microsoft's consumer operating systems from MS-DOS through Windows ME.

Over the years, the file system has been expanded from FAT12 to FAT16 and FAT32. Various features have been added to the file system including subdirectories, extended attributes, and long filenames. Below is more information on the versions of the FAT file system:

## New Technology File System (NTFS)

NTFS is a file system type that is commonly used for Microsoft Windows. It is the standard file system for Windows NT, Windows 2000, Windows XP, Windows Vista and Windows 7. It provides numerous improvements over the FAT file system, including better security and better disk utilization. **NTFS** is a proprietary file system developed by Microsoft Corporation for its Windows line of operating systems, beginning with Windows NT 3.1 and Windows 2000, including Windows XP, Windows Server 2003, and all their successors to date.

NTFS offers several features that are not available with FAT 32: file compression, encryption, permissions, and mirroring drives. NTFS supersedes the FAT file system as the preferred file system for Microsoft's Windows operating systems. NTFS has several technical improvements over FAT and HPFS (High Performance File System), such as the use of advanced data structures to improve performance, reliability, and disk space utilization, plus additional extensions, such as security access control lists (ACL).